



# KubeCon 2020 Workshop

## Apache Cassandra and Kubernetes



# Apache Cassandra™ with Kubernetes

1

Housekeeping and Quiz

2

Cassandra Basics

3

Kubernetes Operators

4

Cass Operator in Deep

5

Grafana | Prometheus

6

Resources



# Disclaimer

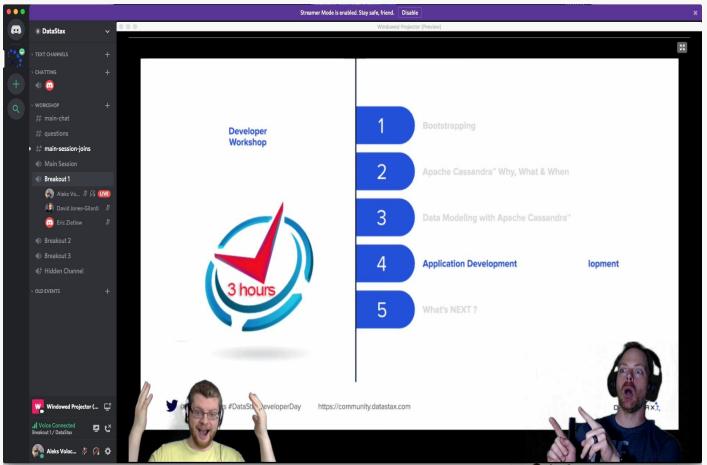
Today,

we won't introduce the basics of Docker and Kubernetes

It's a KubeCon in the end!

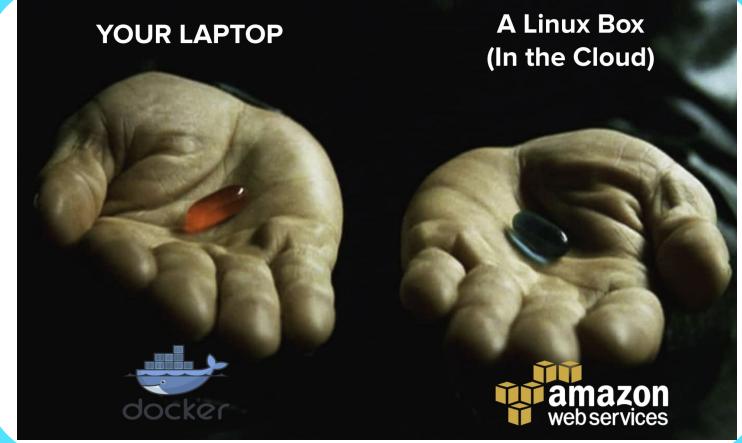


## STREAMS

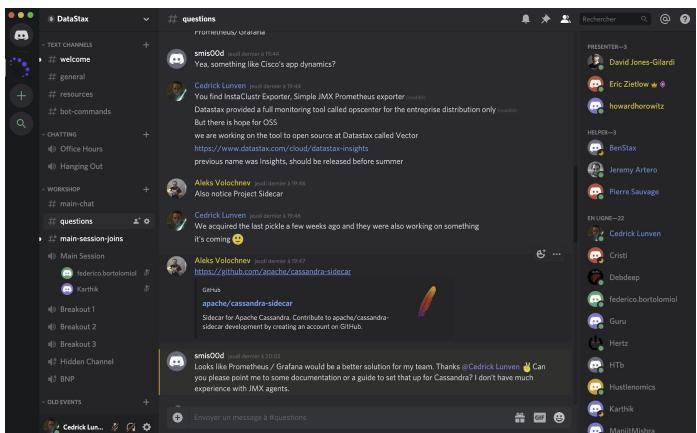


YouTube  
zoom

## RUNTIME



## QUESTIONS



Discord  
zoom

## MATERIALS

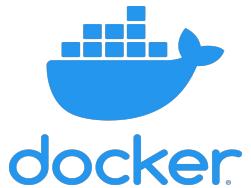
- Slides
- Notebooks
- DockerFile





# YOUR LAPTOP

Your Laptop



docker



kind



kubernetes



# CLOUD INSTANCE



<http://kubecon2020.datastaxtraining.com/>

Cloud Env



kubernetes



# menti.com

# 51 63 695



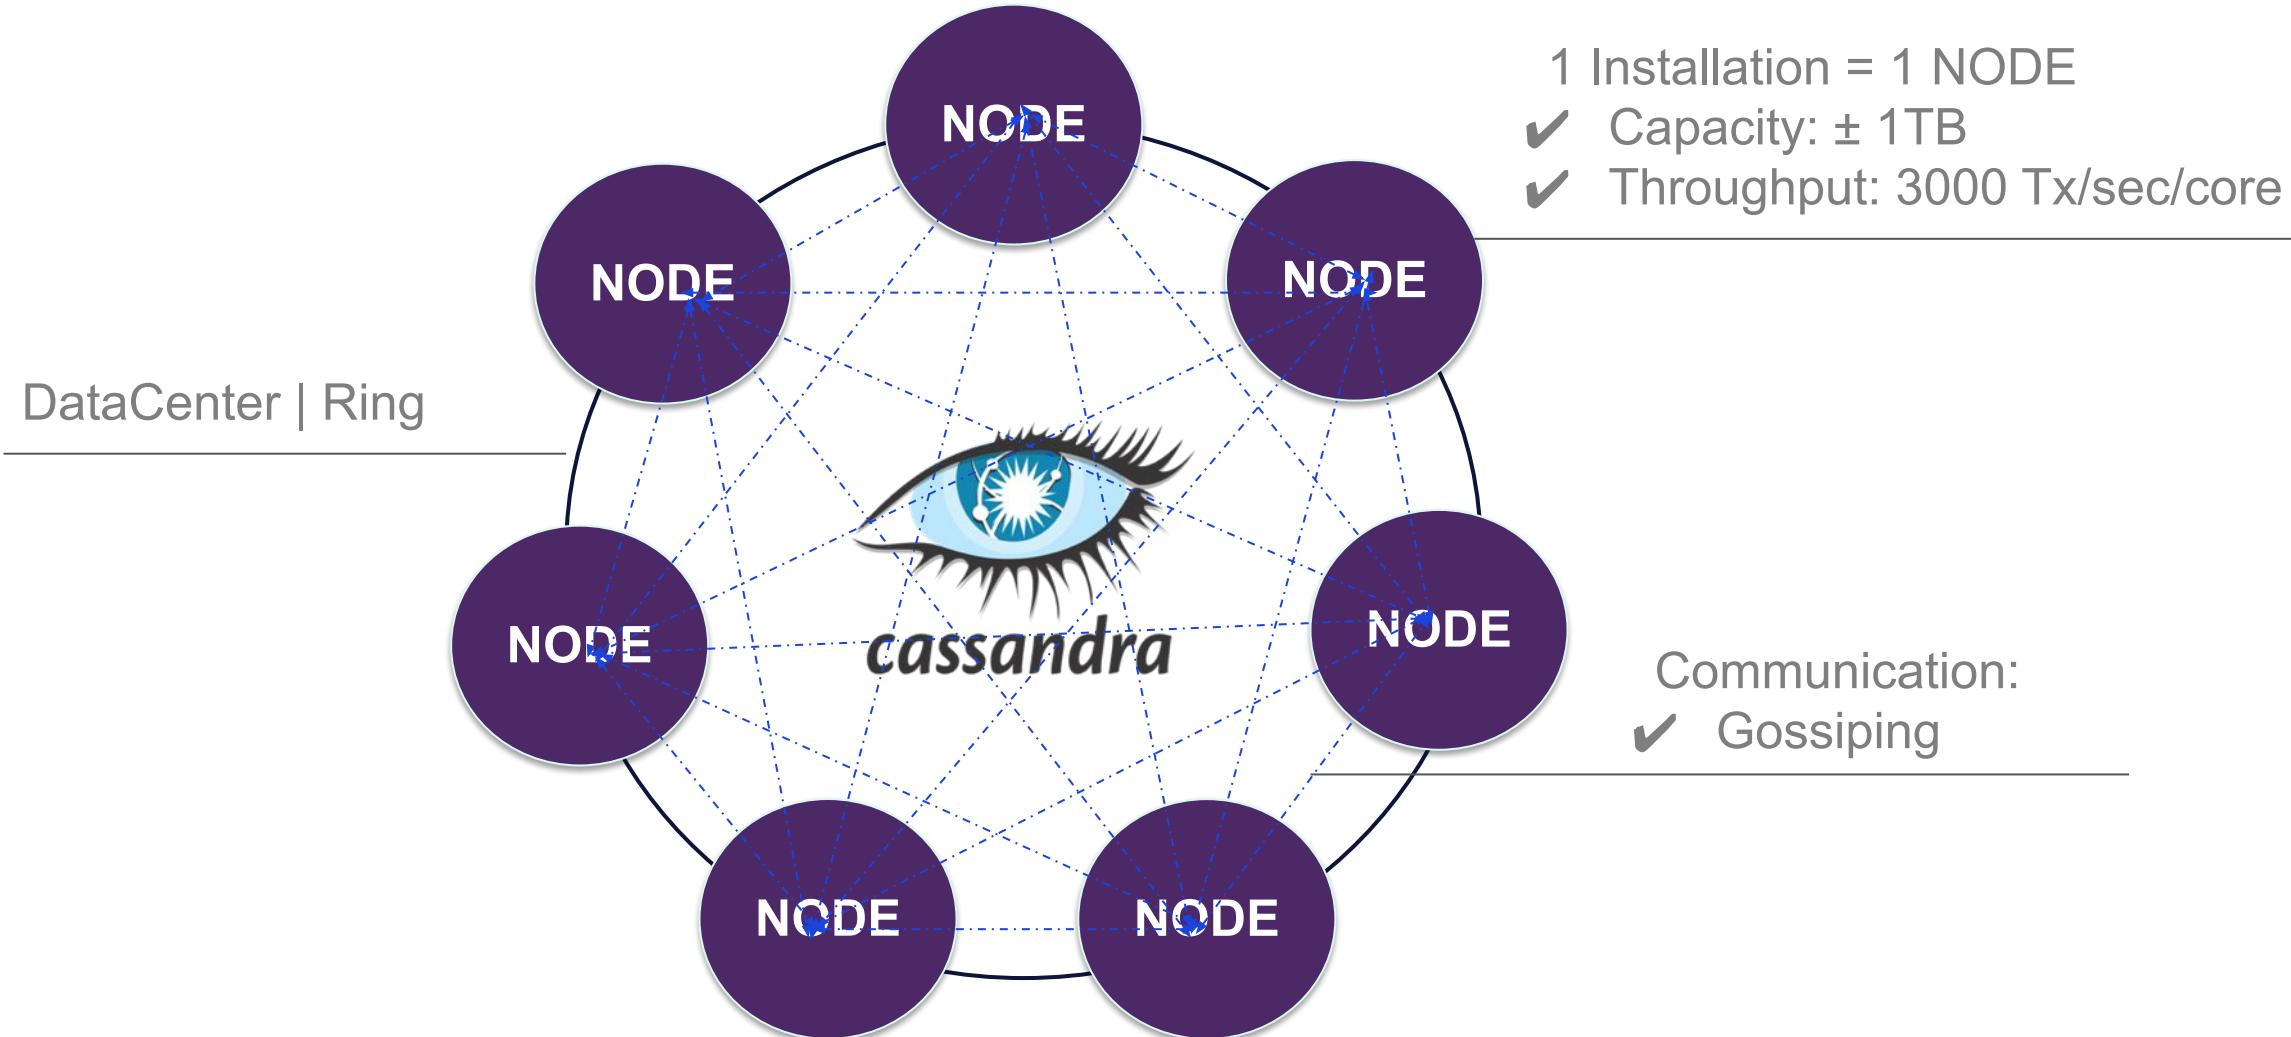
Available on the iPhone  
**App Store**

GET IT ON  
**Google play**

# Apache Cassandra™ with Kubernetes

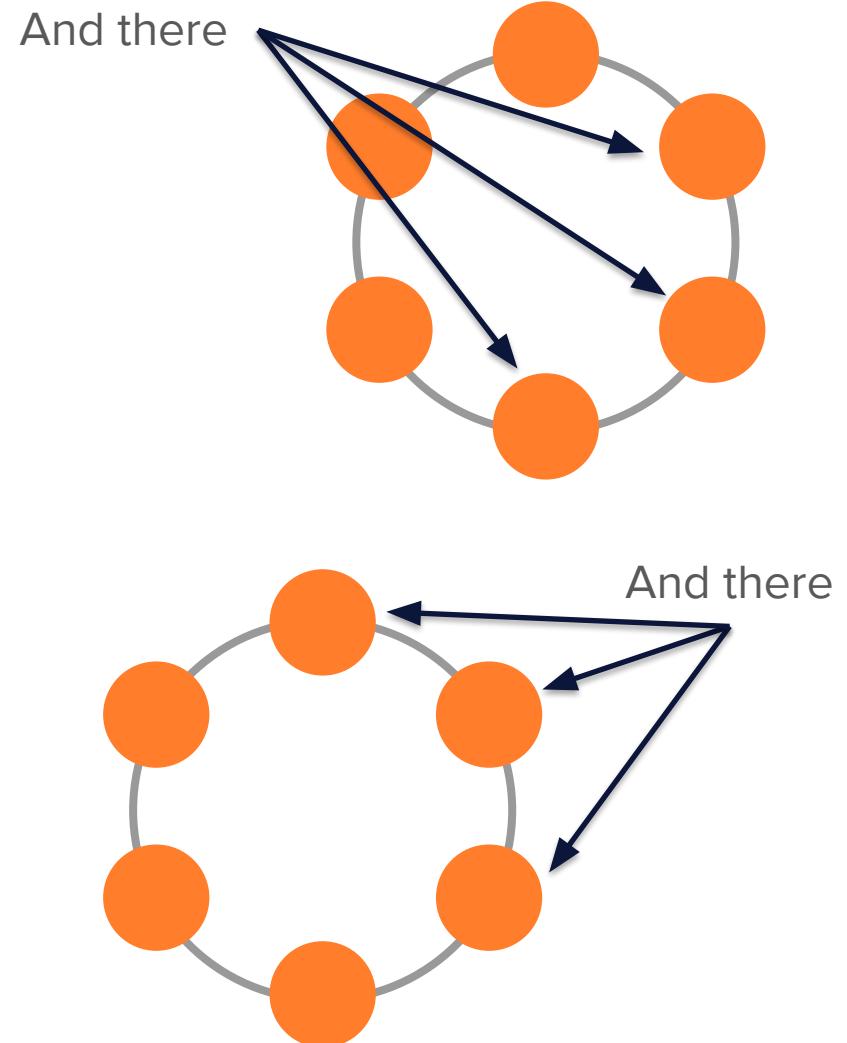
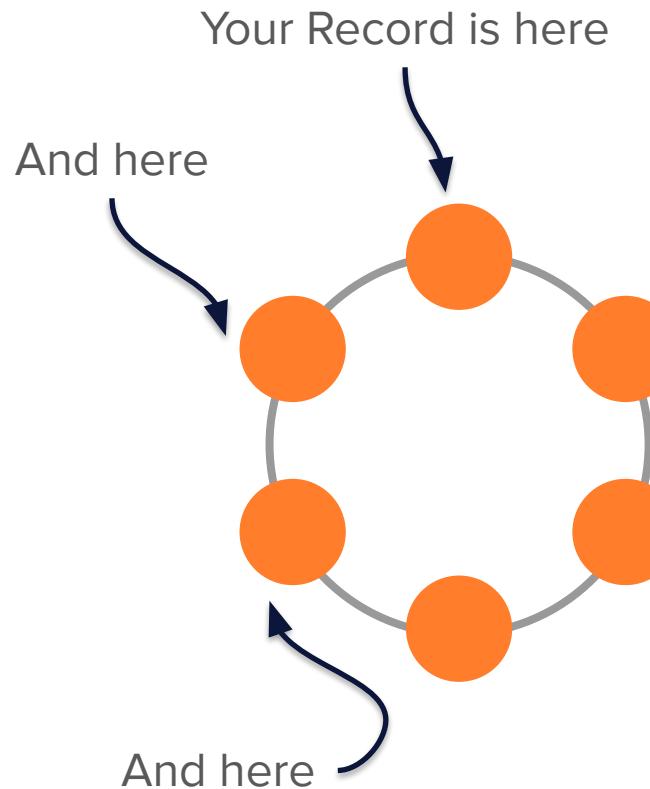
- 1 Housekeeping and Quizz
- 2 Cassandra Basics
- 3 Kubernetes Operators
- 4 Cass Operator in Deep
- 5 Grafana | Prometheus
- 6 Resources

# Apache Cassandra™ = NoSQL Distributed Database



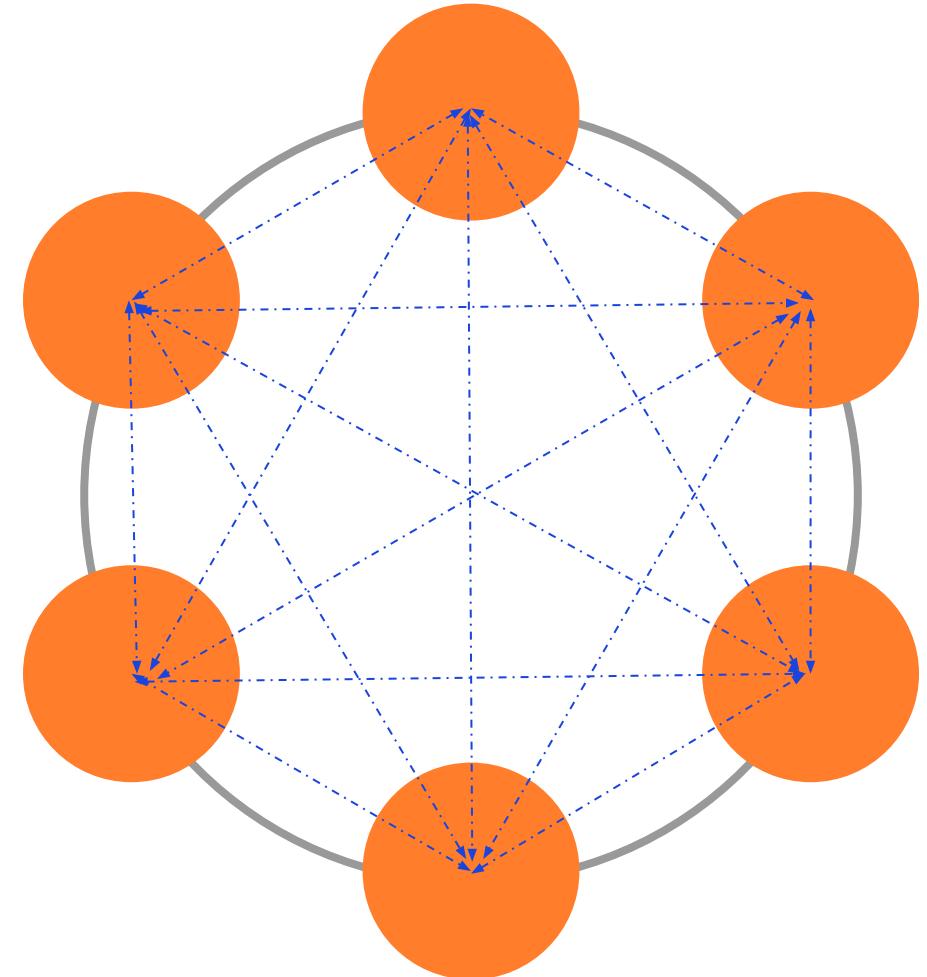
# Cassandra Features

- Distributed
  - Responsive
  - Scalable
- Replicated
  - Available
- Masterless
  - No SPoF



# Cassandra is Distributed and Highly Available

- Nodes organized in a ring
- Nodes monitor each other's health status using a gossip protocol
- Cassandra routes requests away from nodes that are down or slow



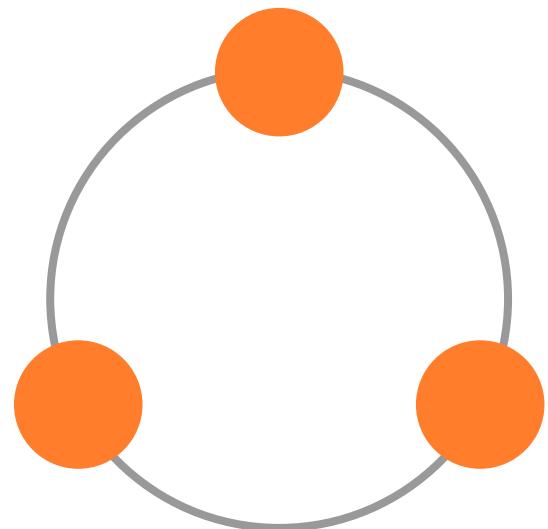
@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>

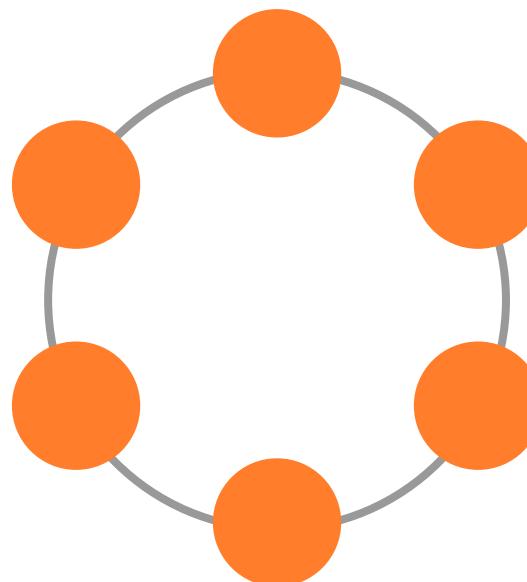


# Horizontal vs. Vertical Scaling

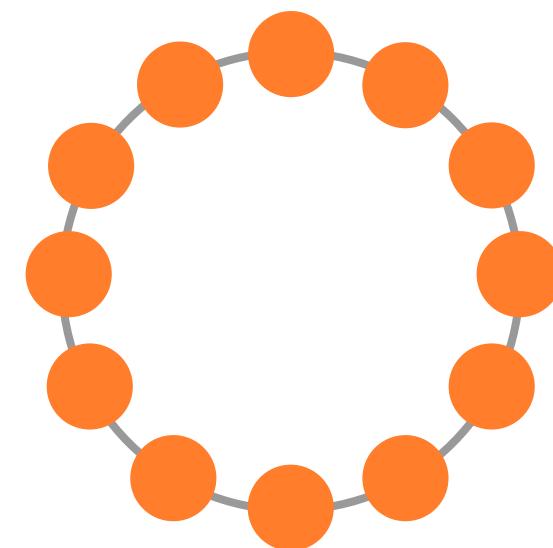
- Vertical scaling requires one large expensive machine
- Horizontal scaling requires multiple less-expensive commodity hardware



100,000 transactions/second



200,000 transactions/second



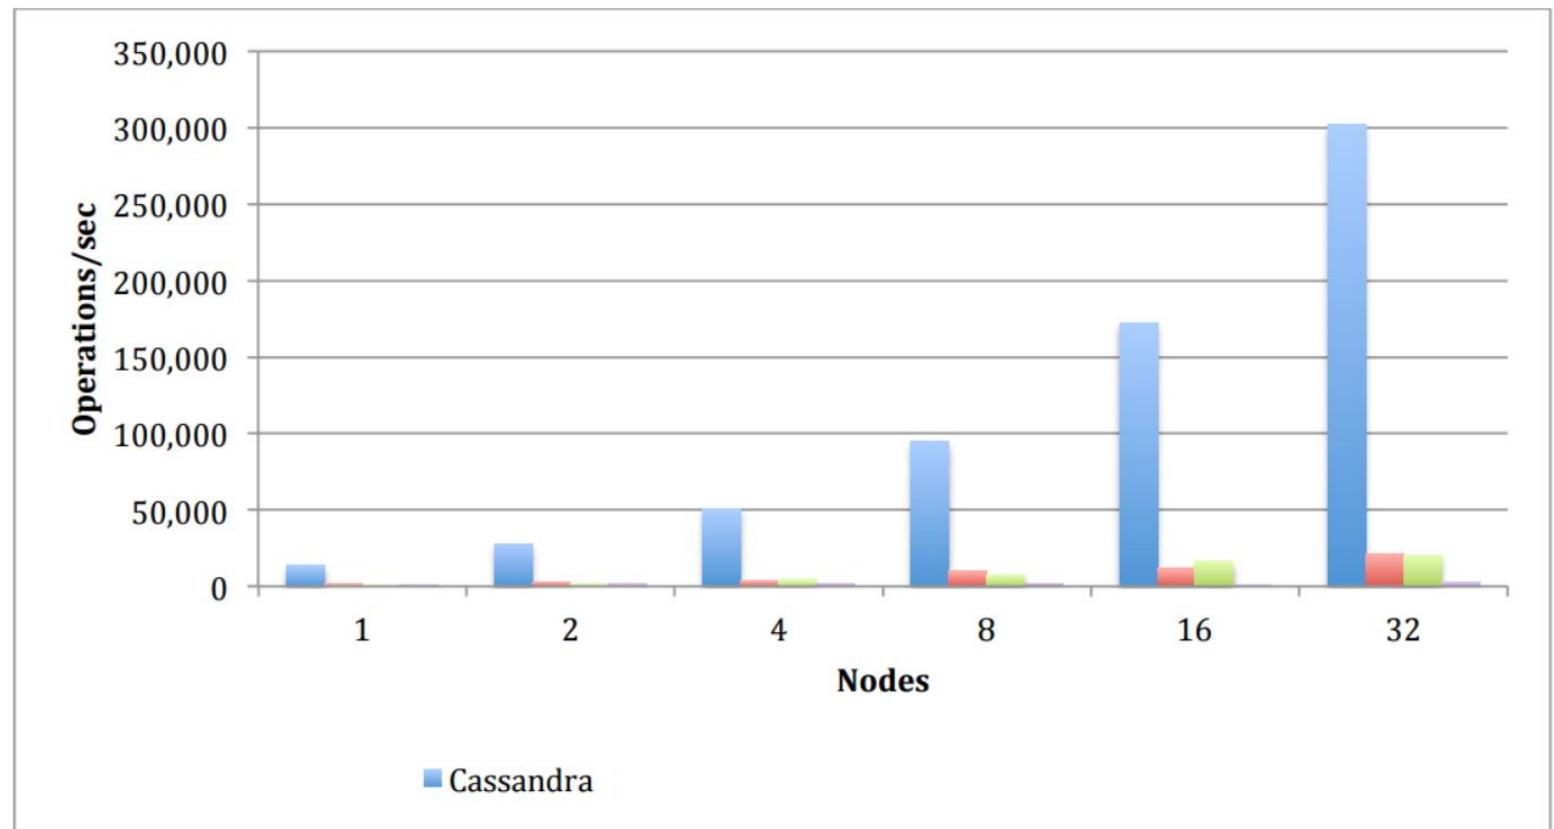
400,000 transactions/second



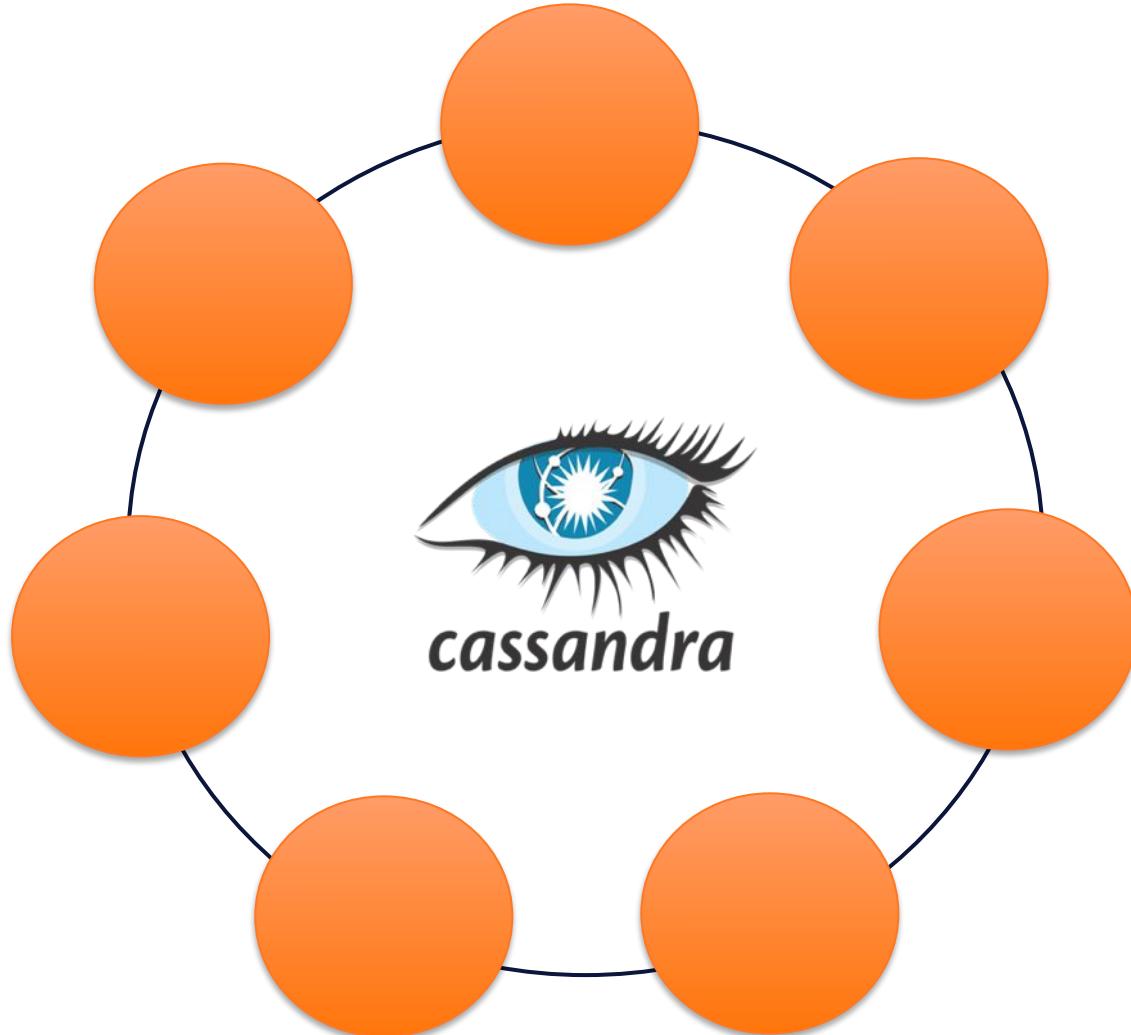
# Scales Linearly

- Need more capacity?
- Need more throughput?
- Add nodes!

Balanced Read/Write Mix



# Data is Distributed?



Город	Имя	Телефон
Moscow	Ivan	8.000.000
Moscow	Anna	4.000.000
SPb	Aleksandr	2.230.000
Berlin	Felix	3.350.000
London	John	9.200.000
Sydney	Abigail	4.900.000
Berlin	Sarah	500.000
Toronto	Andrew	6.200.000
Toronto	Peter	4.200.000
Paris	Cedrick	1.100.000
Paris	Evelyn	37.430.000
Mumbai	Prabhakaran	20.200.000

Partition Key

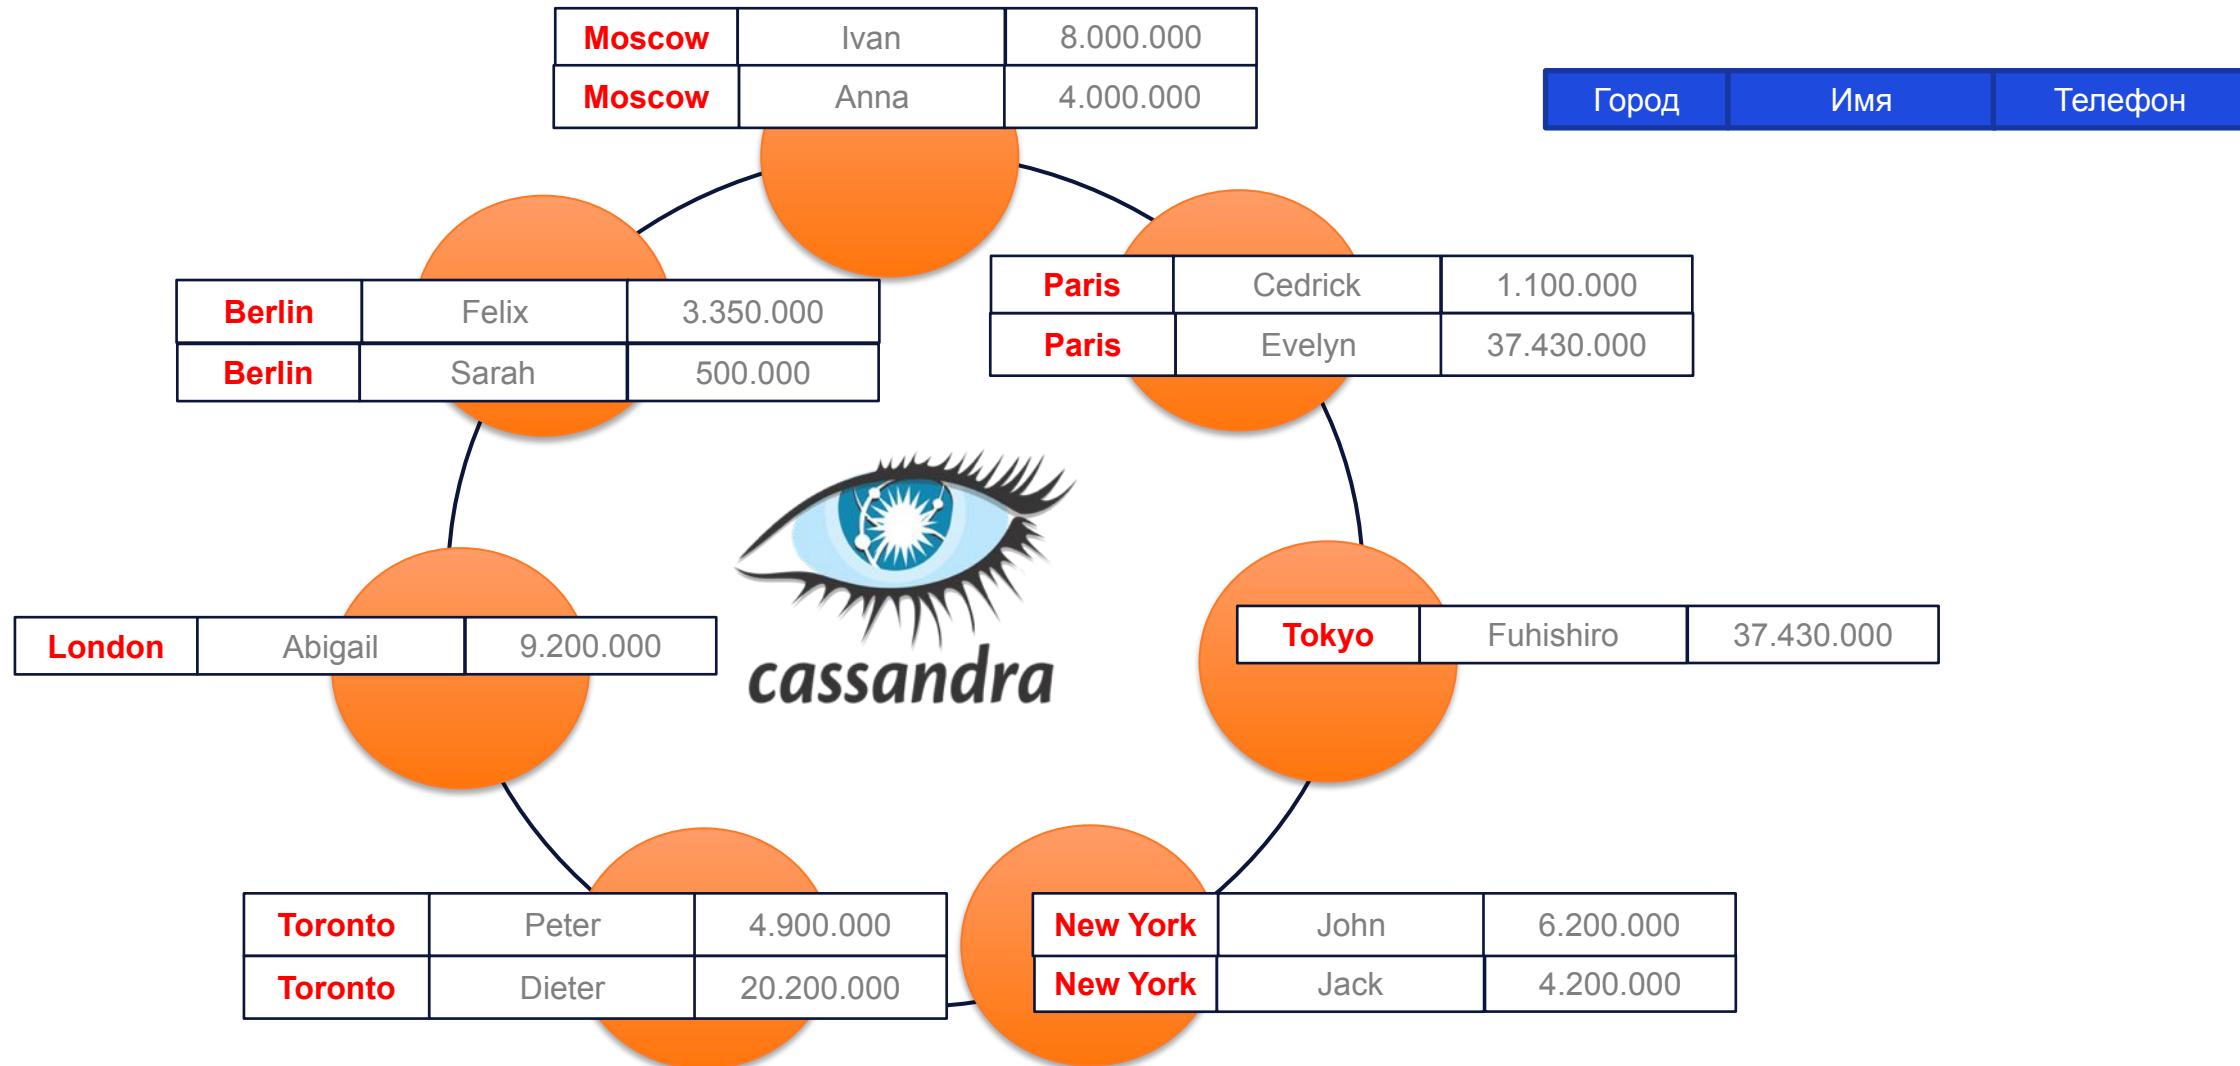


@DataStaxDevs #DataStaxDeveloperDay

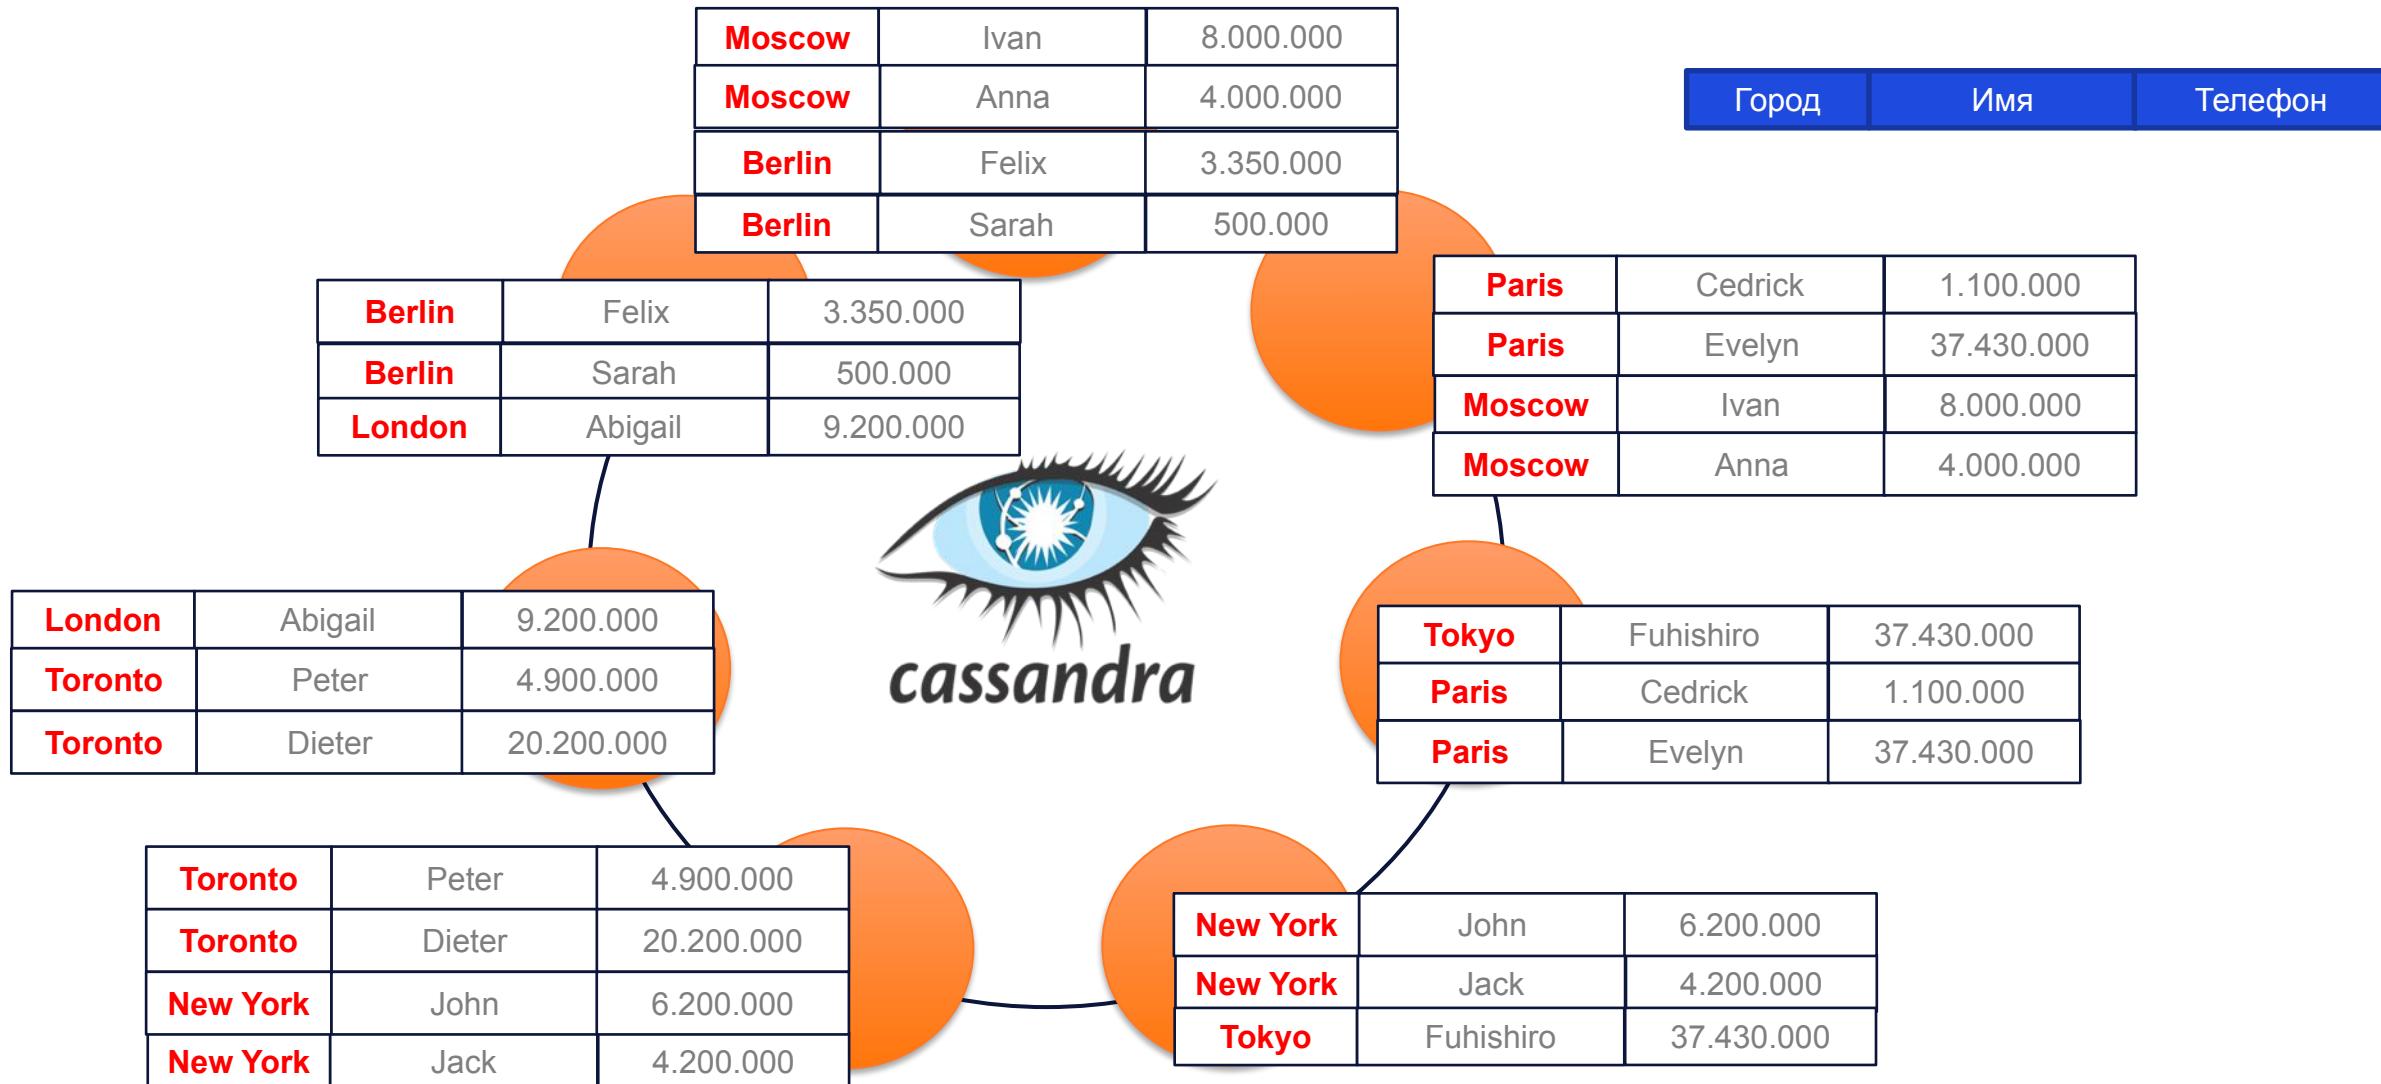
<https://community.datastax.com>



# Data is Distributed

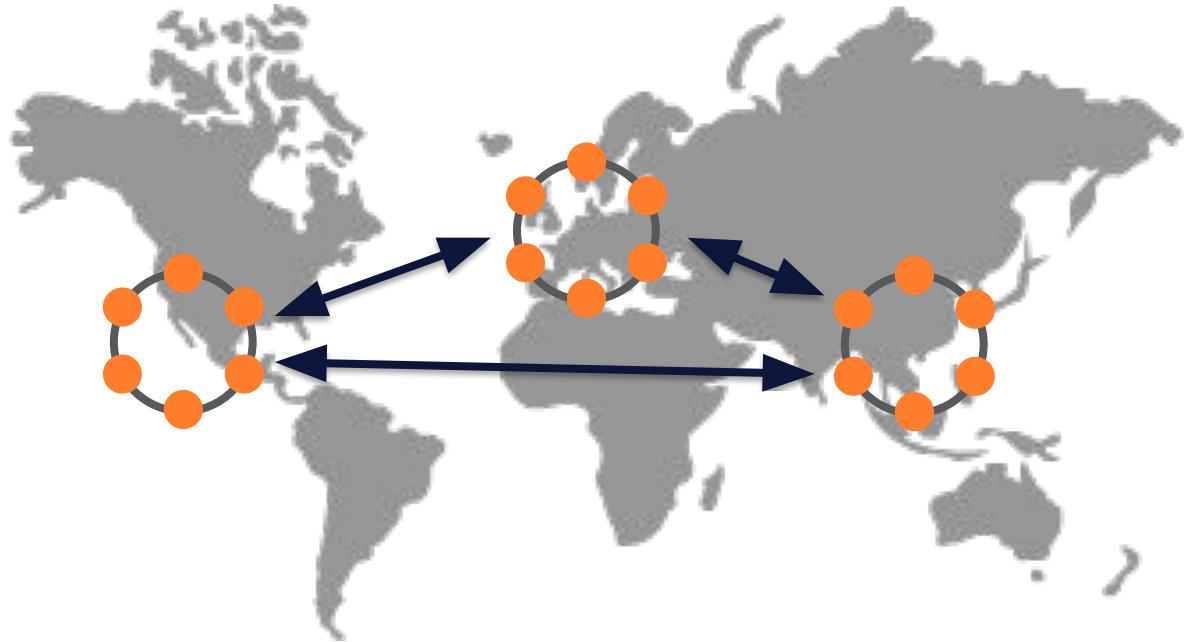


# Data is Replicated

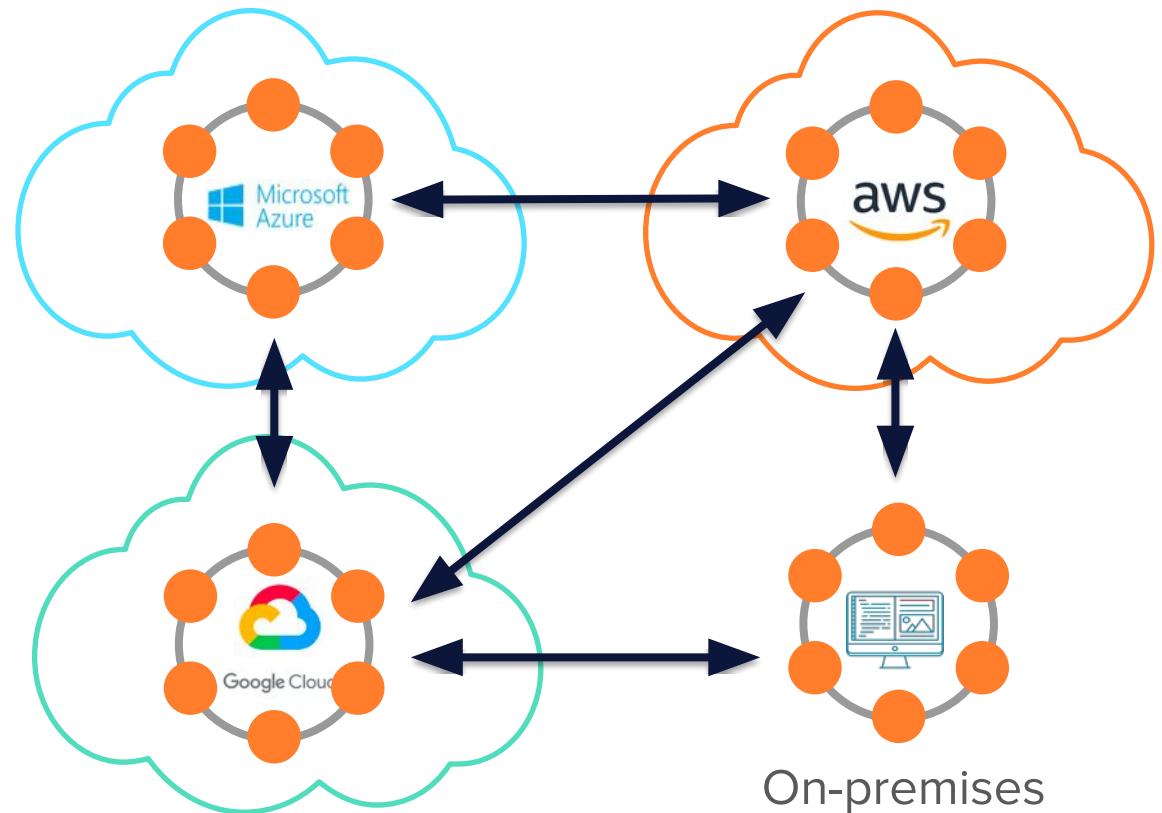


# Data Distributed Everywhere

- Geographic Distribution



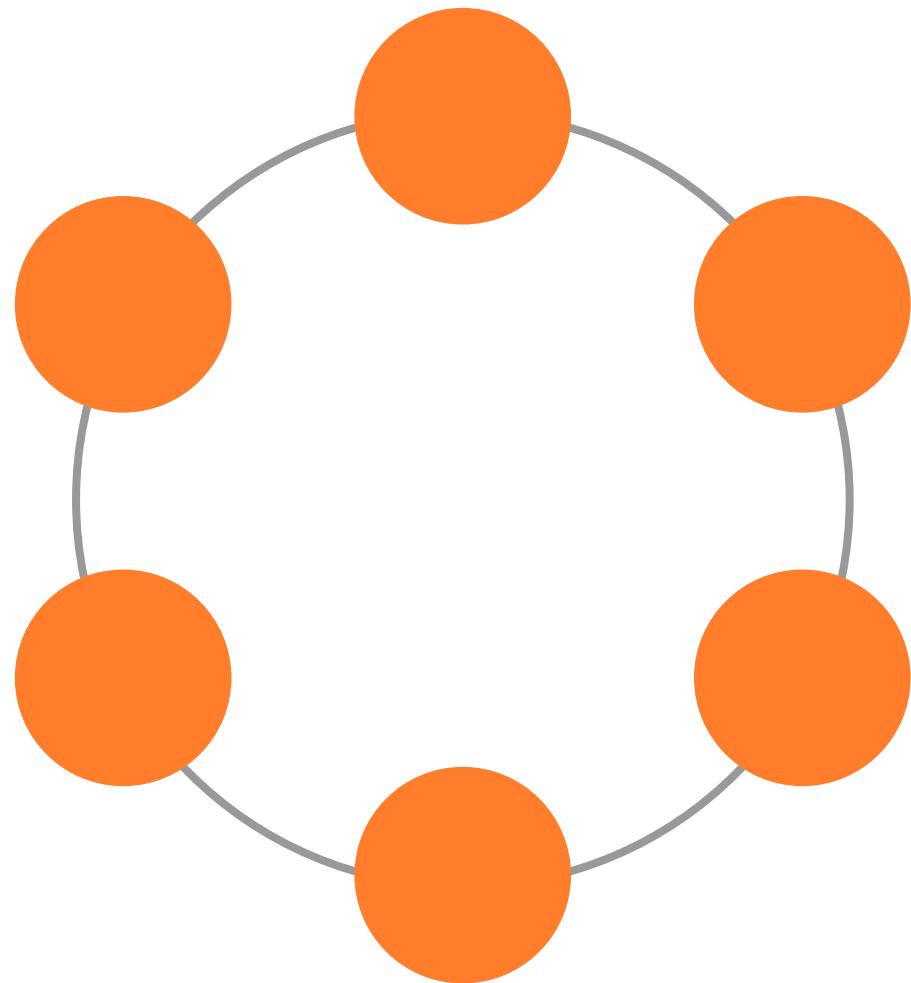
- Hybrid-Cloud and Multi-Cloud



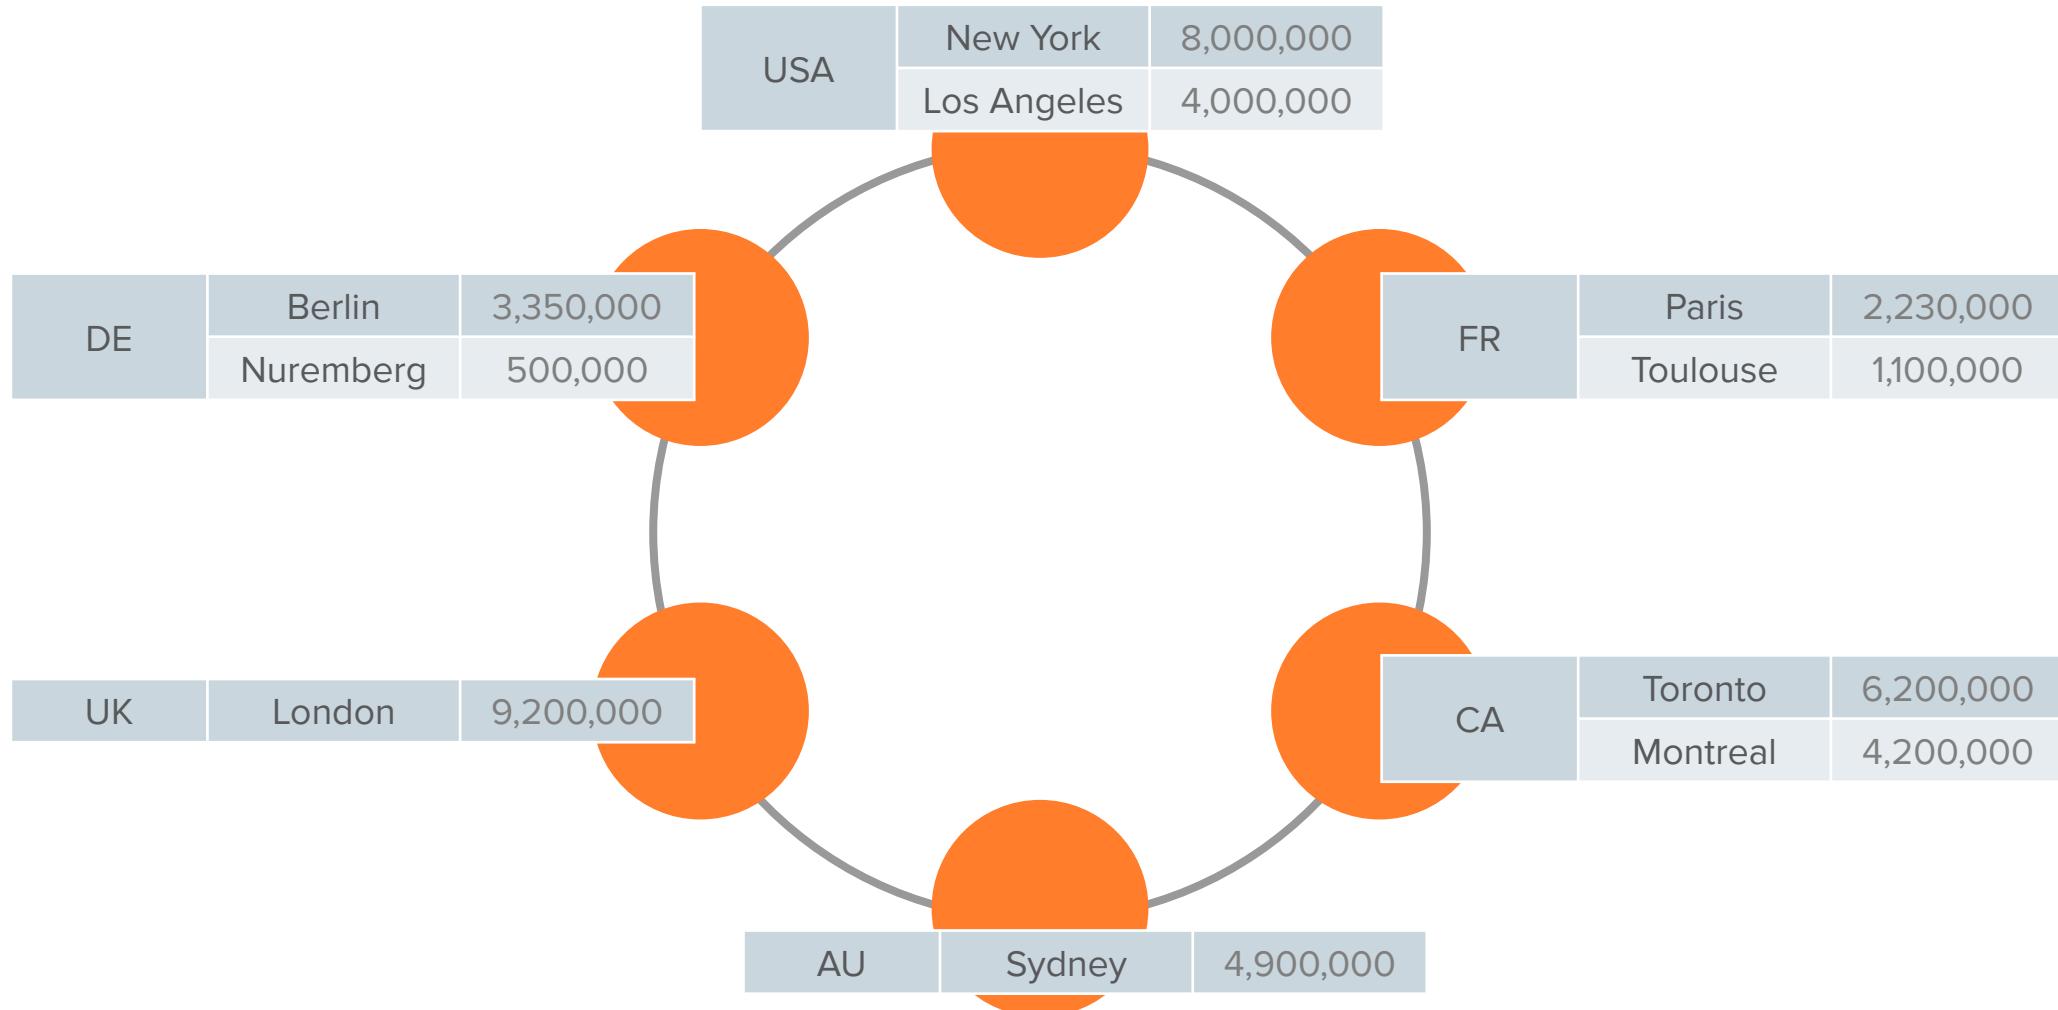
# How the Ring Works

- How is the data distributed?

Country	City	Population
USA	New York	8,000,000
	Los Angeles	4,000,000
DE	Berlin	3,350,000
	Nuremberg	500,000
FR	Paris	2,230,000
	Toulouse	1,100,000
CA	Toronto	6,200,000
	Montreal	4,200,000
UK	London	9,200,000
AU	Sydney	4,900,000



# How the Ring Works

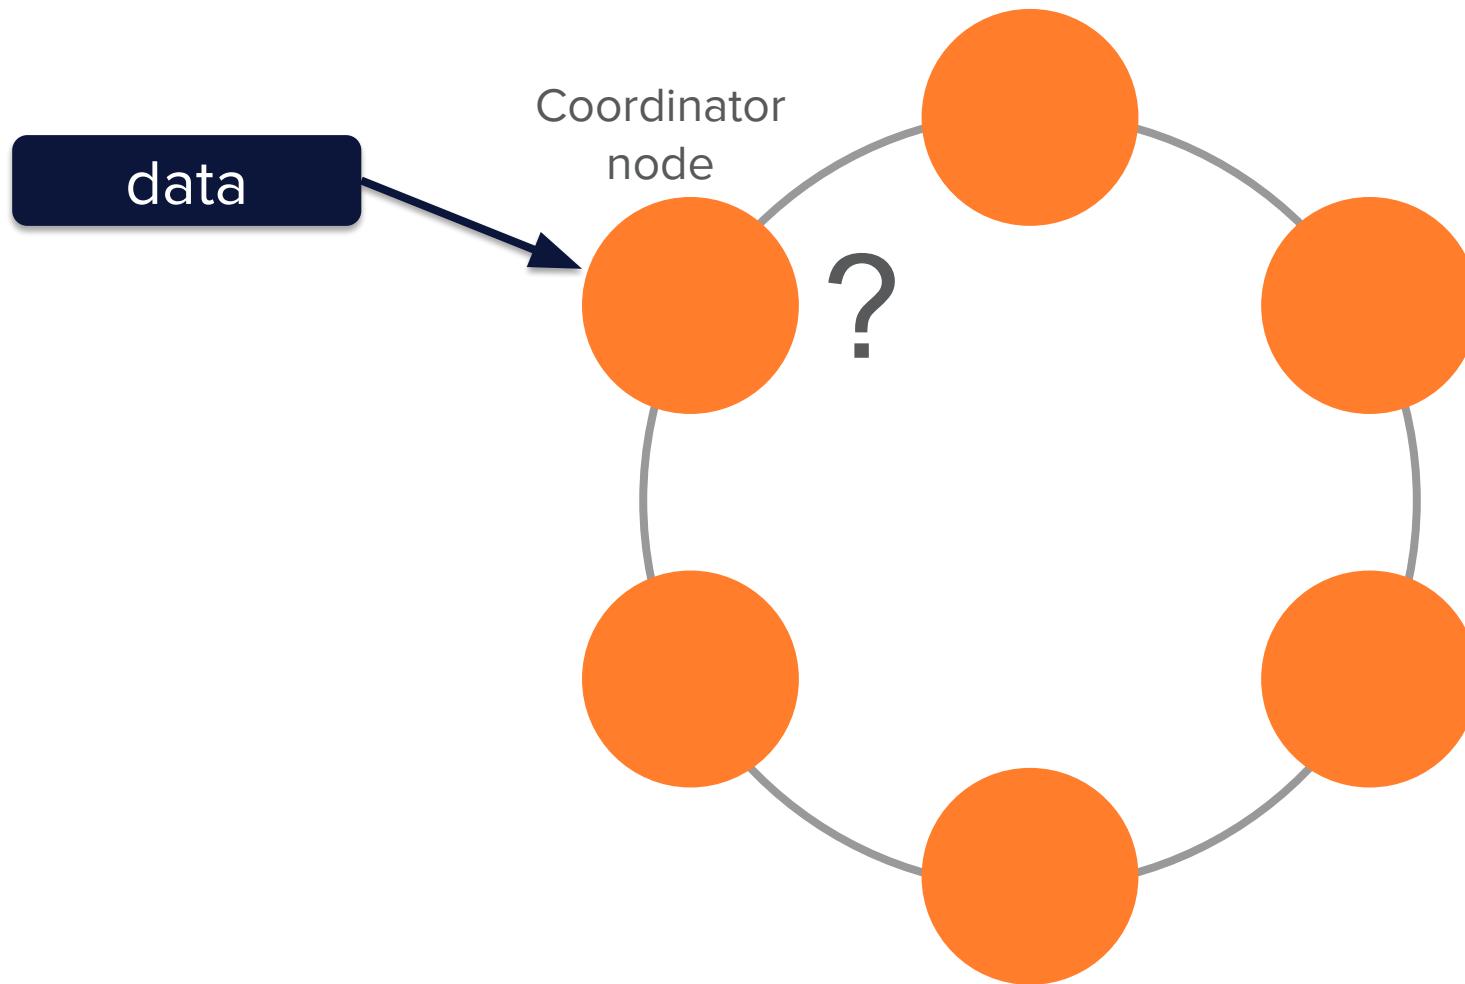


@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>



# How the Ring Works

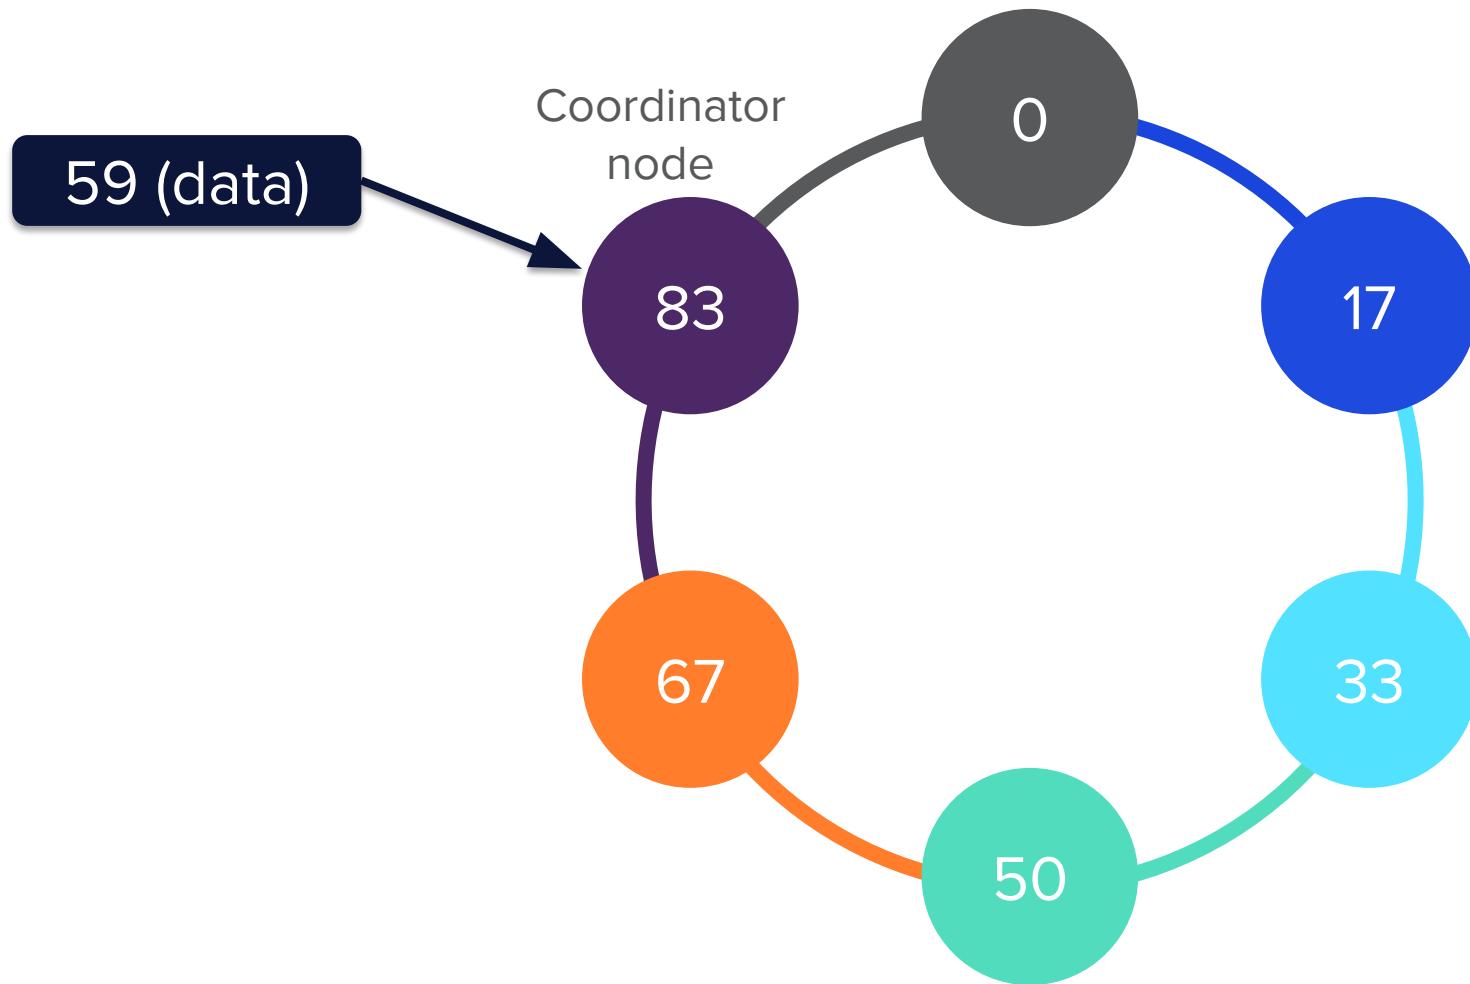


@DataStaxDevs #DataStaxDeveloperDay

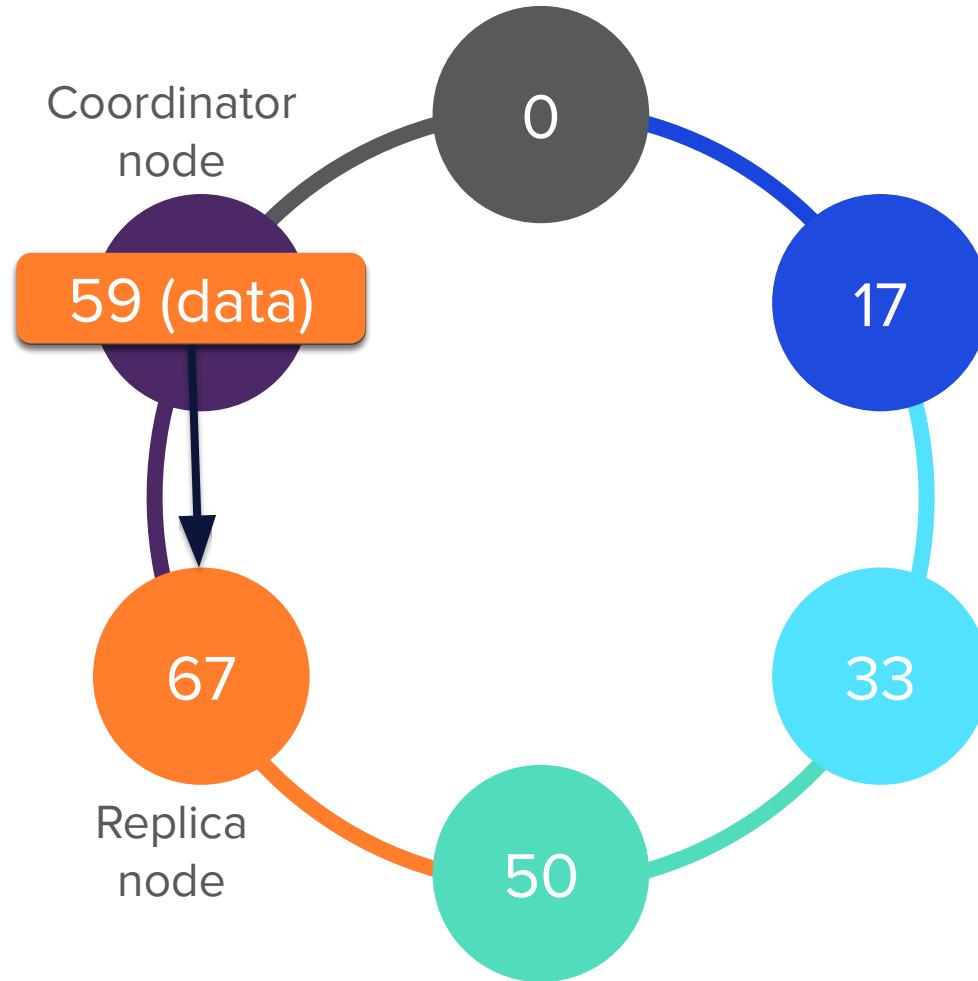
<https://community.datastax.com>



# How the Ring Works



# How the Ring Works

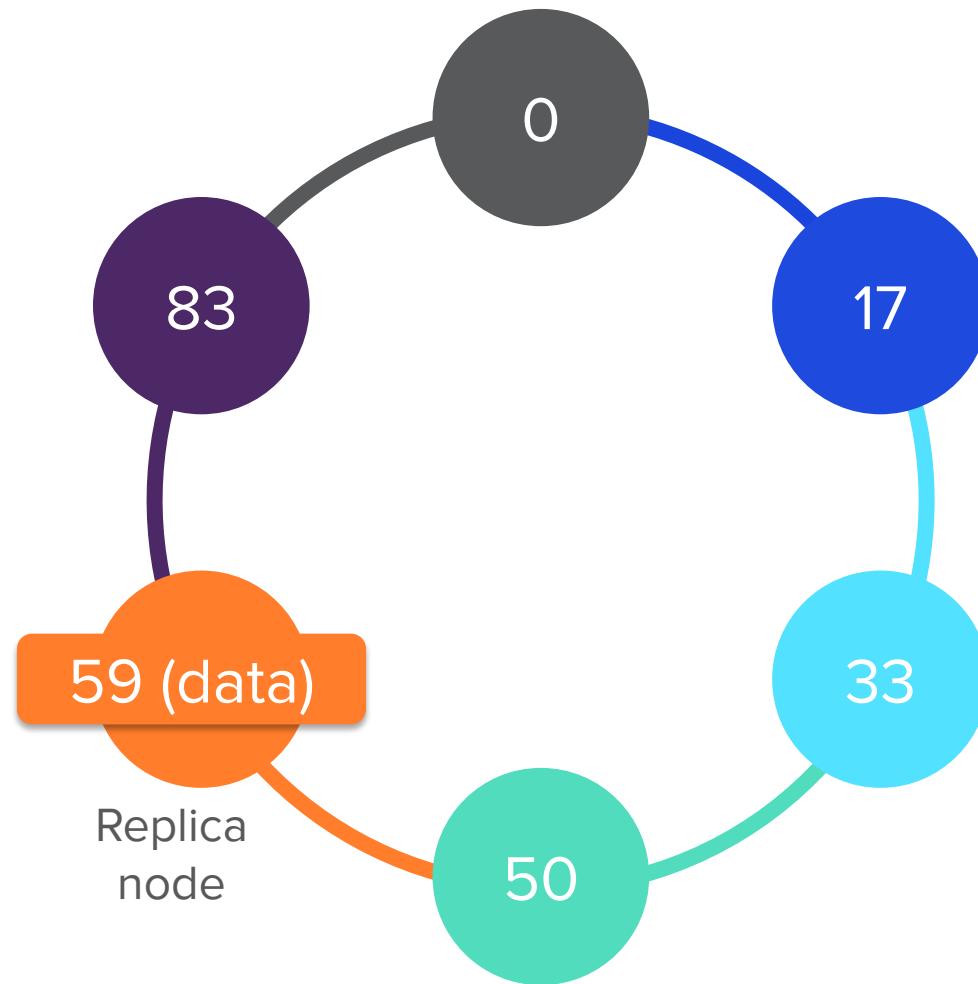


@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>

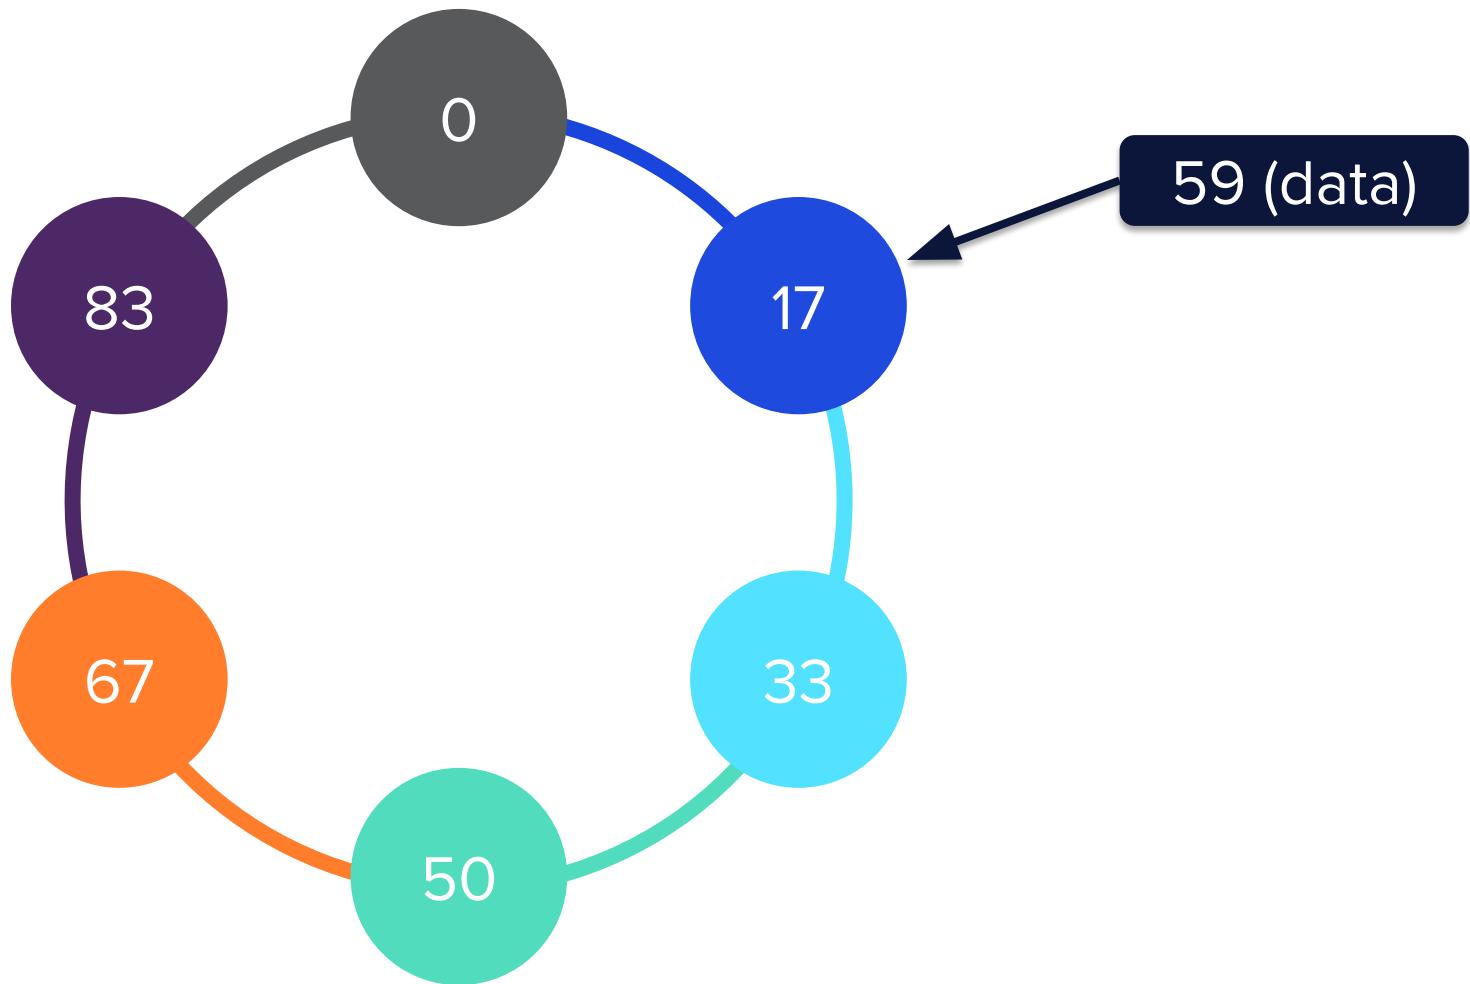


# How the Ring Works



# Replication within the Ring

RF = 1



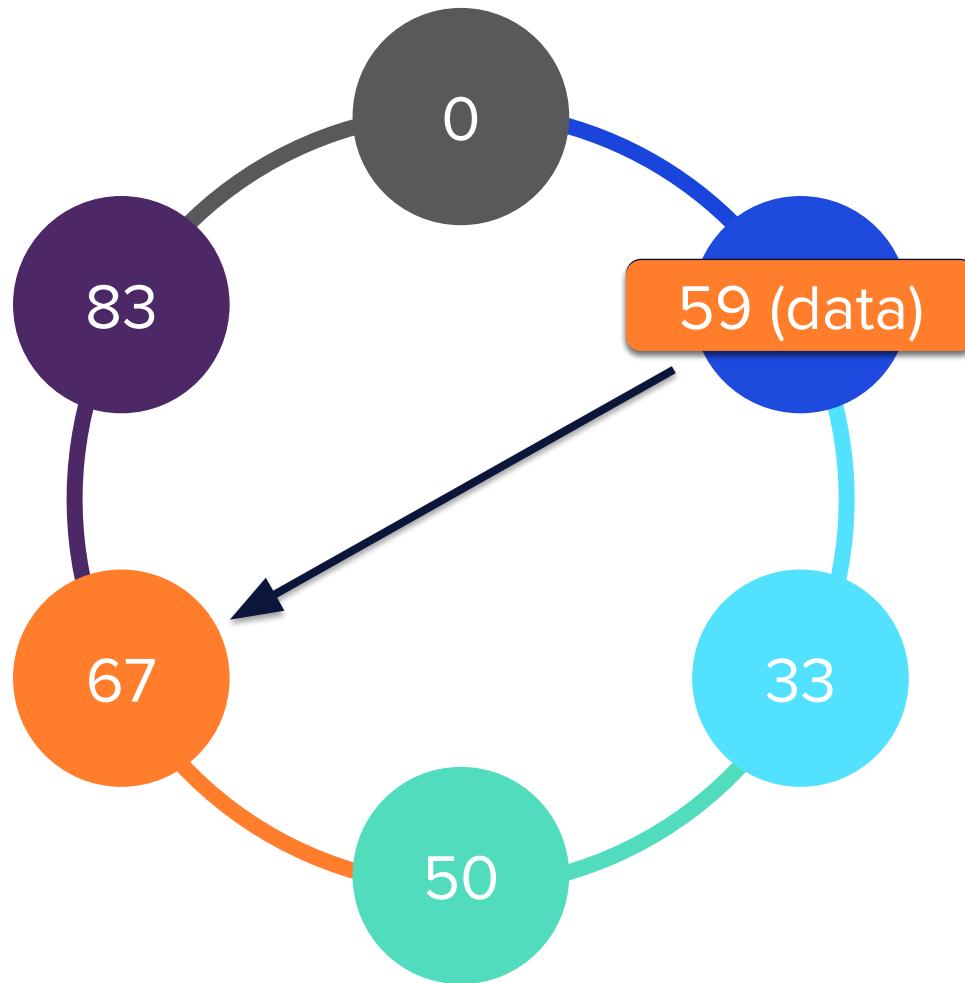
@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>



# Replication within the Ring

RF = 1



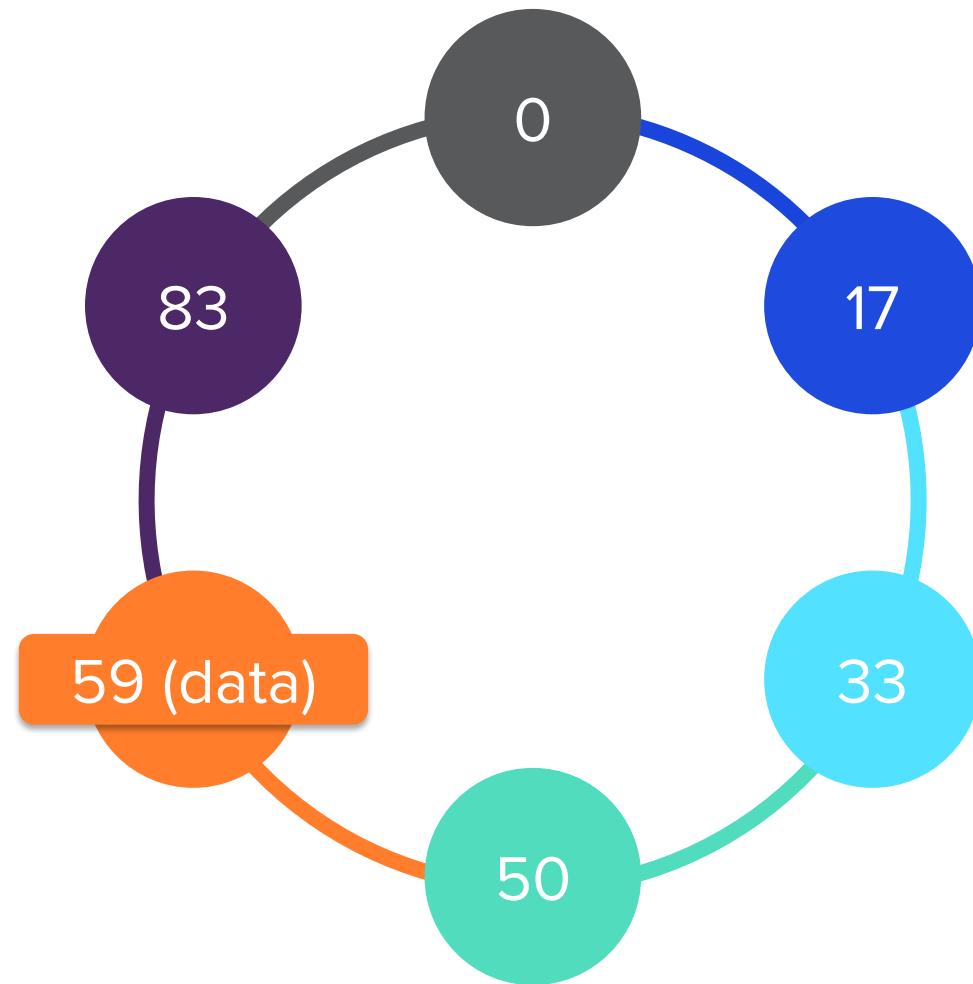
@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>



# Replication within the Ring

RF = 1



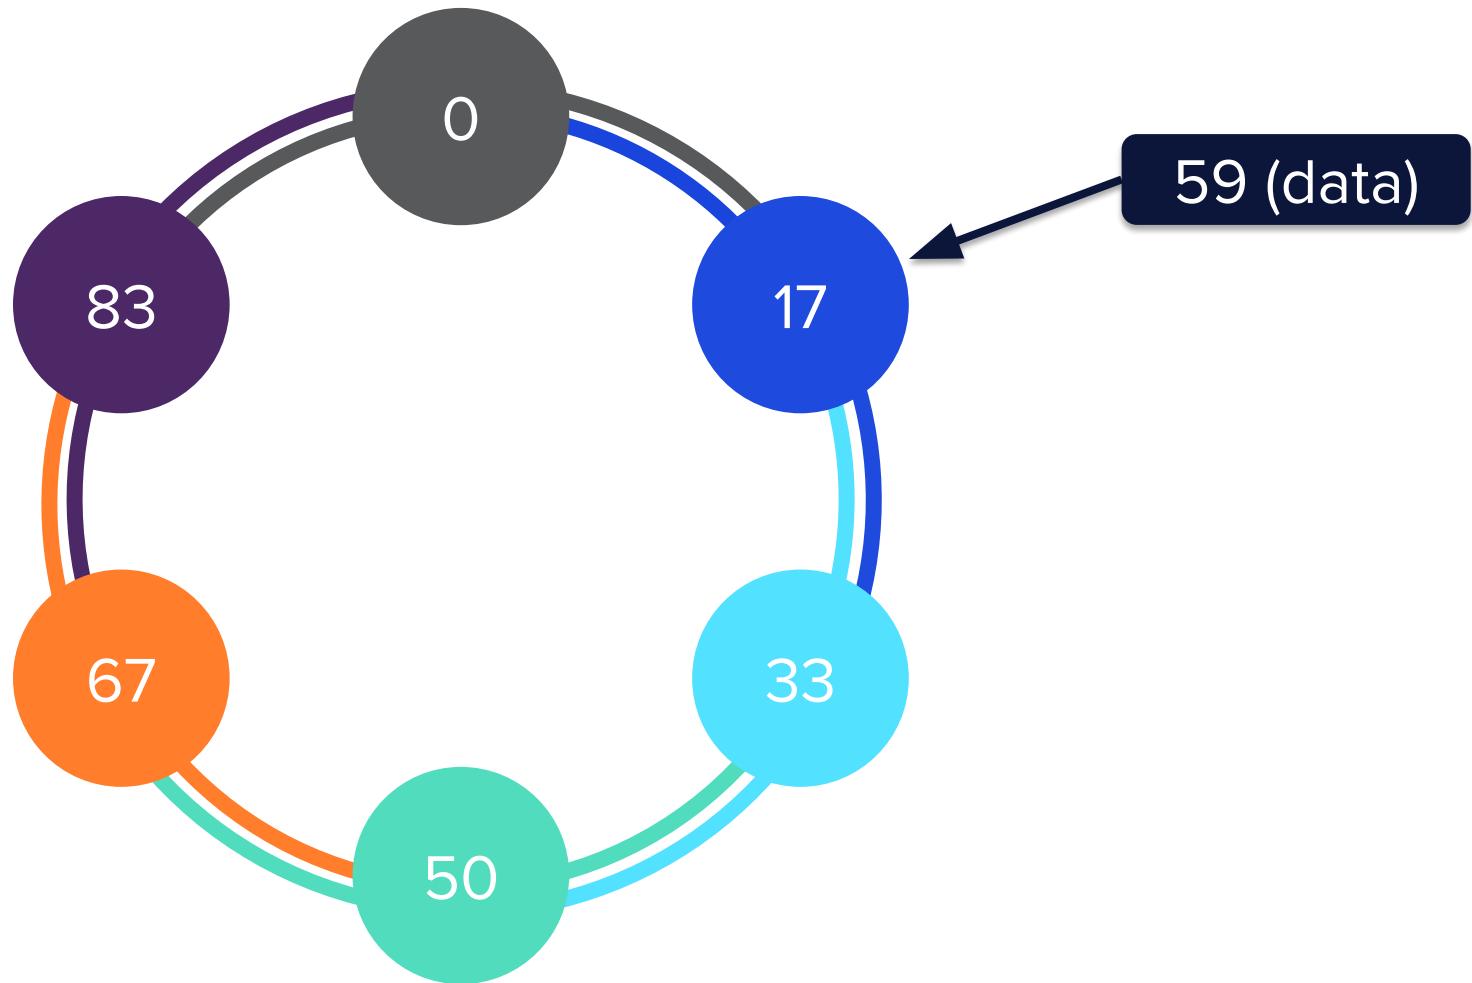
@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>



# Replication within the Ring

RF = 2



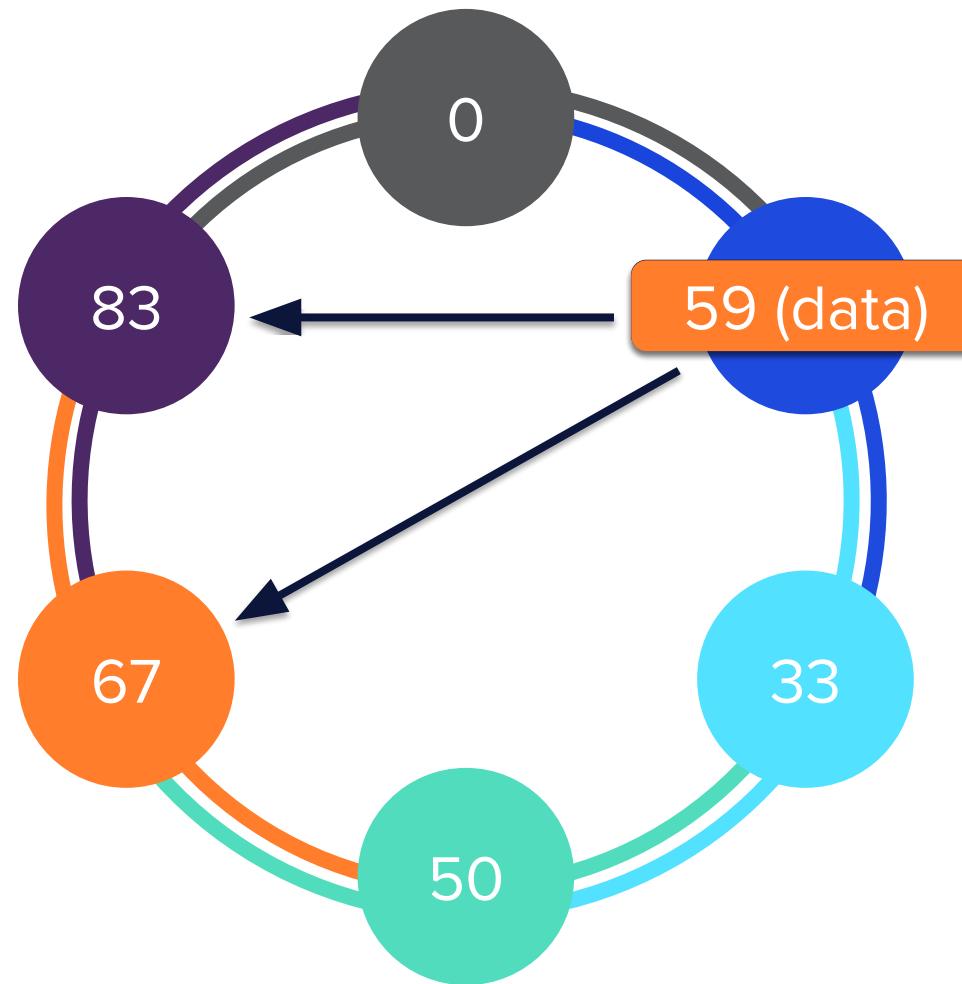
@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>



# Replication within the Ring

RF = 2



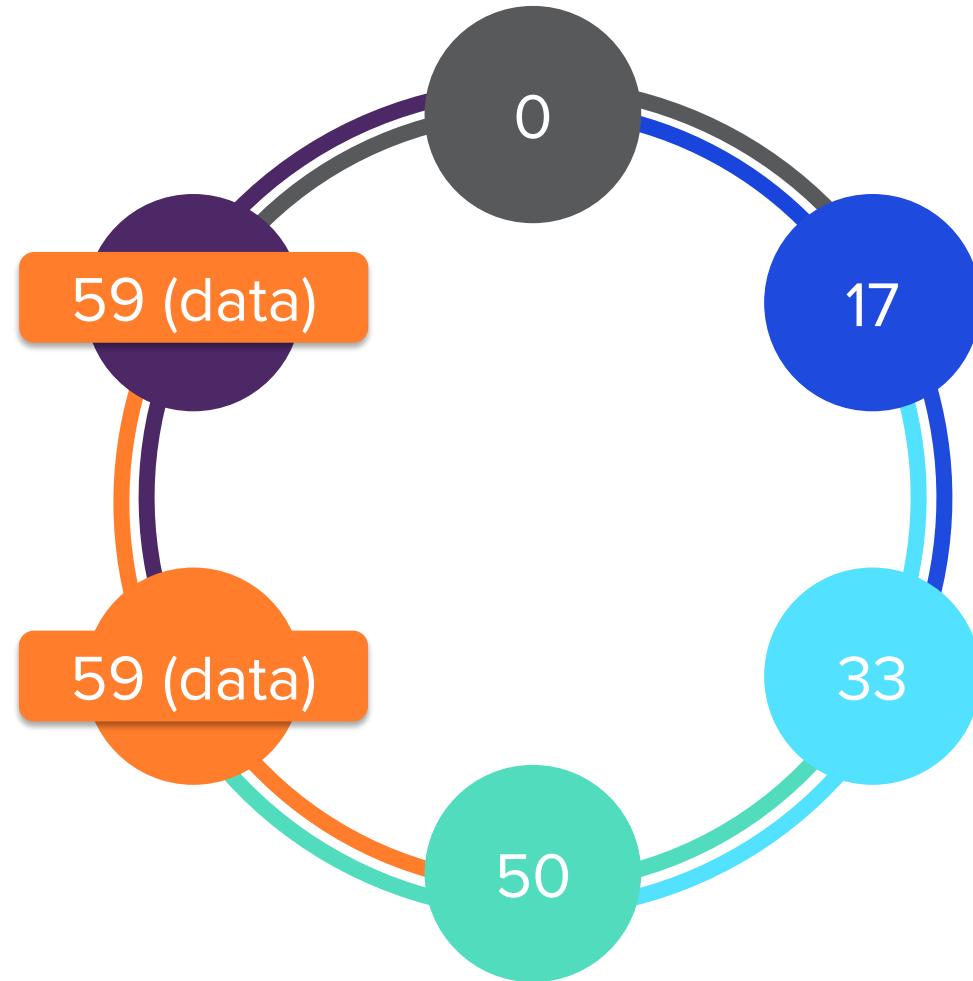
@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>



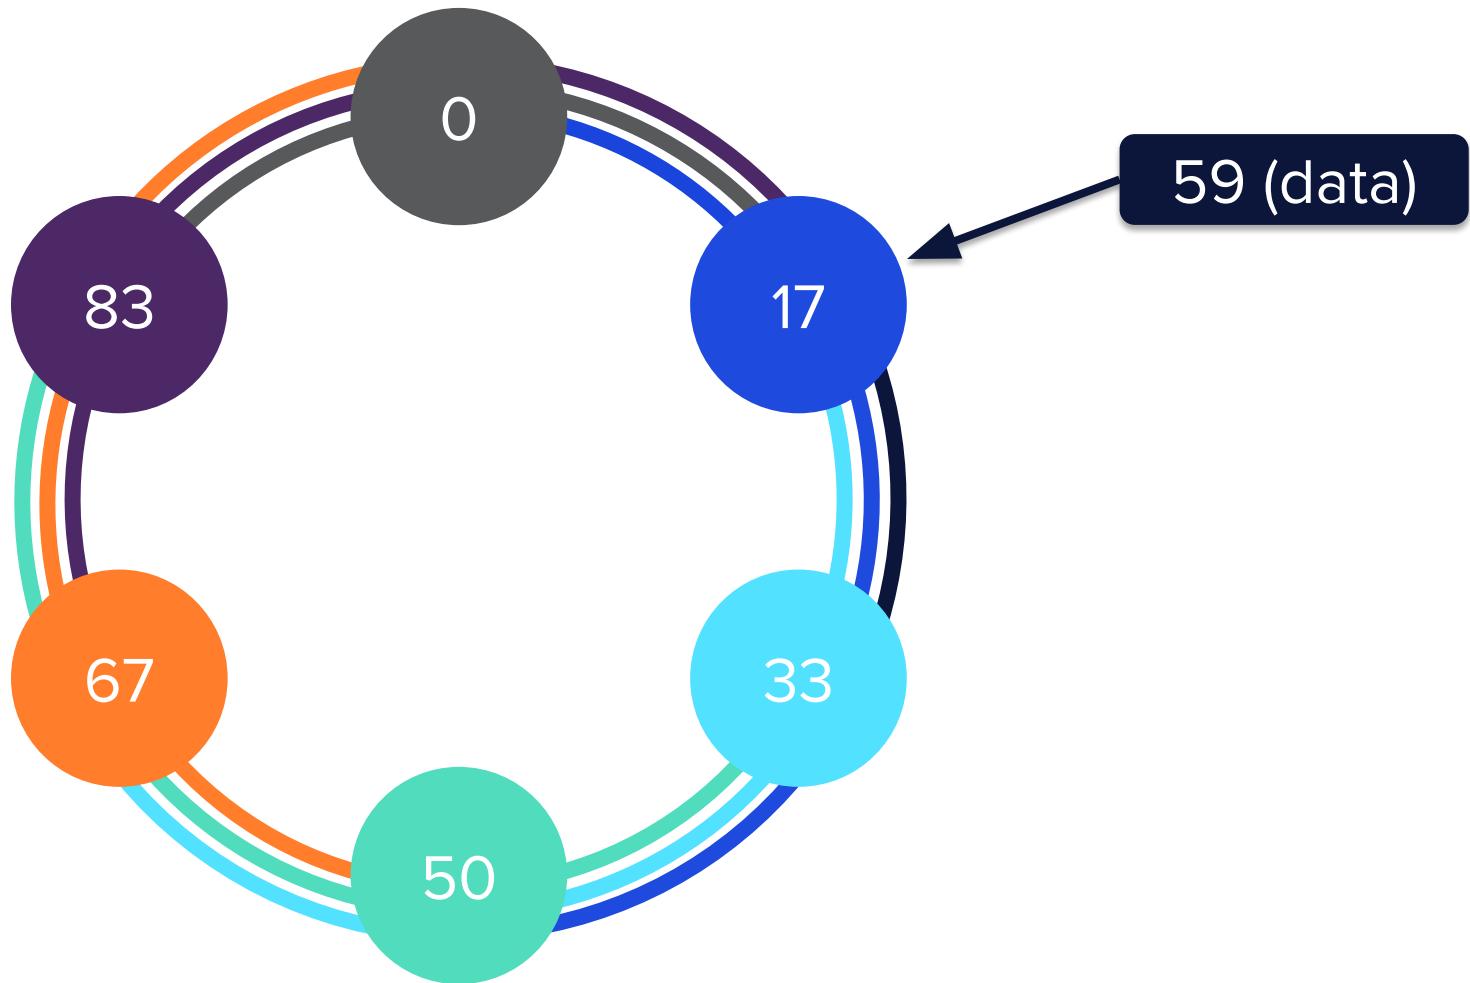
# Replication within the Ring

RF = 2



# Replication within the Ring

RF = 3



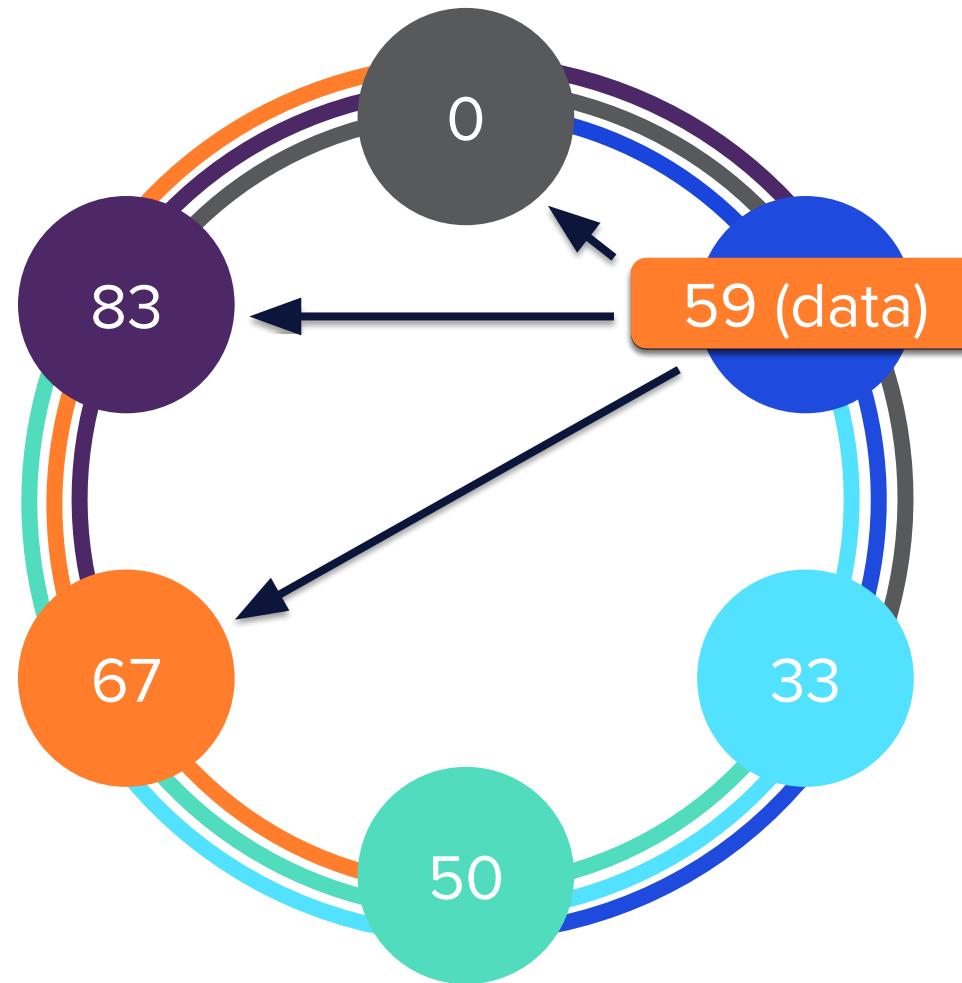
@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>



# Replication within the Ring

RF = 3



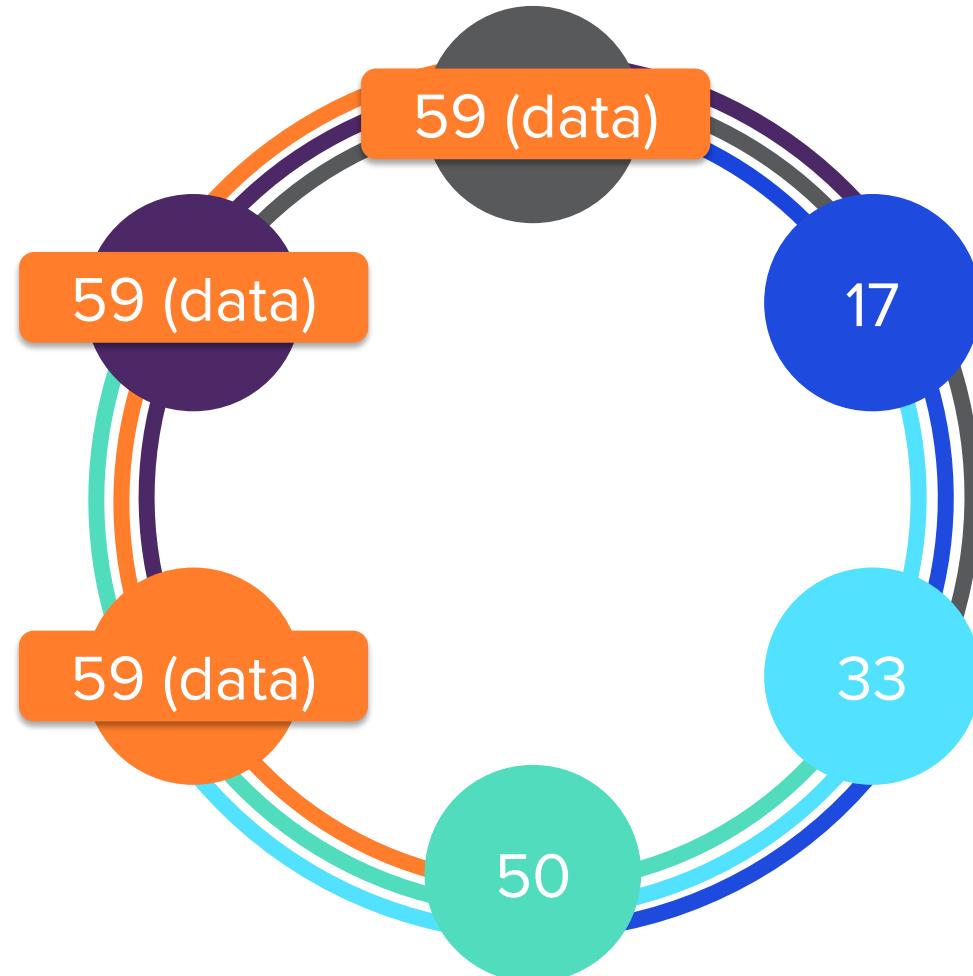
@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>



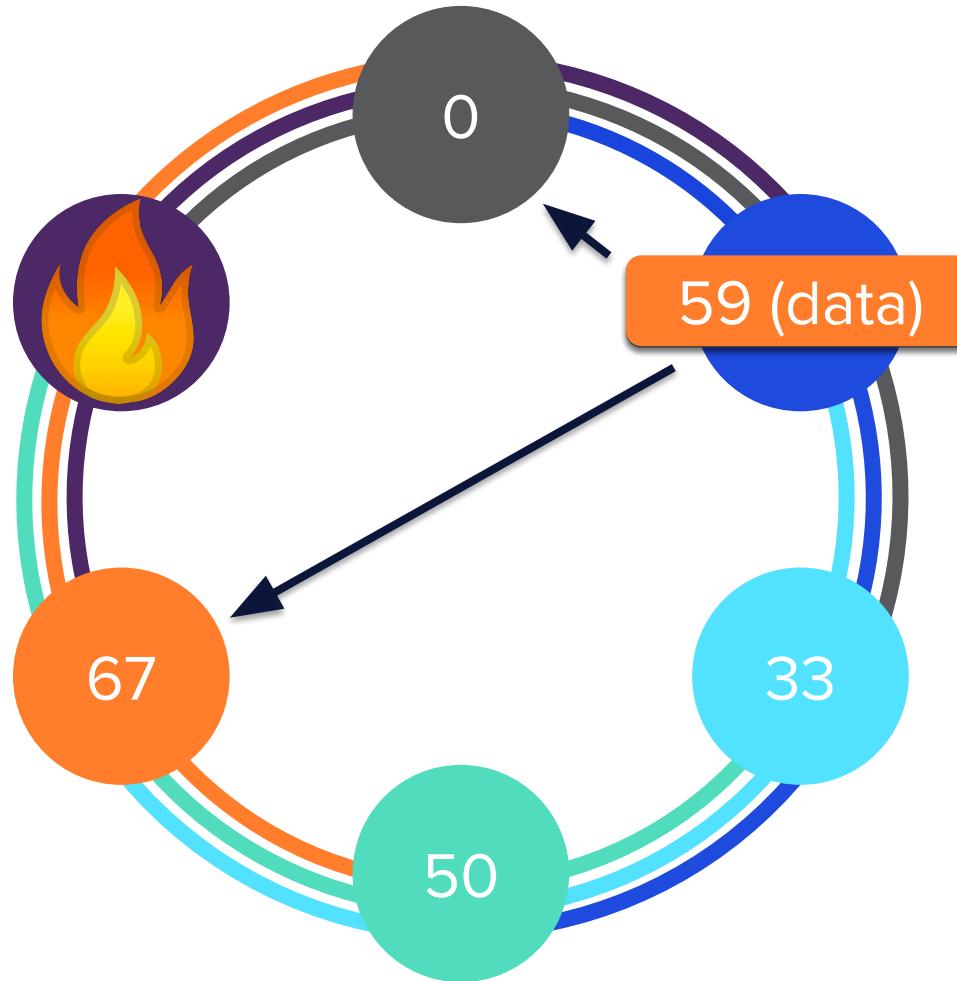
# Replication within the Ring

RF = 3



# Node Failure

RF = 3



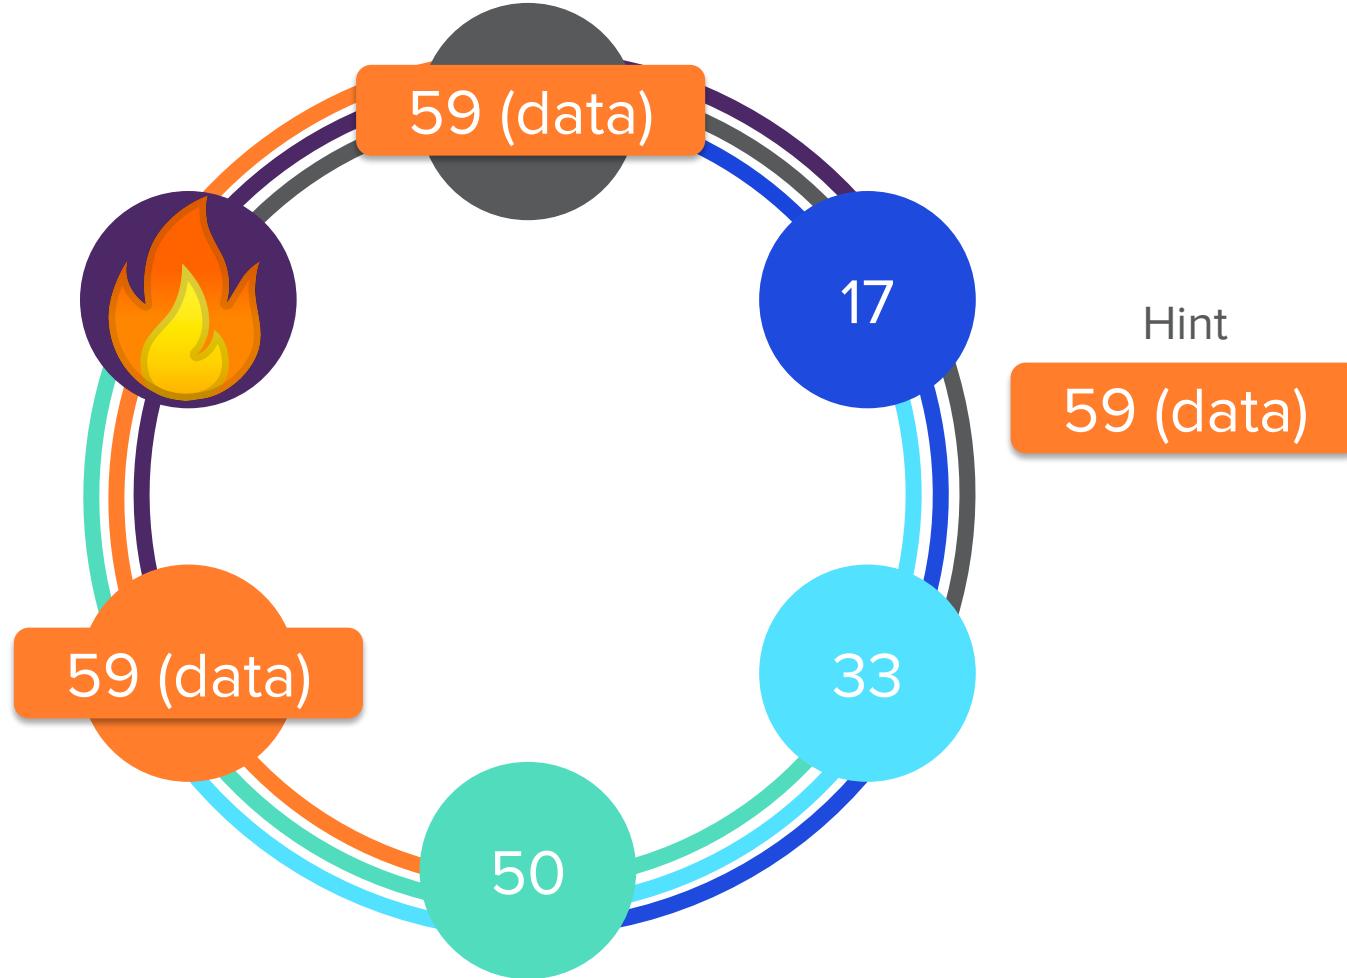
@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>



# Node Failure

RF = 3



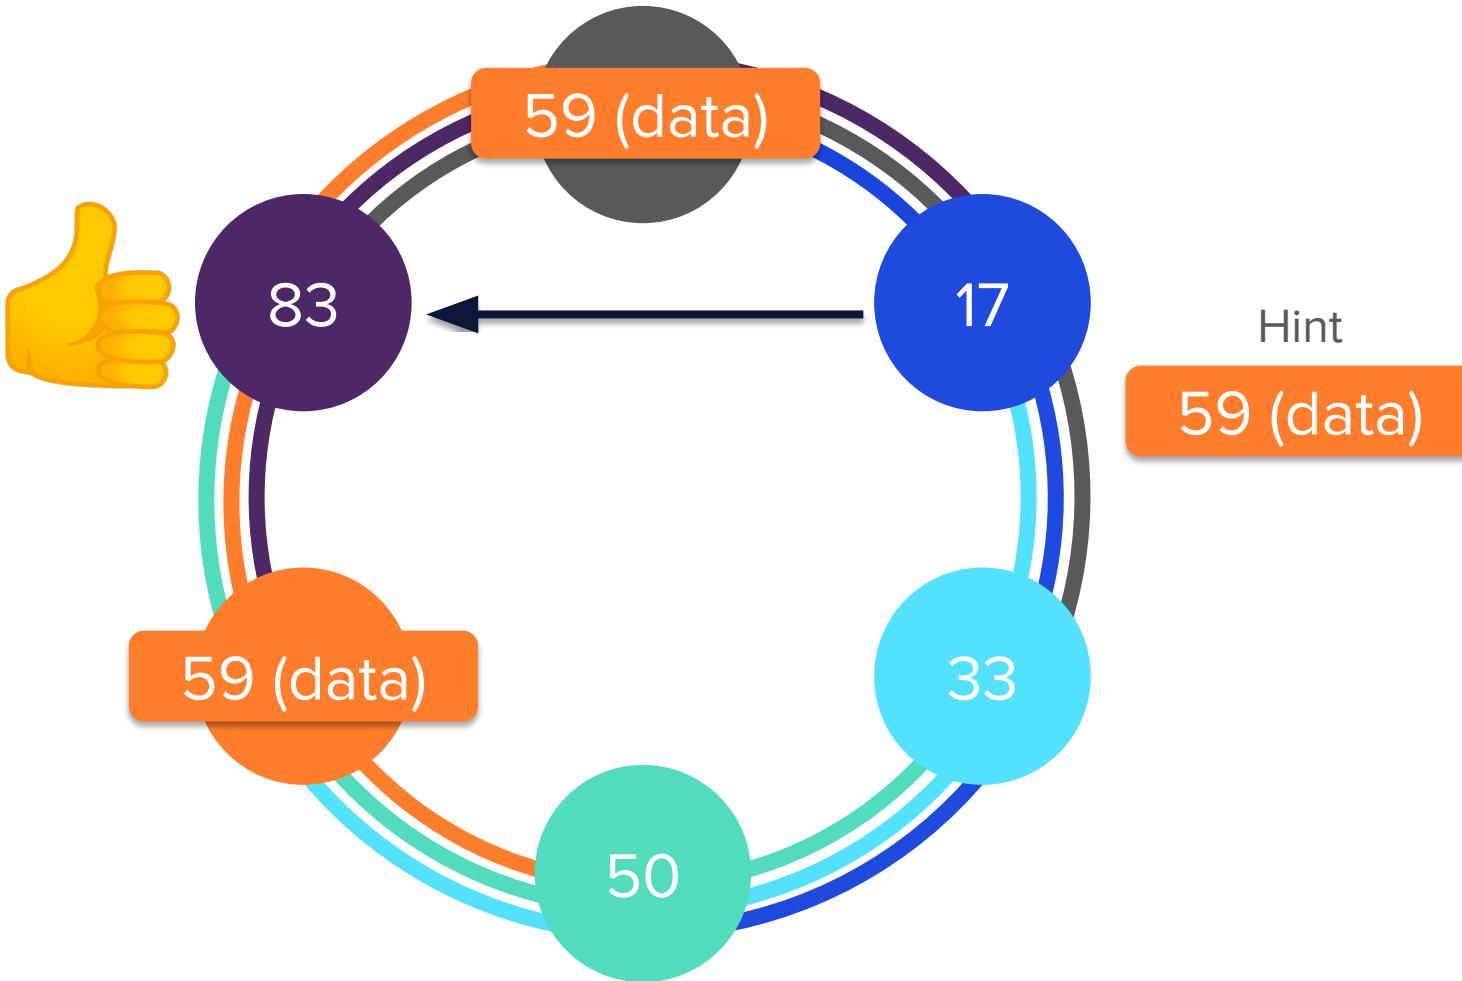
@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>



# Node Failure

RF = 3



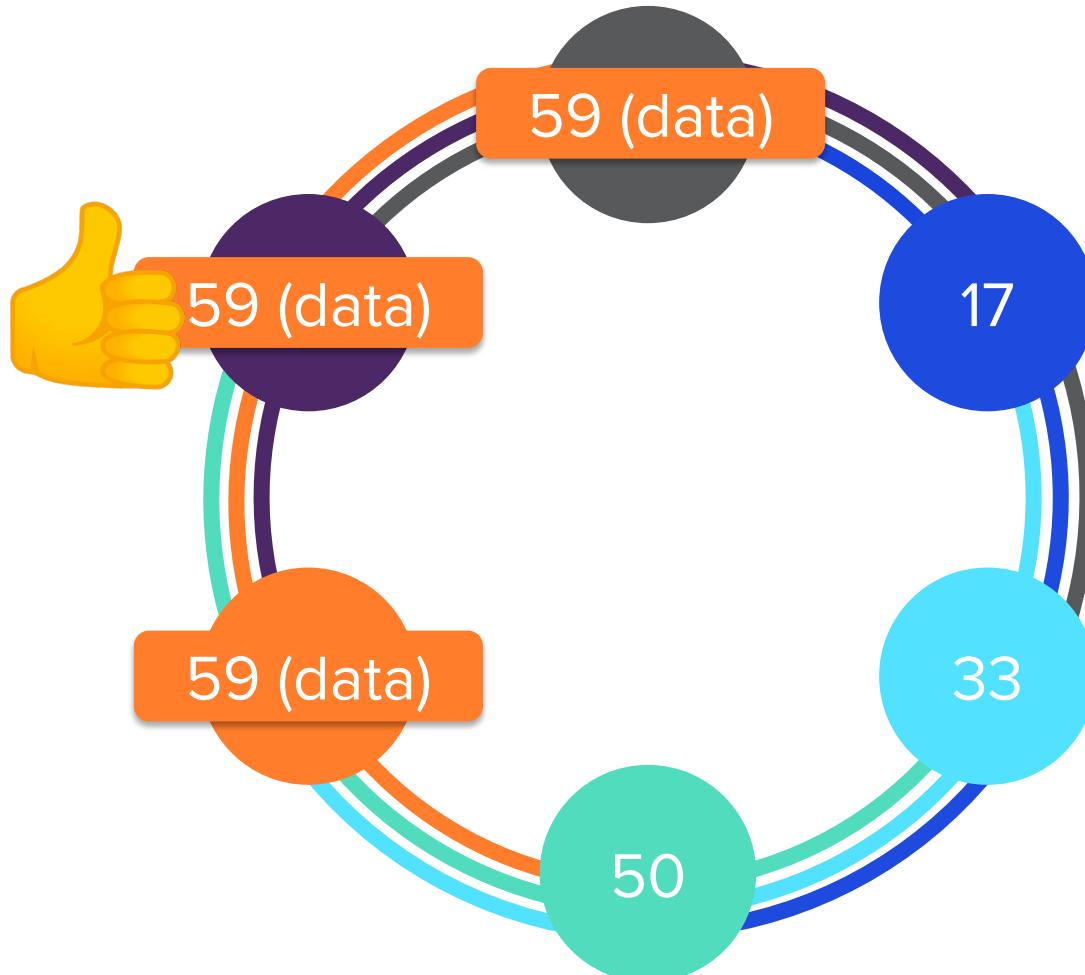
@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>



# Node Failure – Recovered!

RF = 3

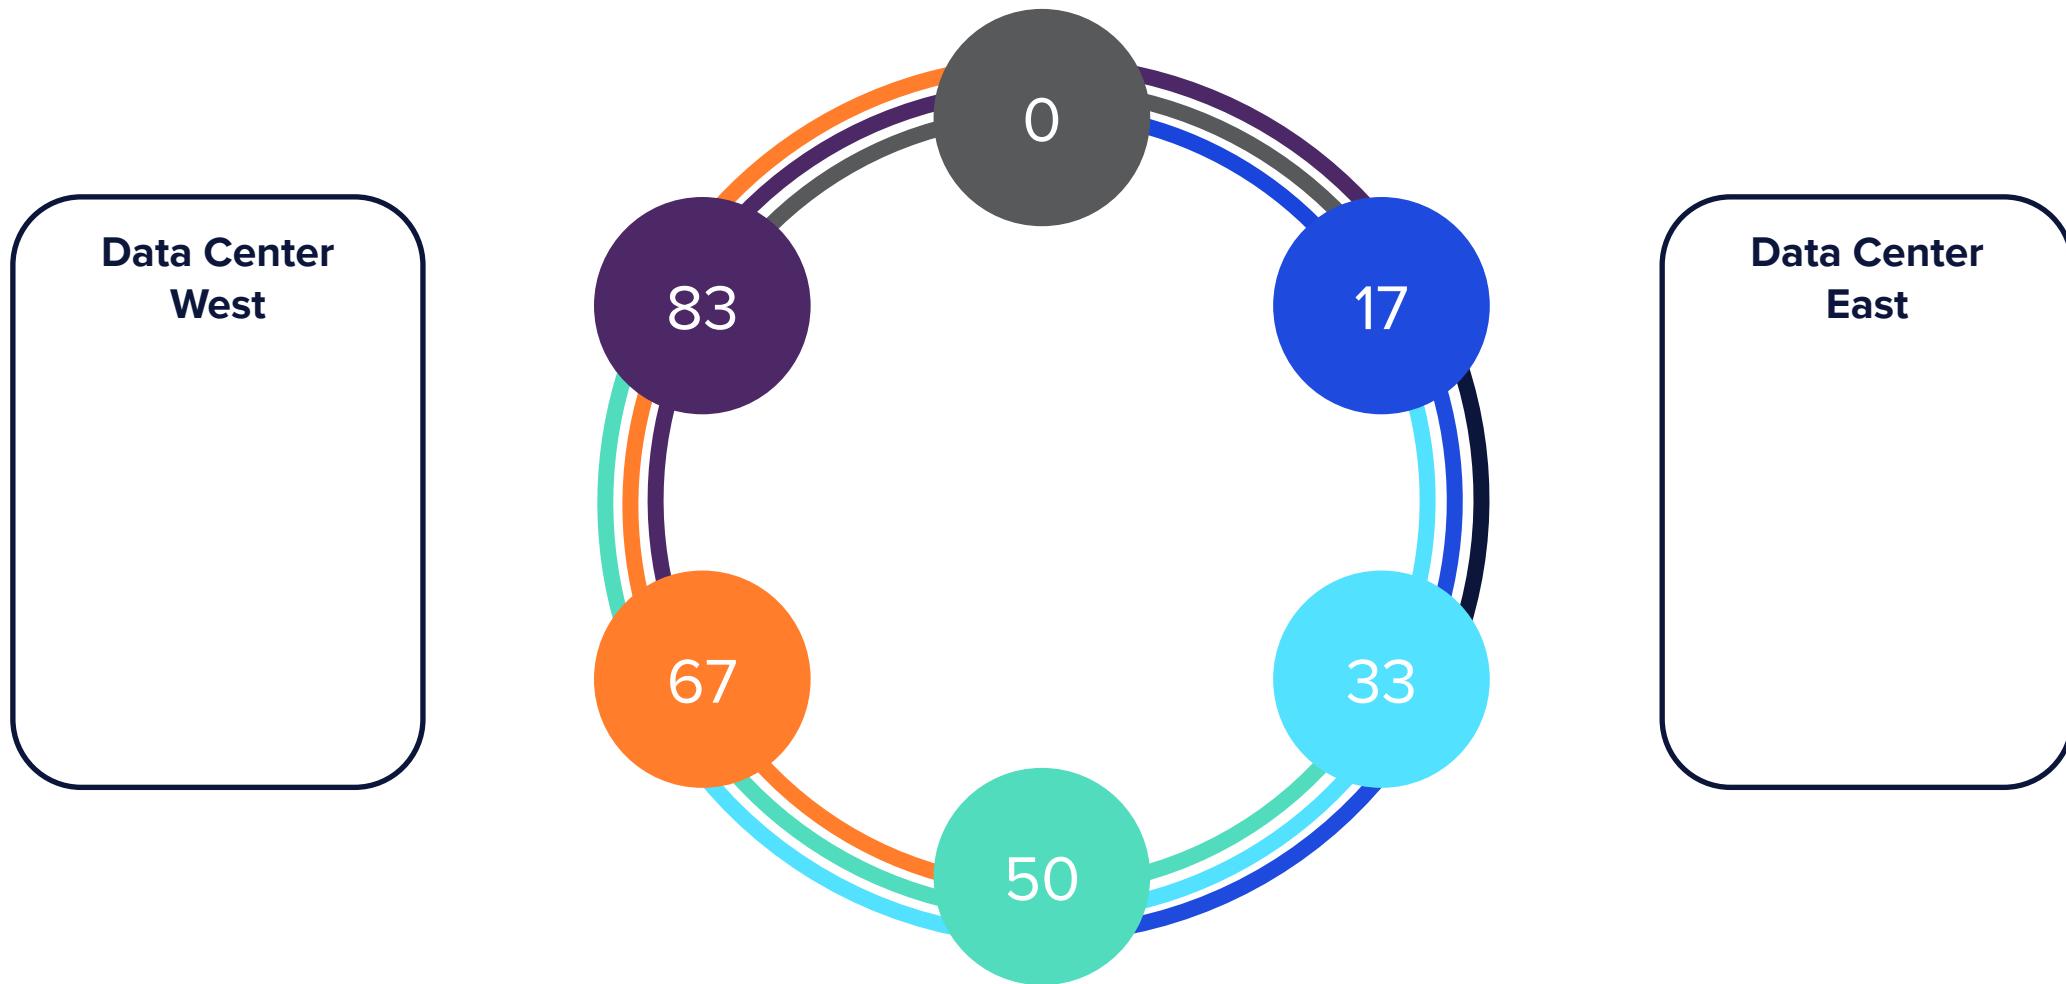


@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>



# Multi-Data Center Replication

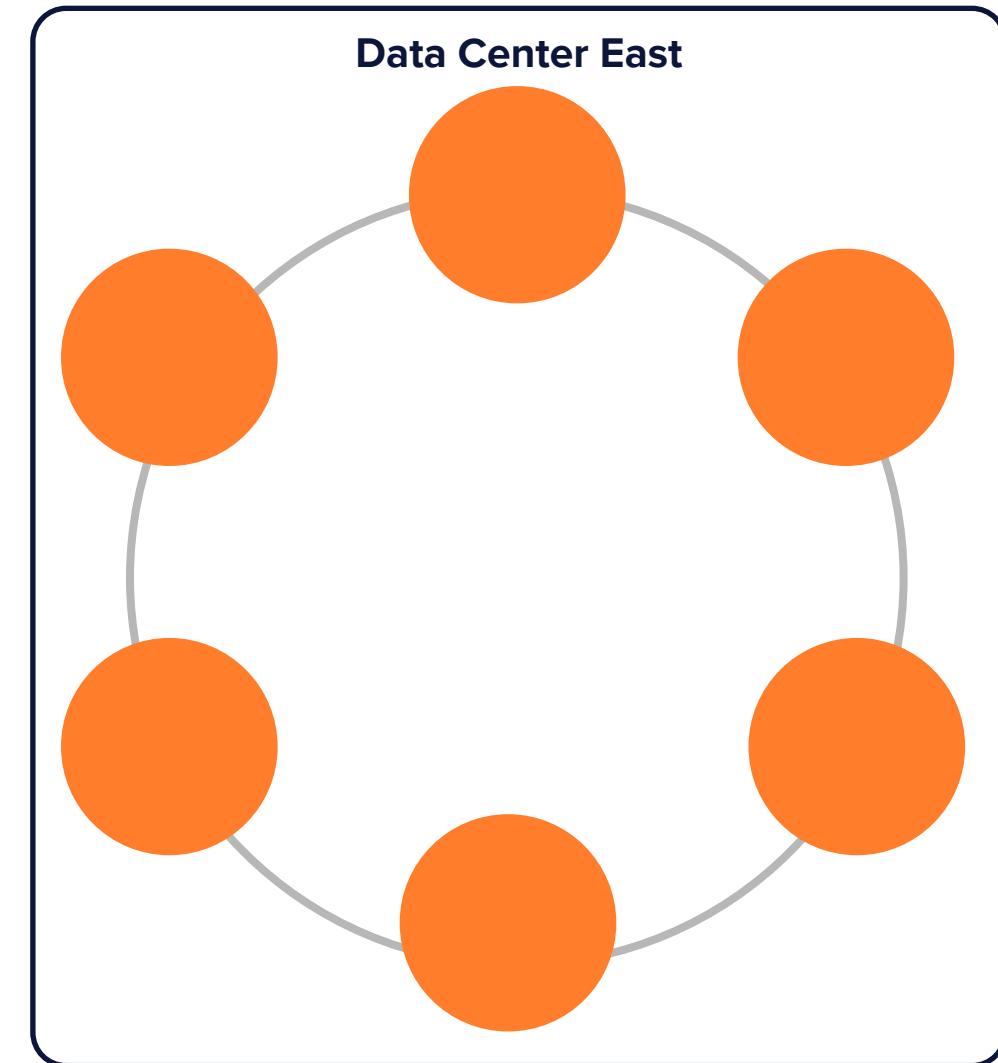
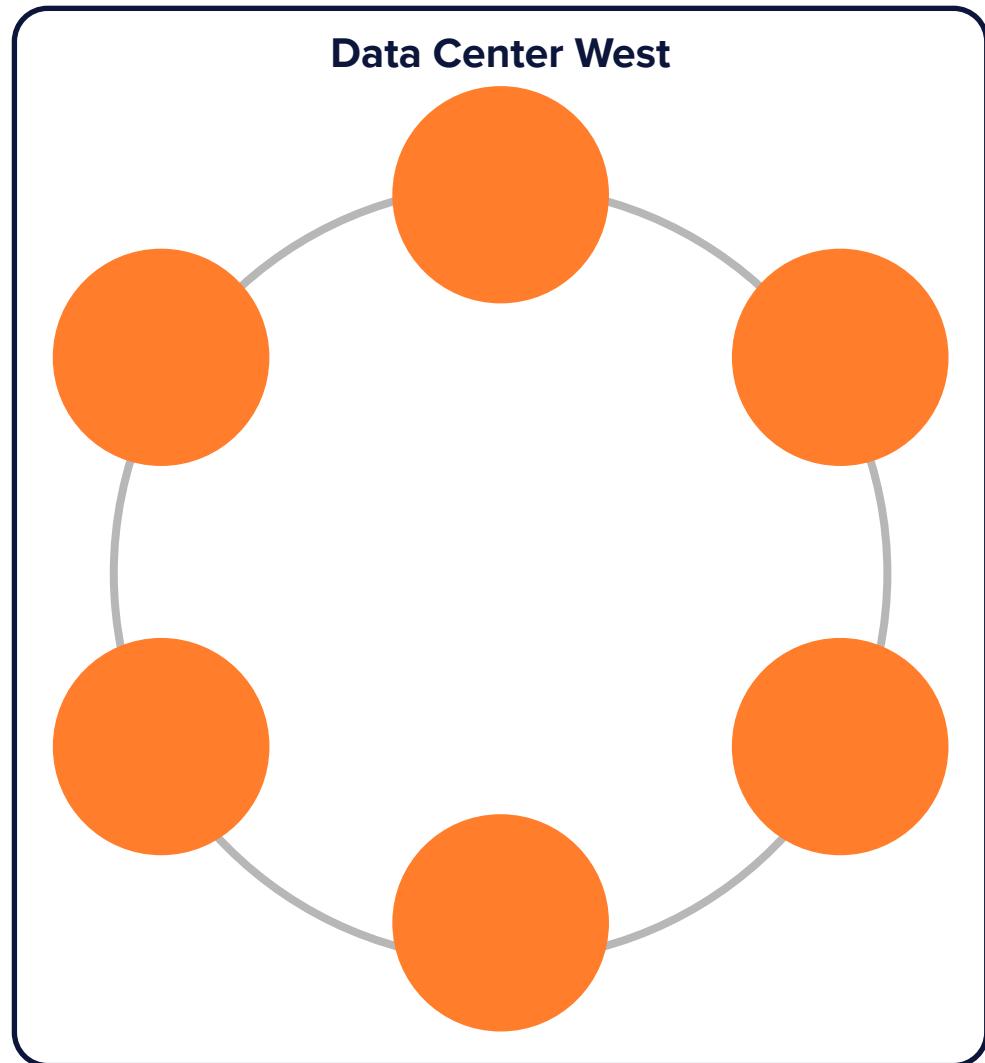


@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>



# Multi-Data Center Replication (RF=2 in each DC)

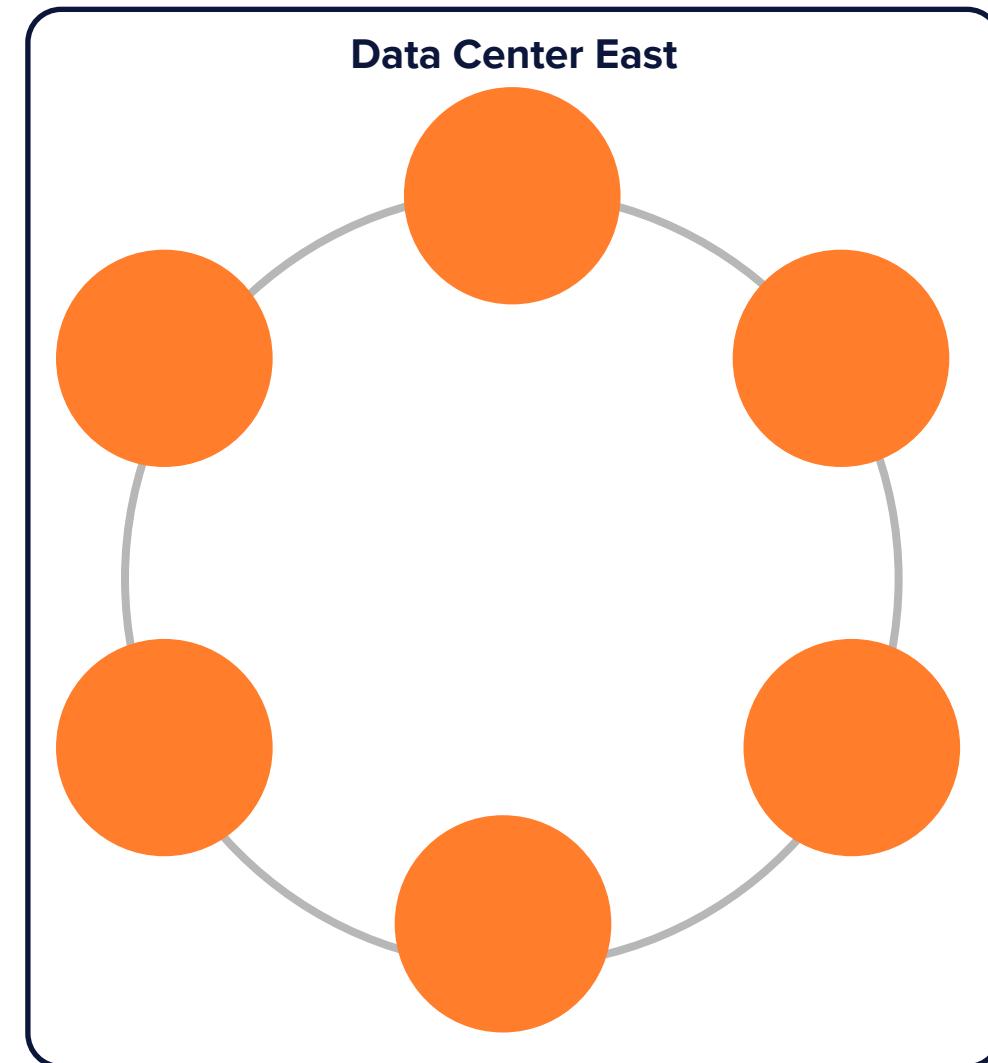
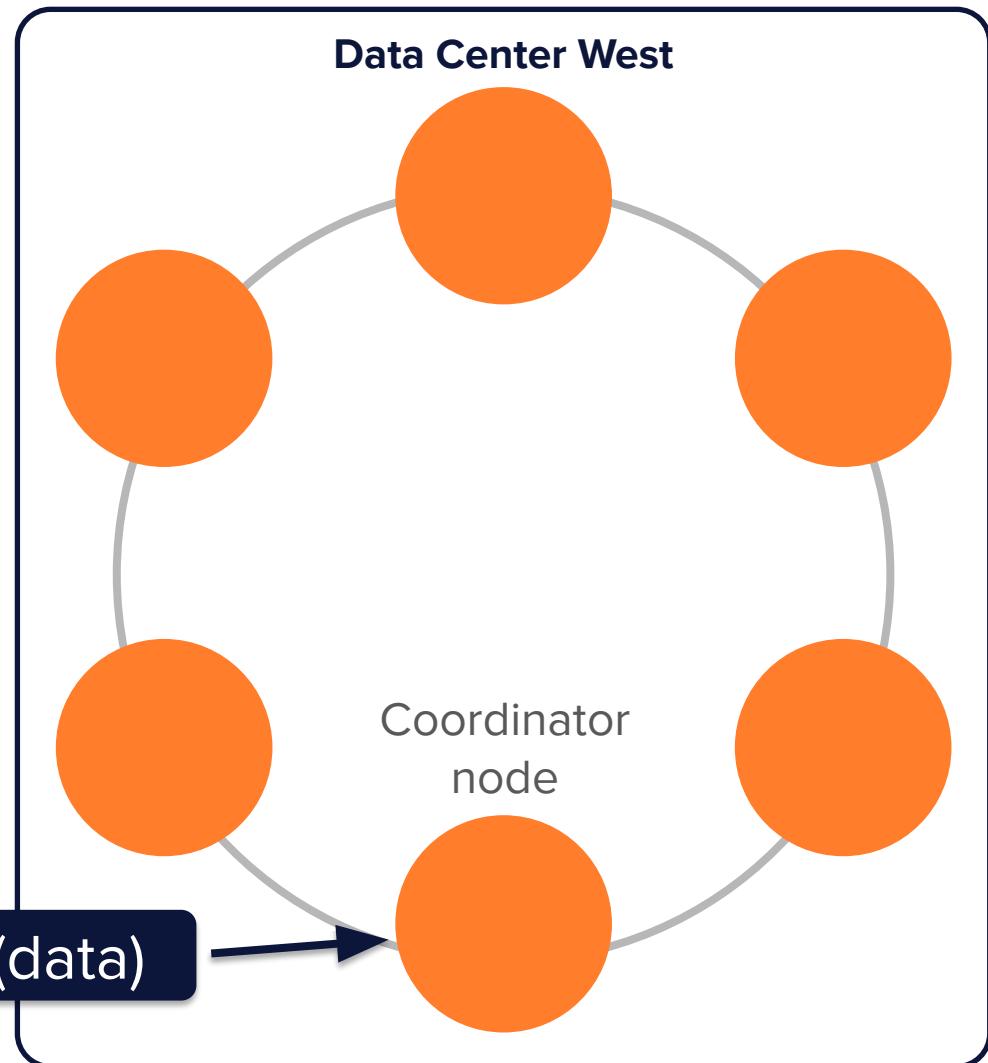


@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>



# Multi-Data Center Replication (RF=2 in each DC)

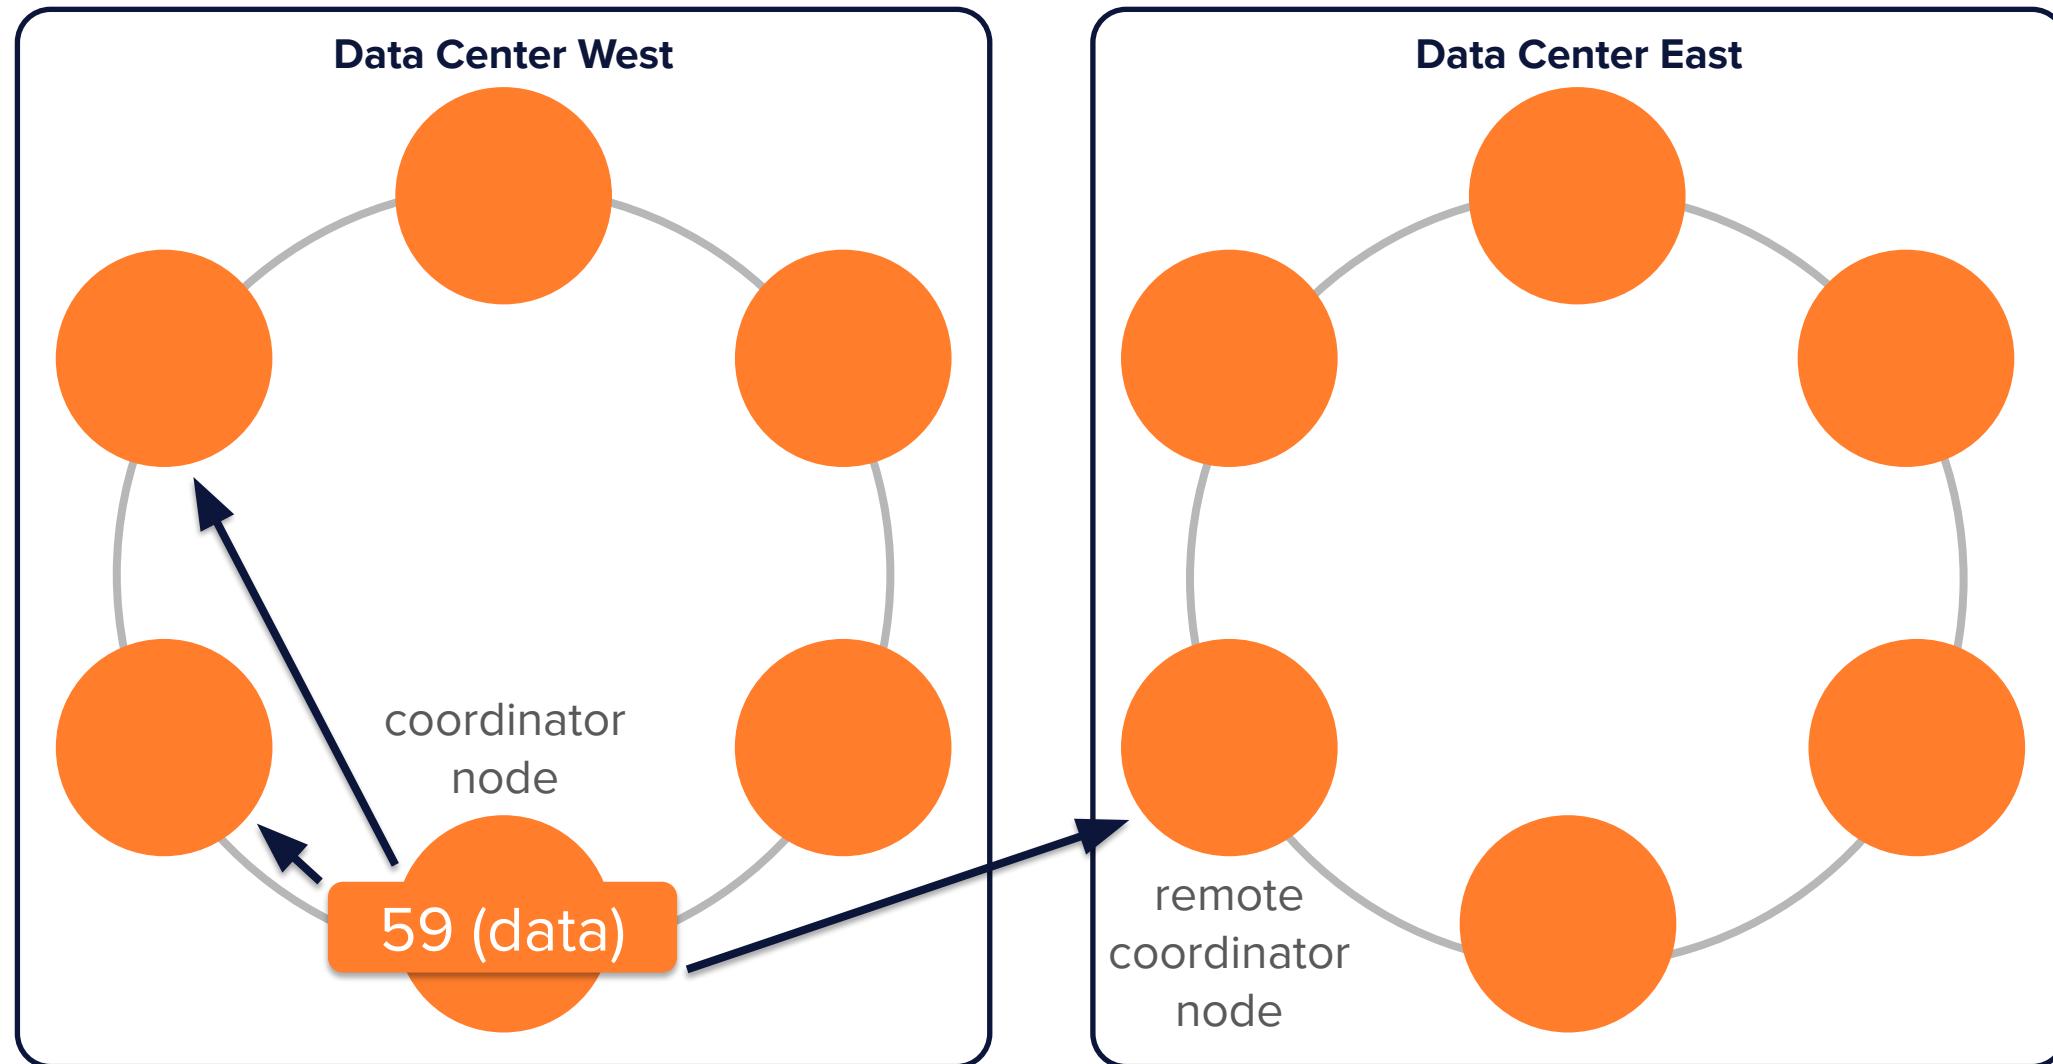


@DataStaxDevs #DataStaxDeveloperDay

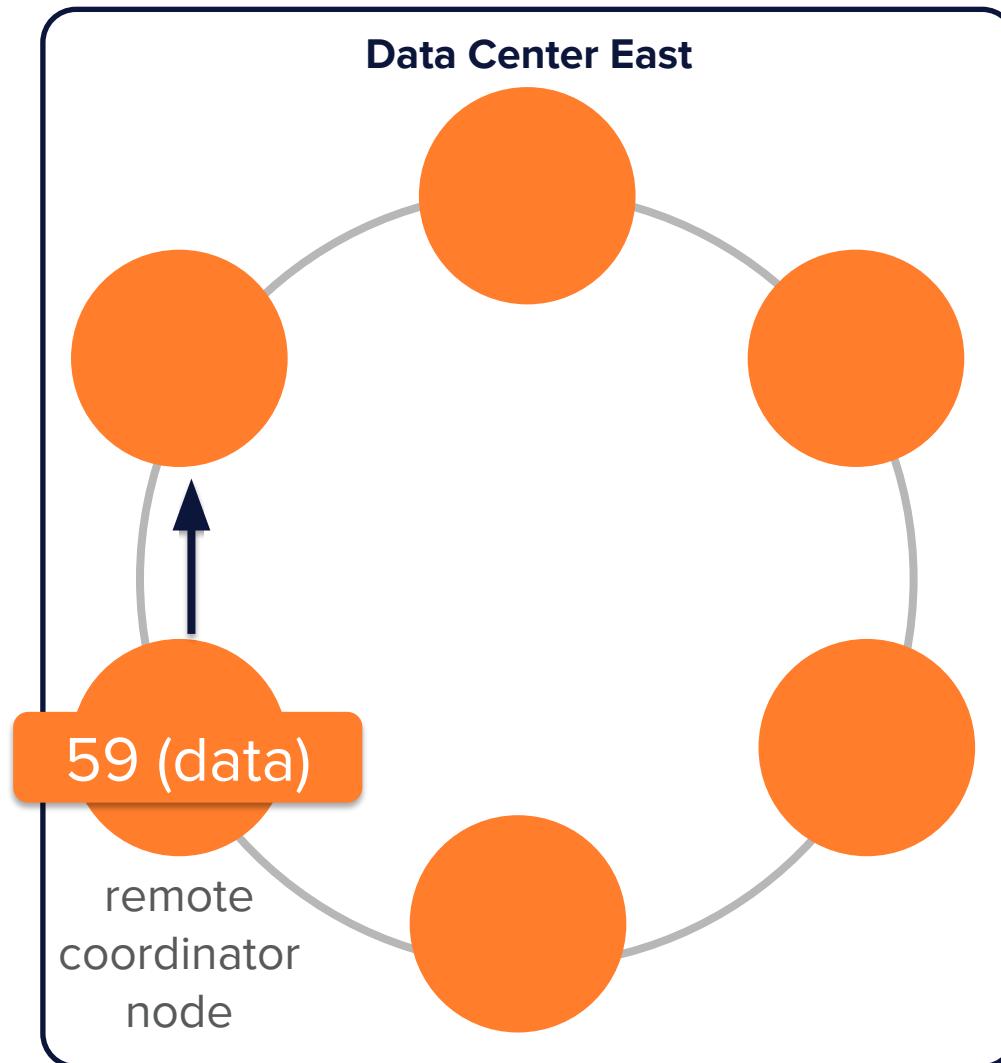
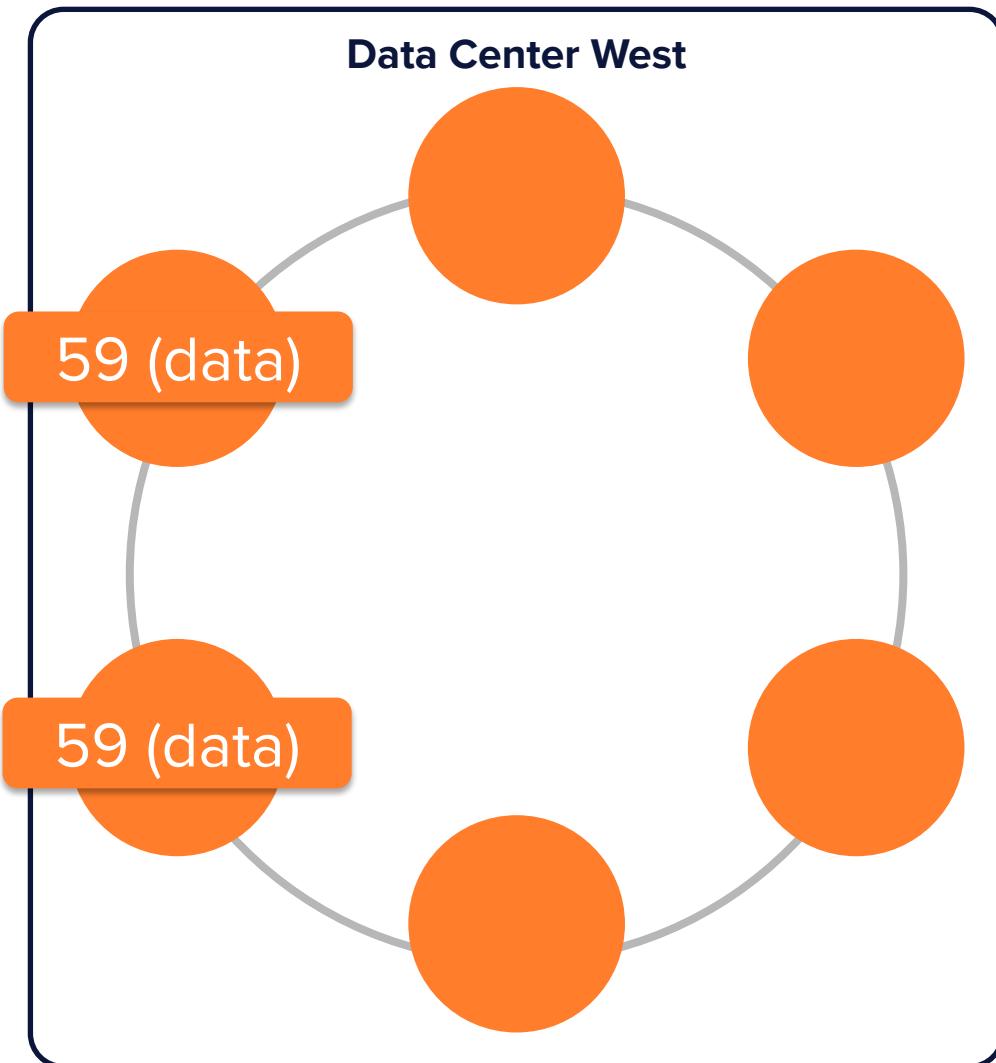
<https://community.datastax.com>



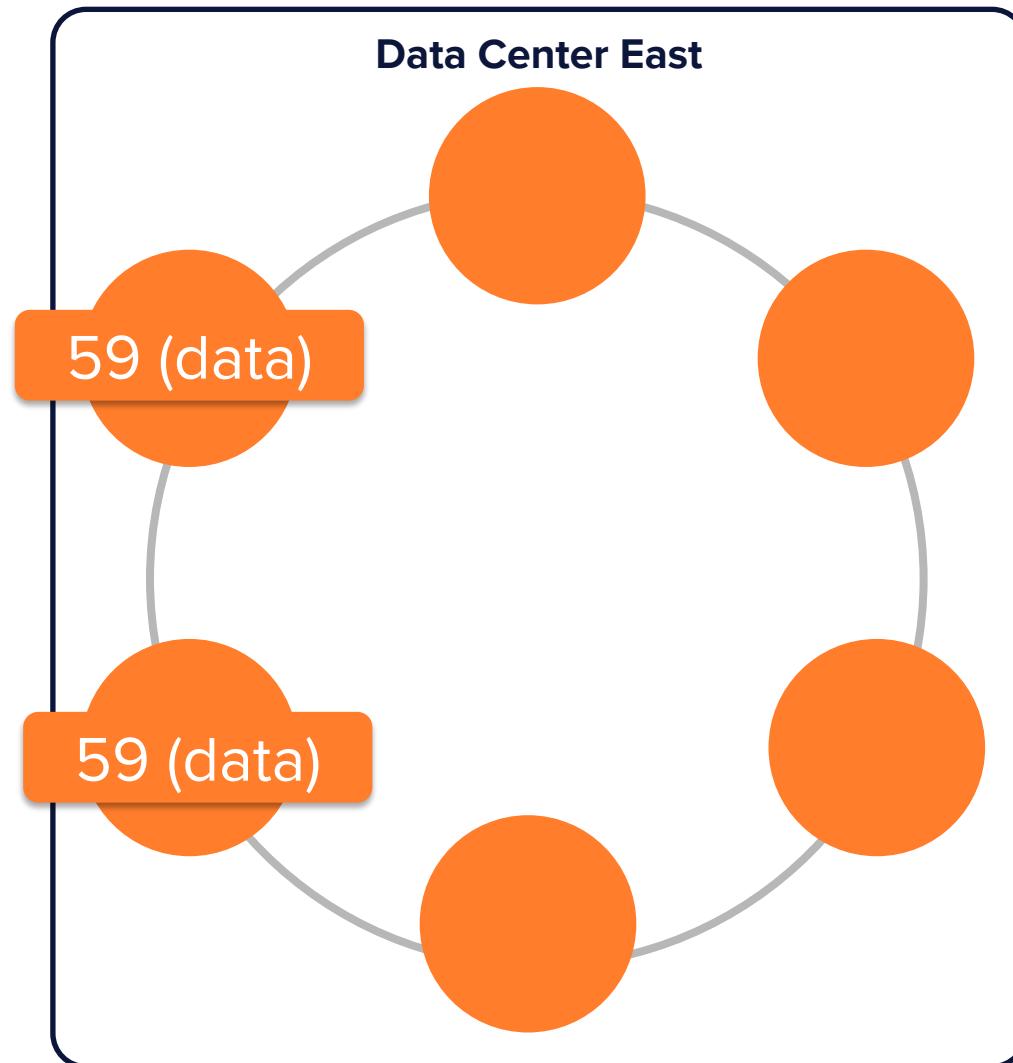
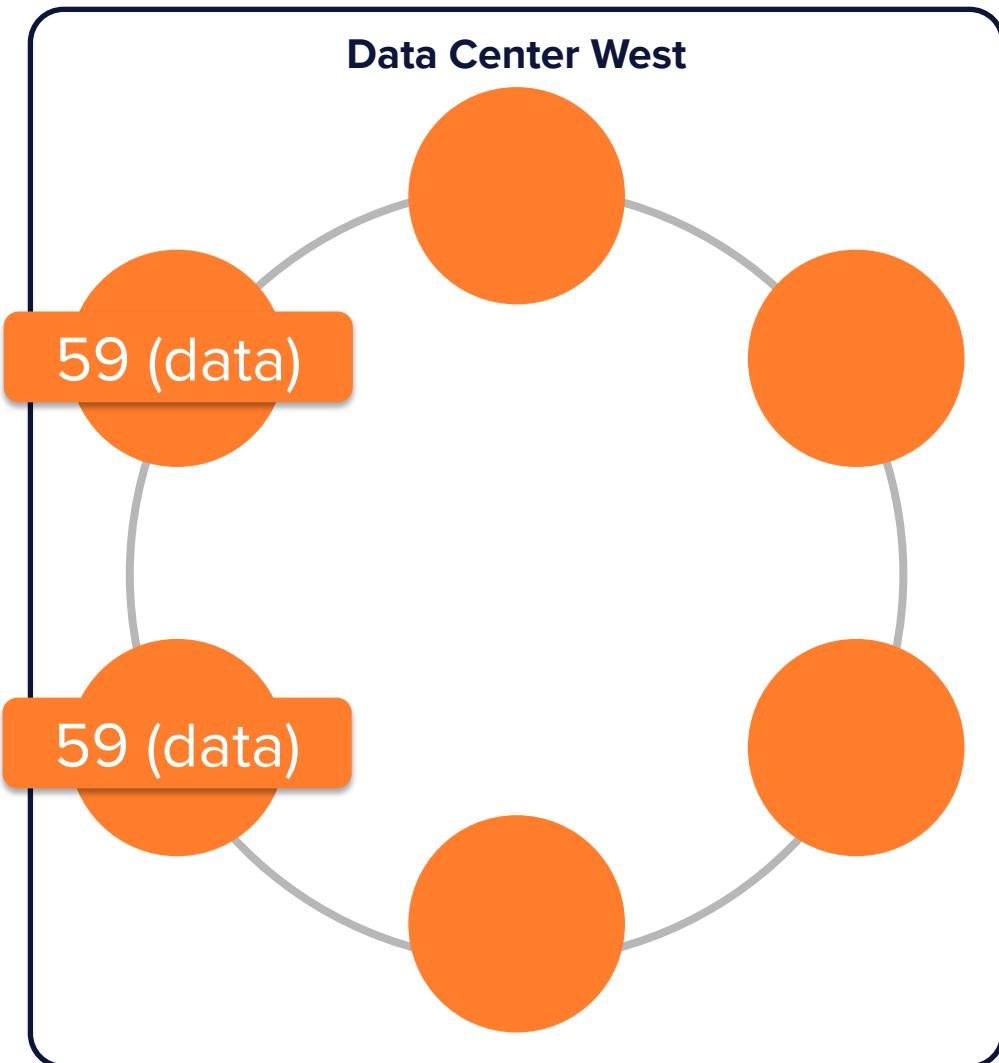
# Multi-Data Center Replication (RF=2 in each DC)



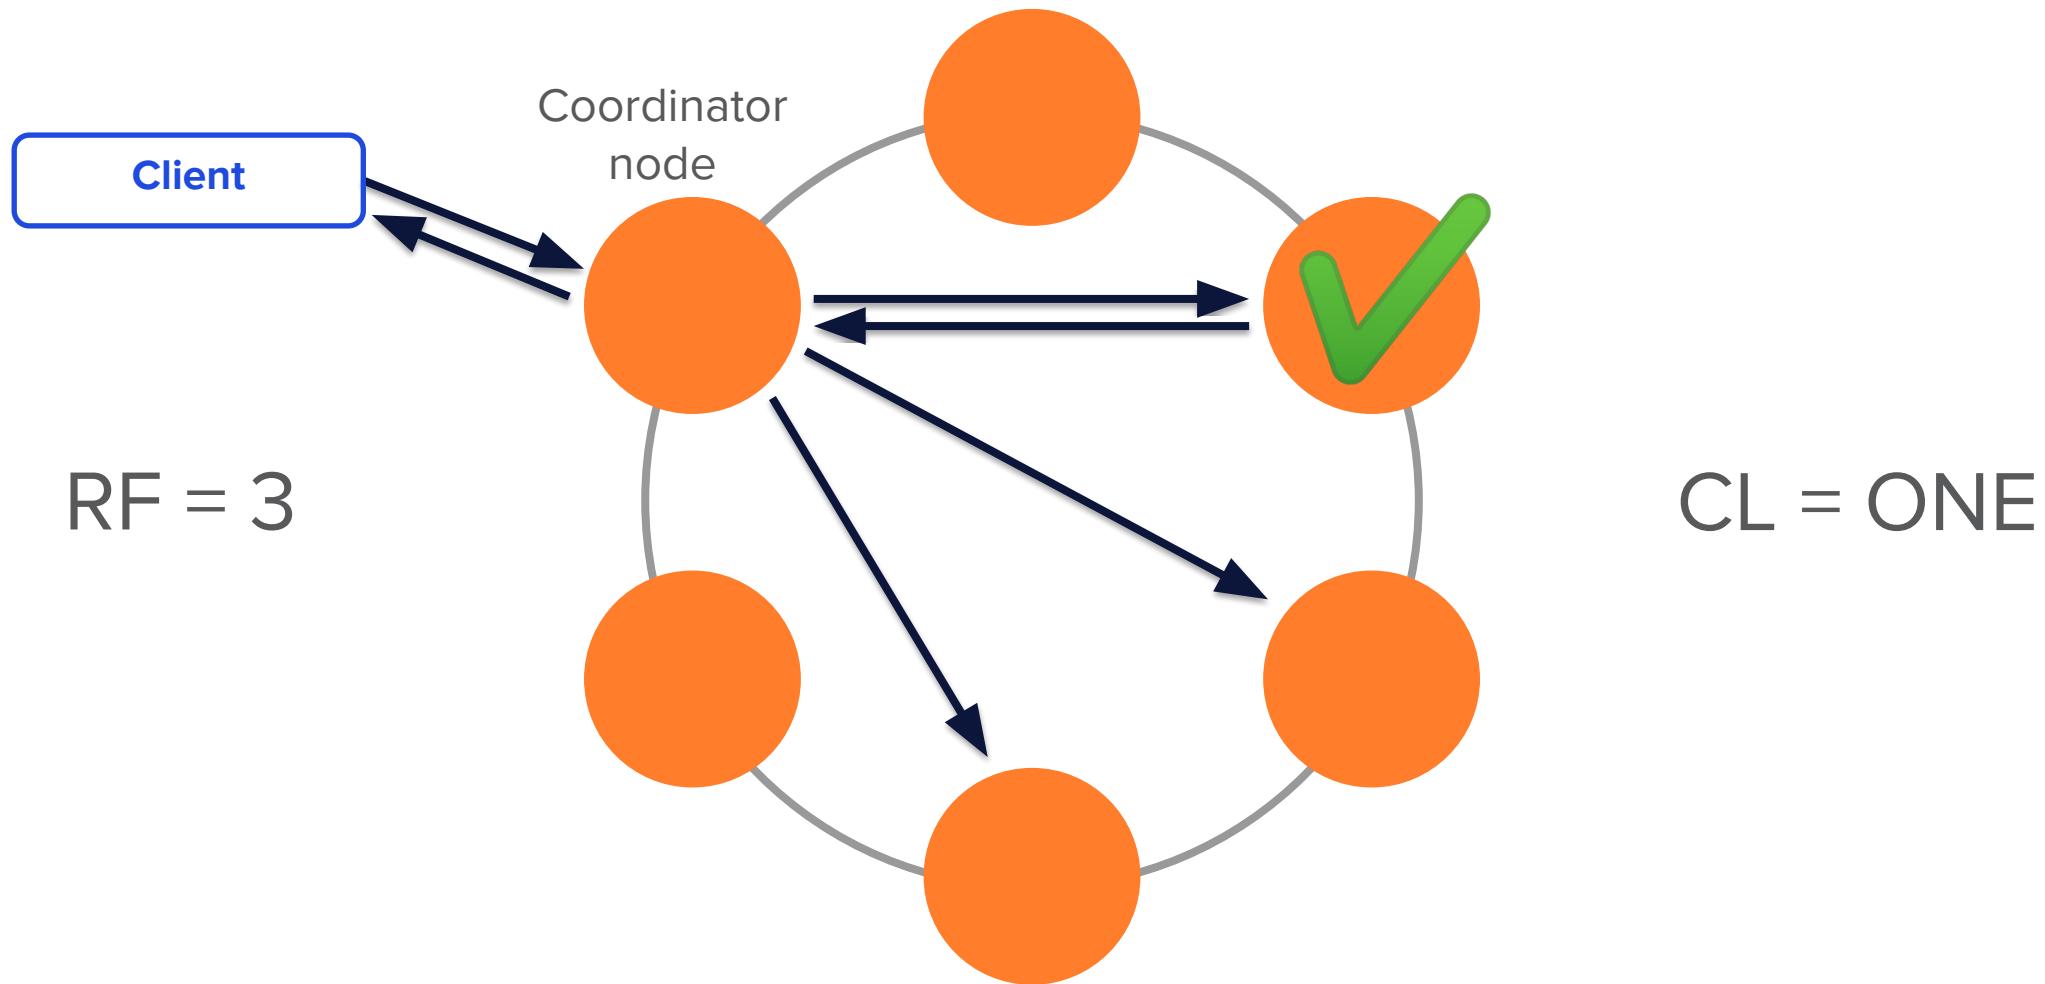
# Multi-Data Center Replication (RF=2 in each DC)



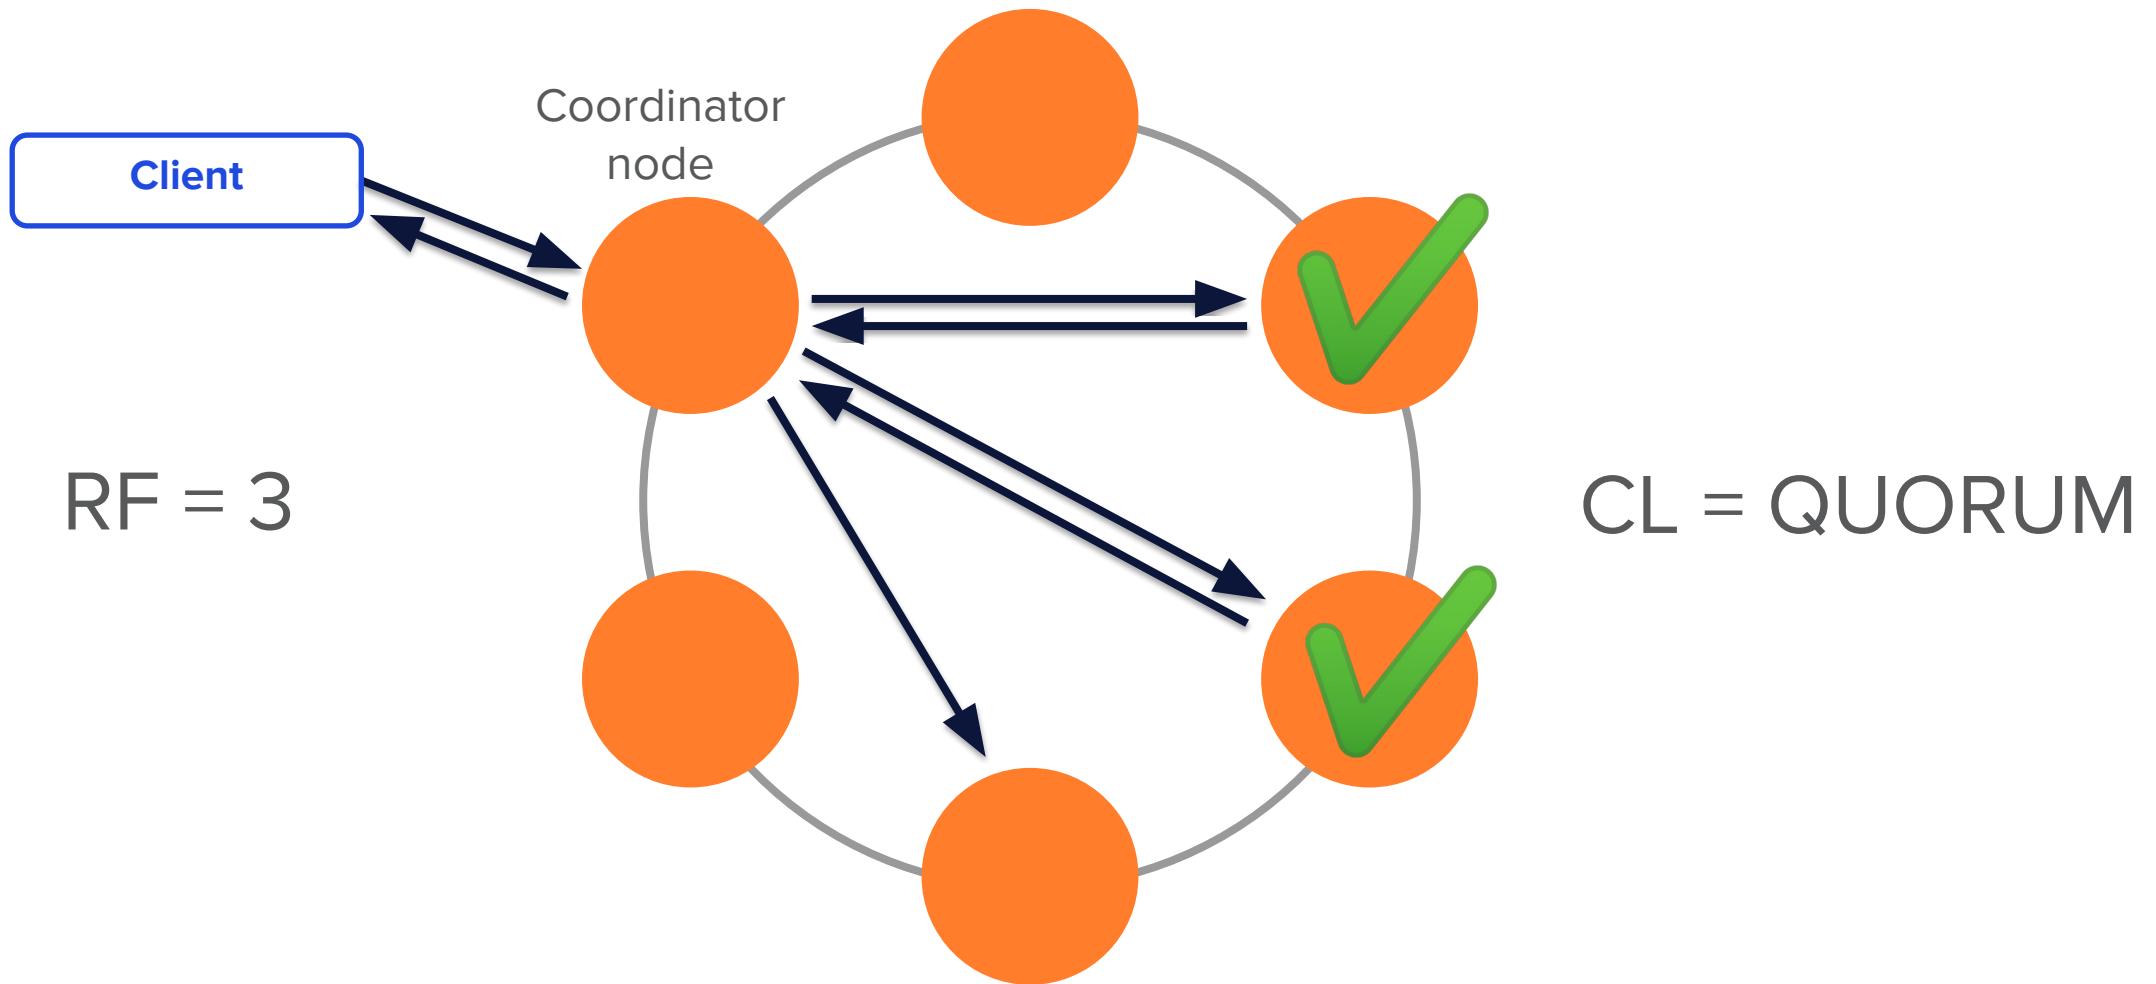
# Multi-Data Center Replication (RF=2 in each DC)



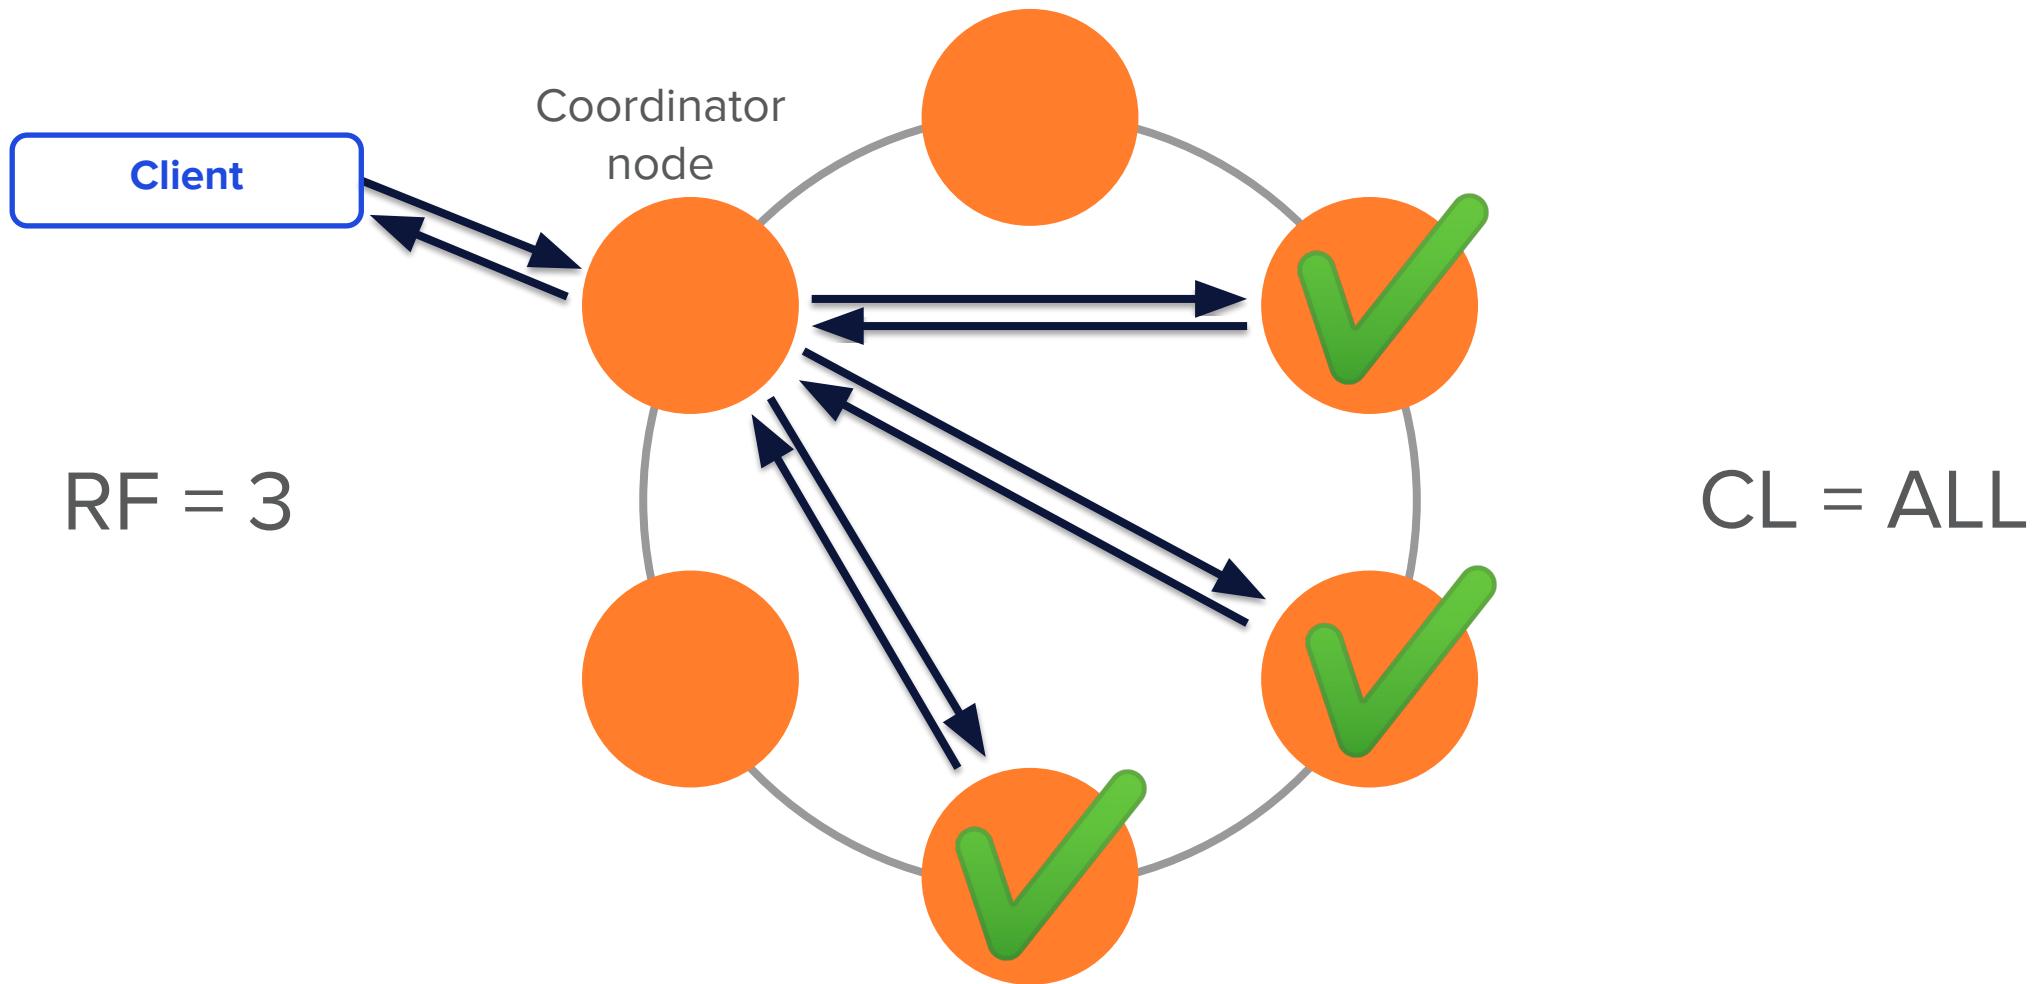
# Consistency Levels



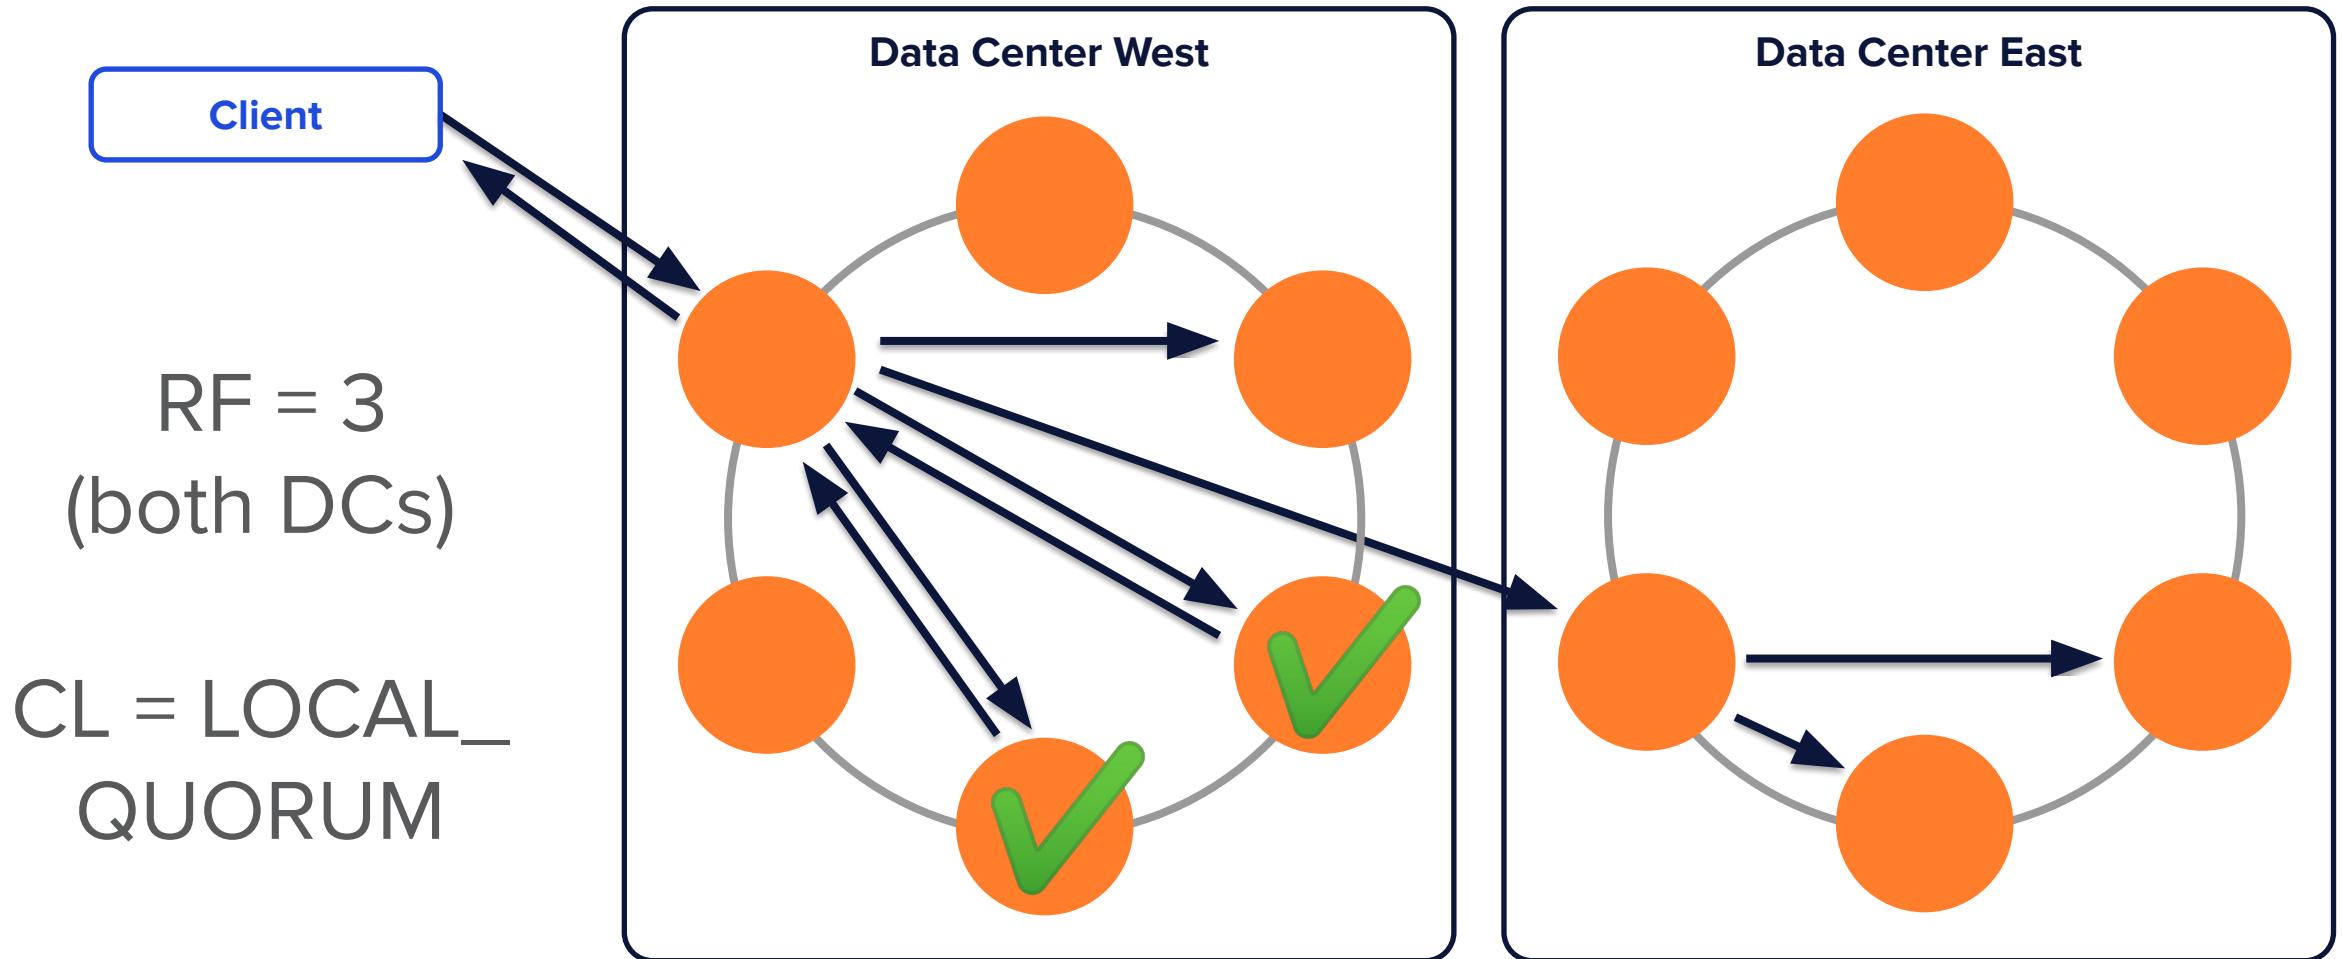
# Consistency Levels



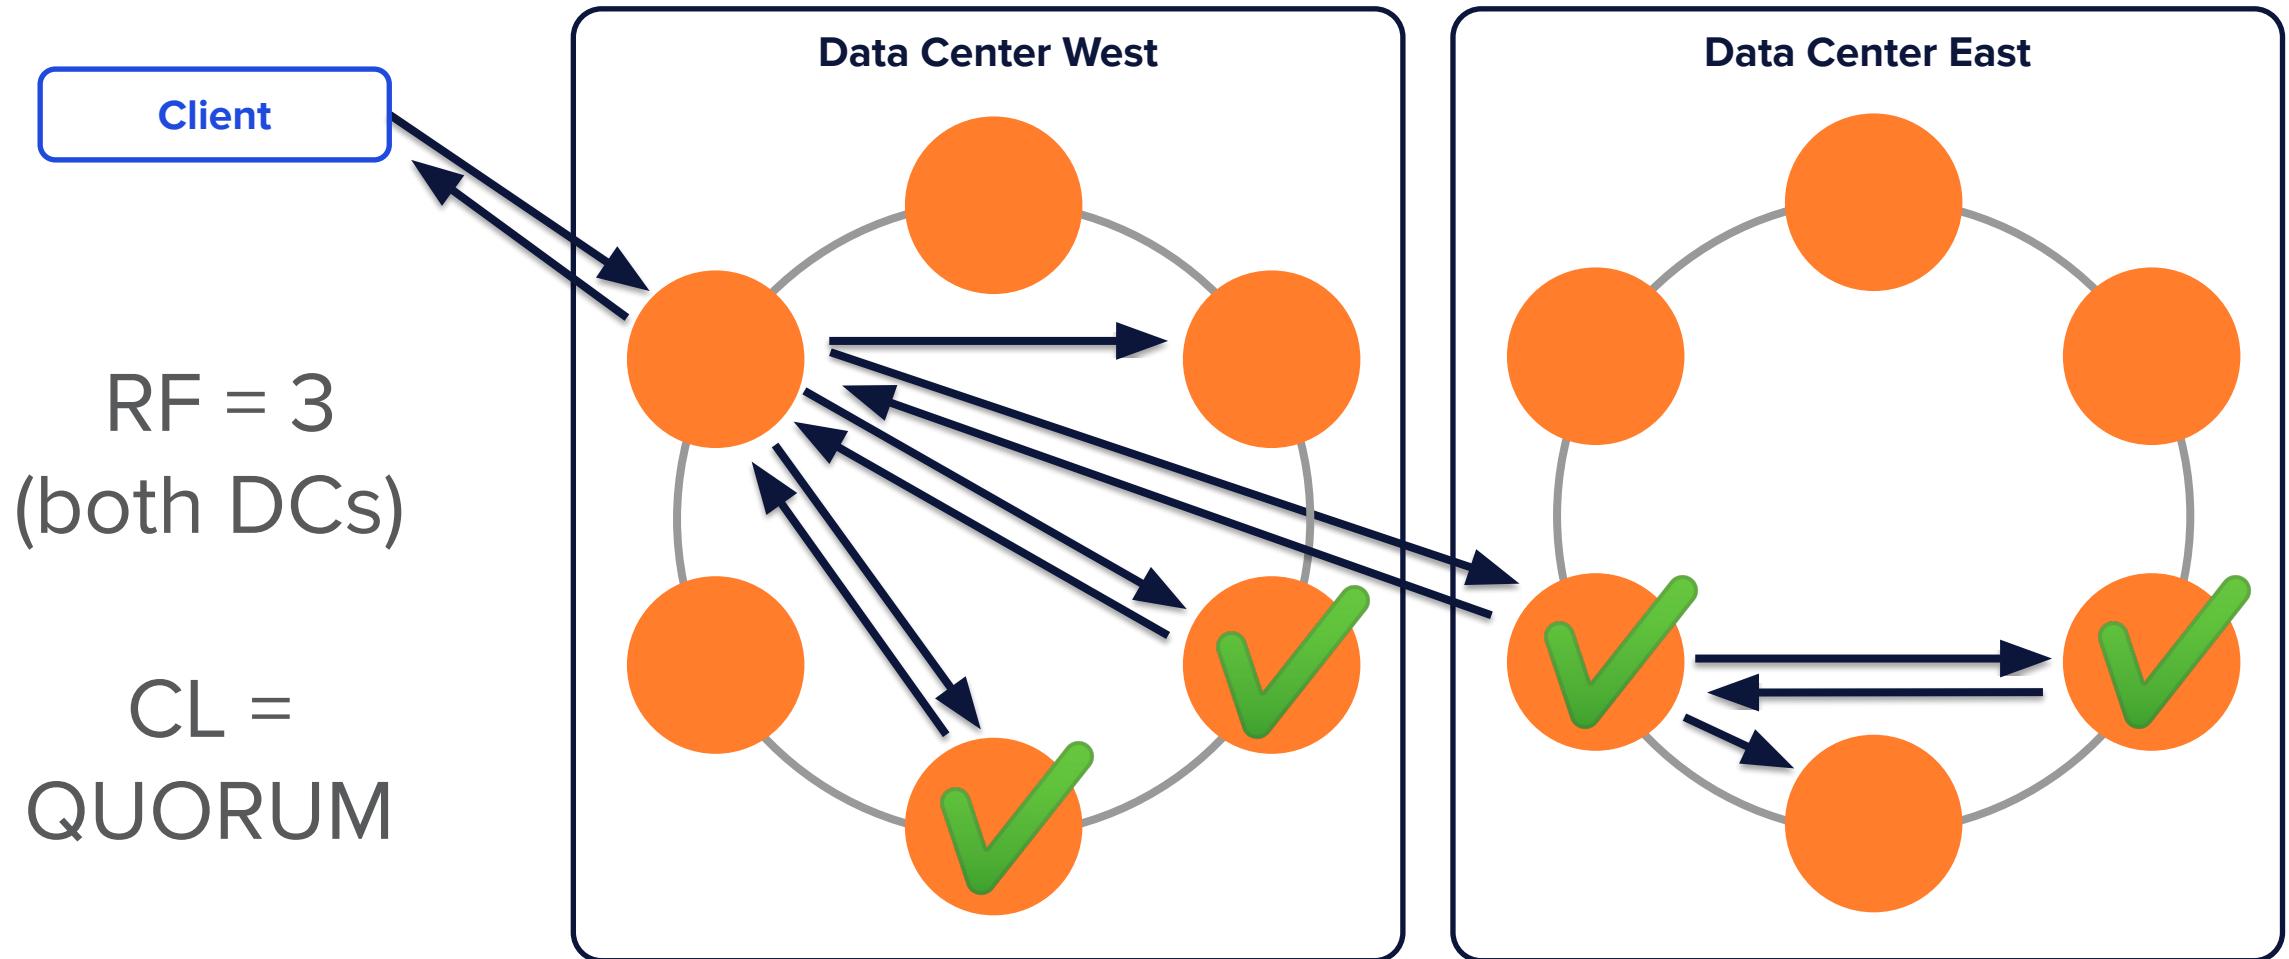
# Consistency Levels



# Consistency Across Data Centers



# Consistency Across Data Centers



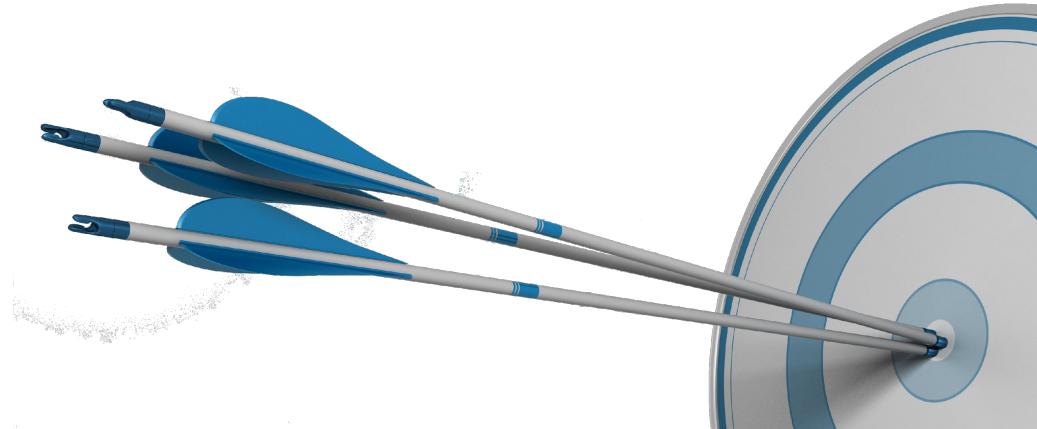
@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>

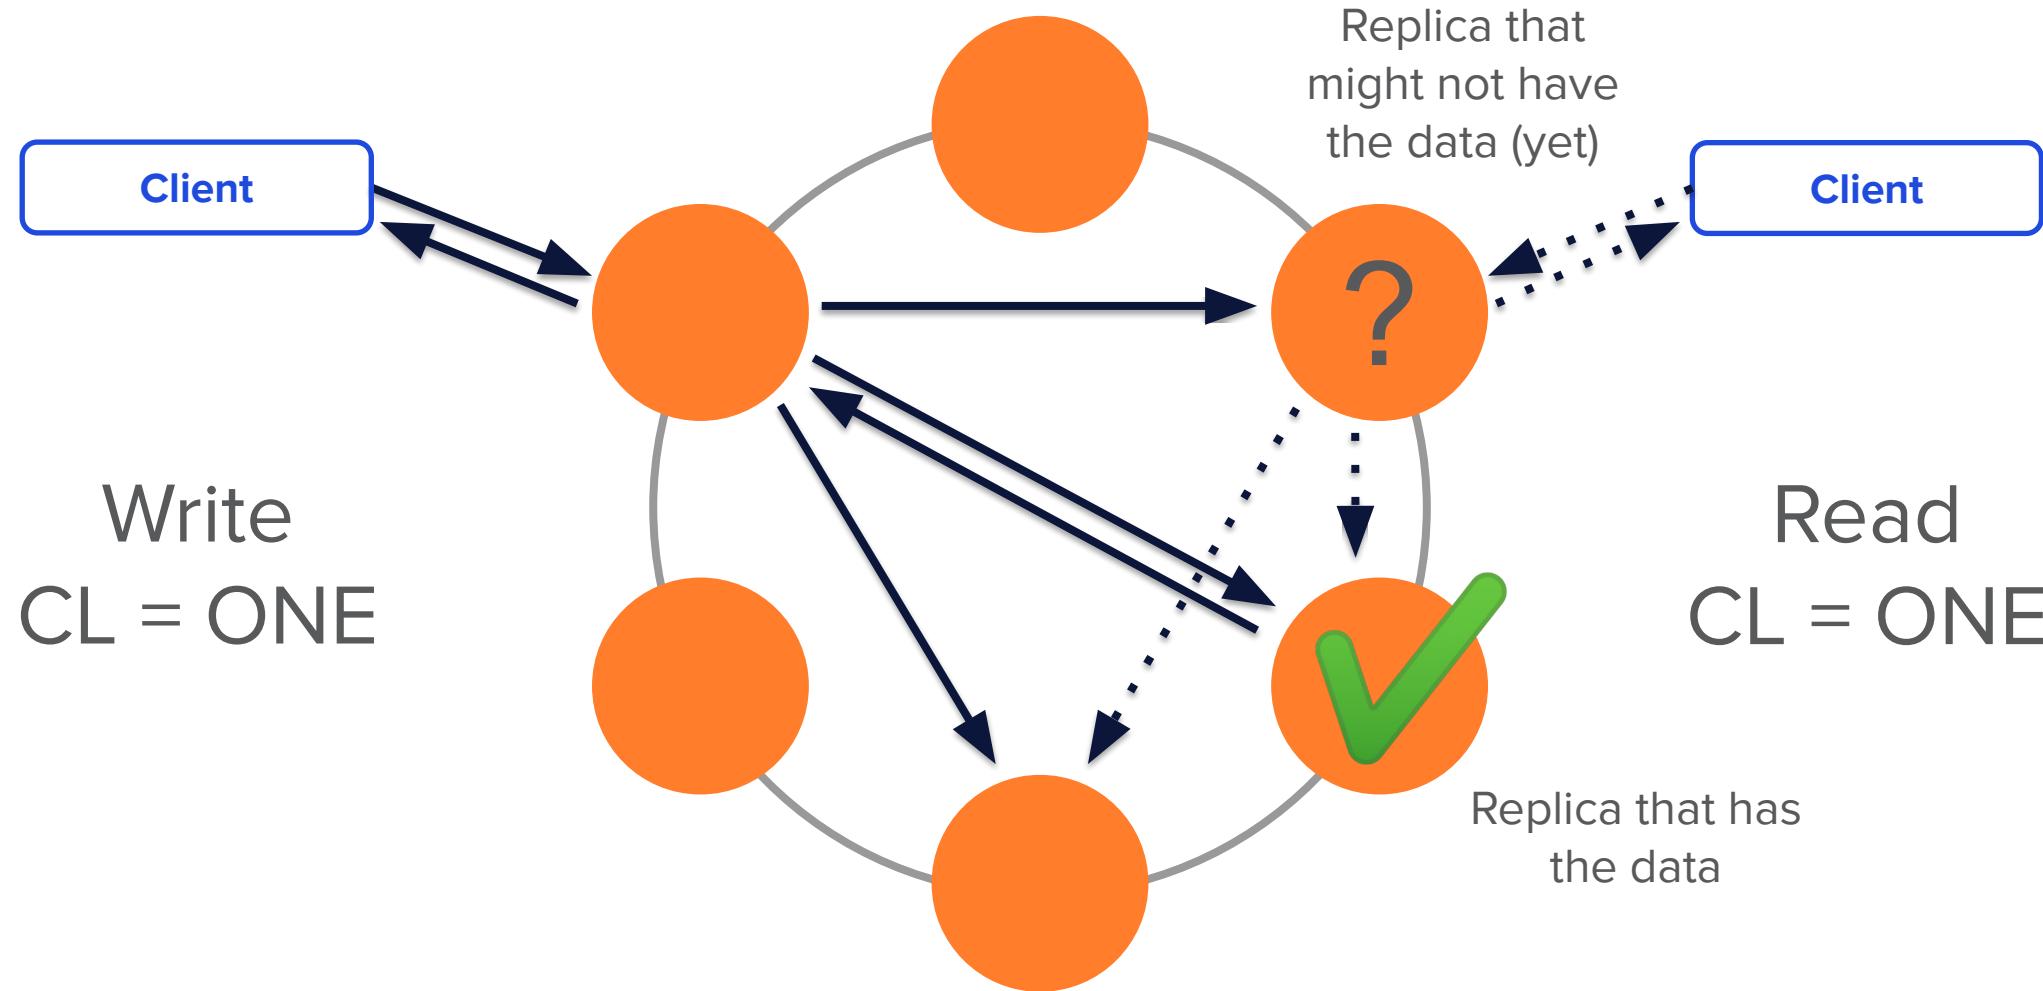


# Cassandra Data Availability Concepts

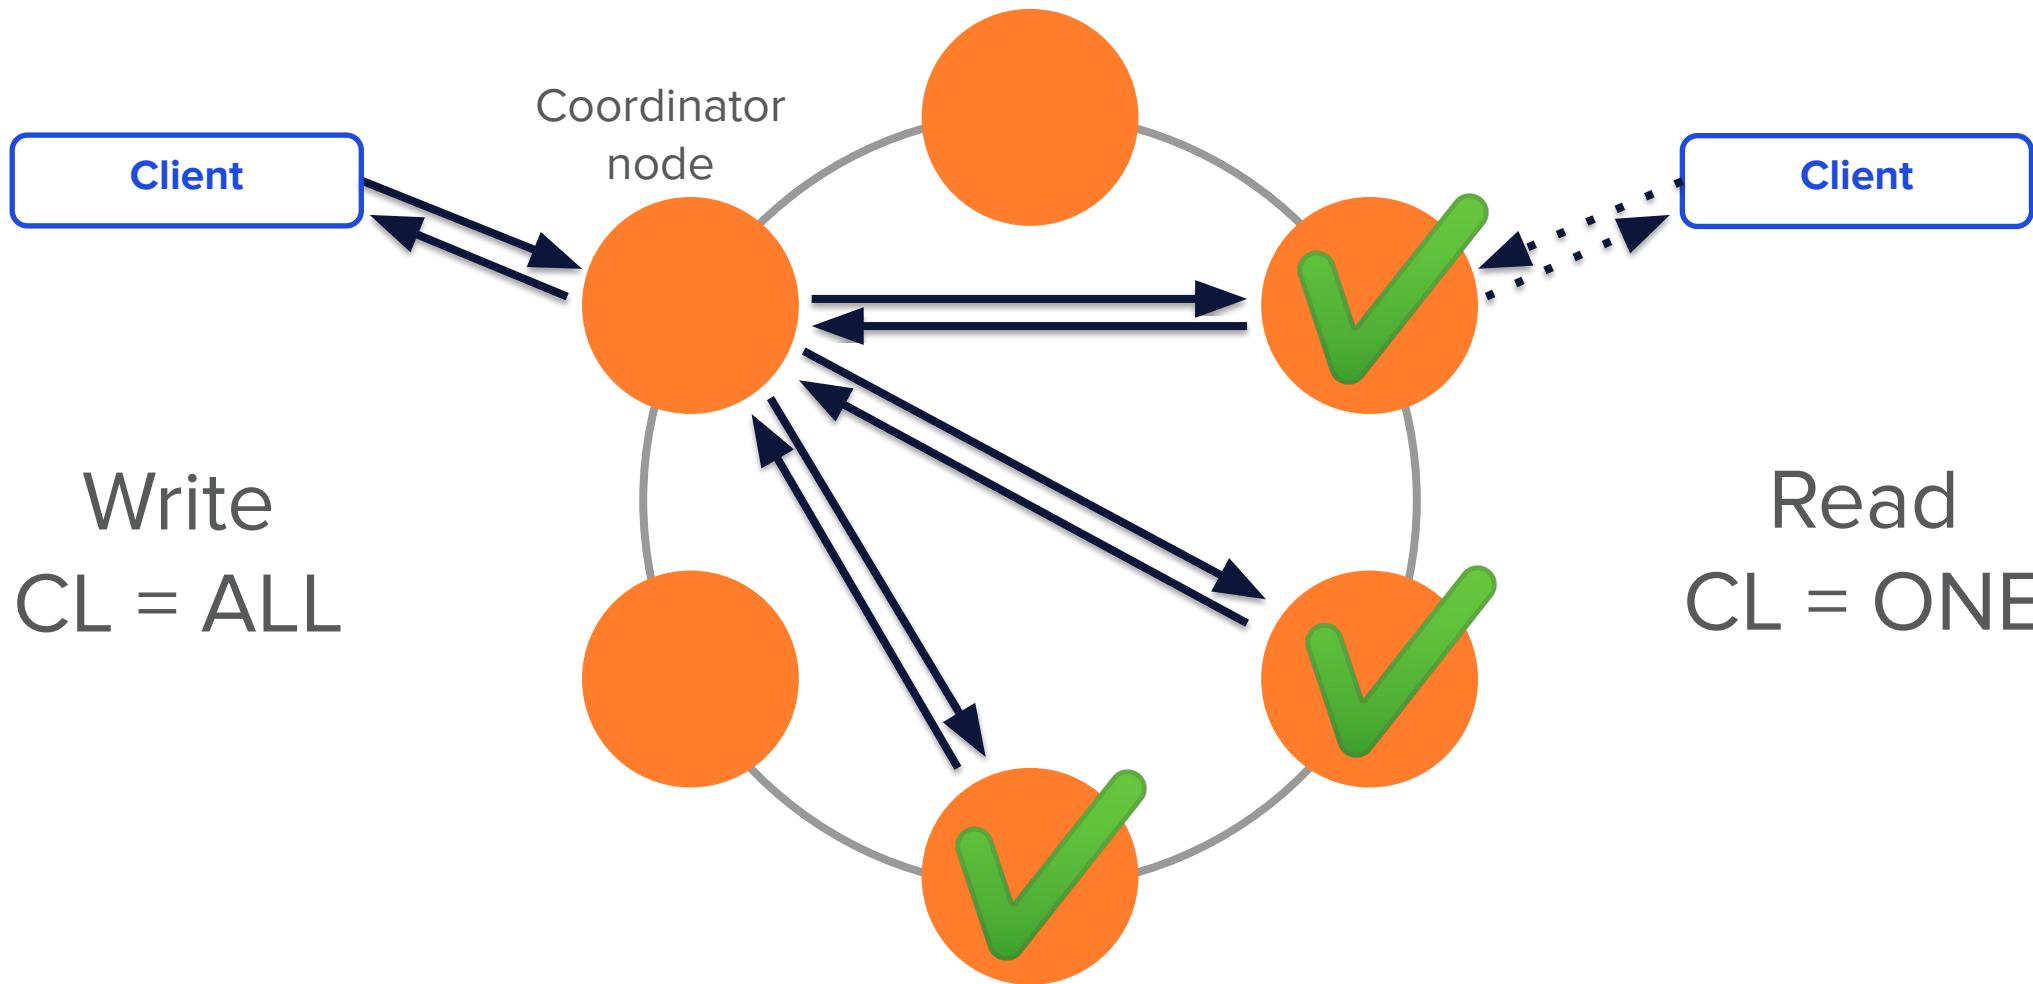
- Replication Factor – the number of copies of your data
  - Increasing replication increases chances of availability
  - Increasing replication increases chances of inconsistency
- Consistency Level – the number of acknowledged copies read/written
- Immediate Consistency: number of writes + number of reads > replication factor
- EXAMPLE:
  - Replication Factor = 3
  - Read/Write Consistency Level = QUORUM
  - $2 + 2 = 4 > 3$  



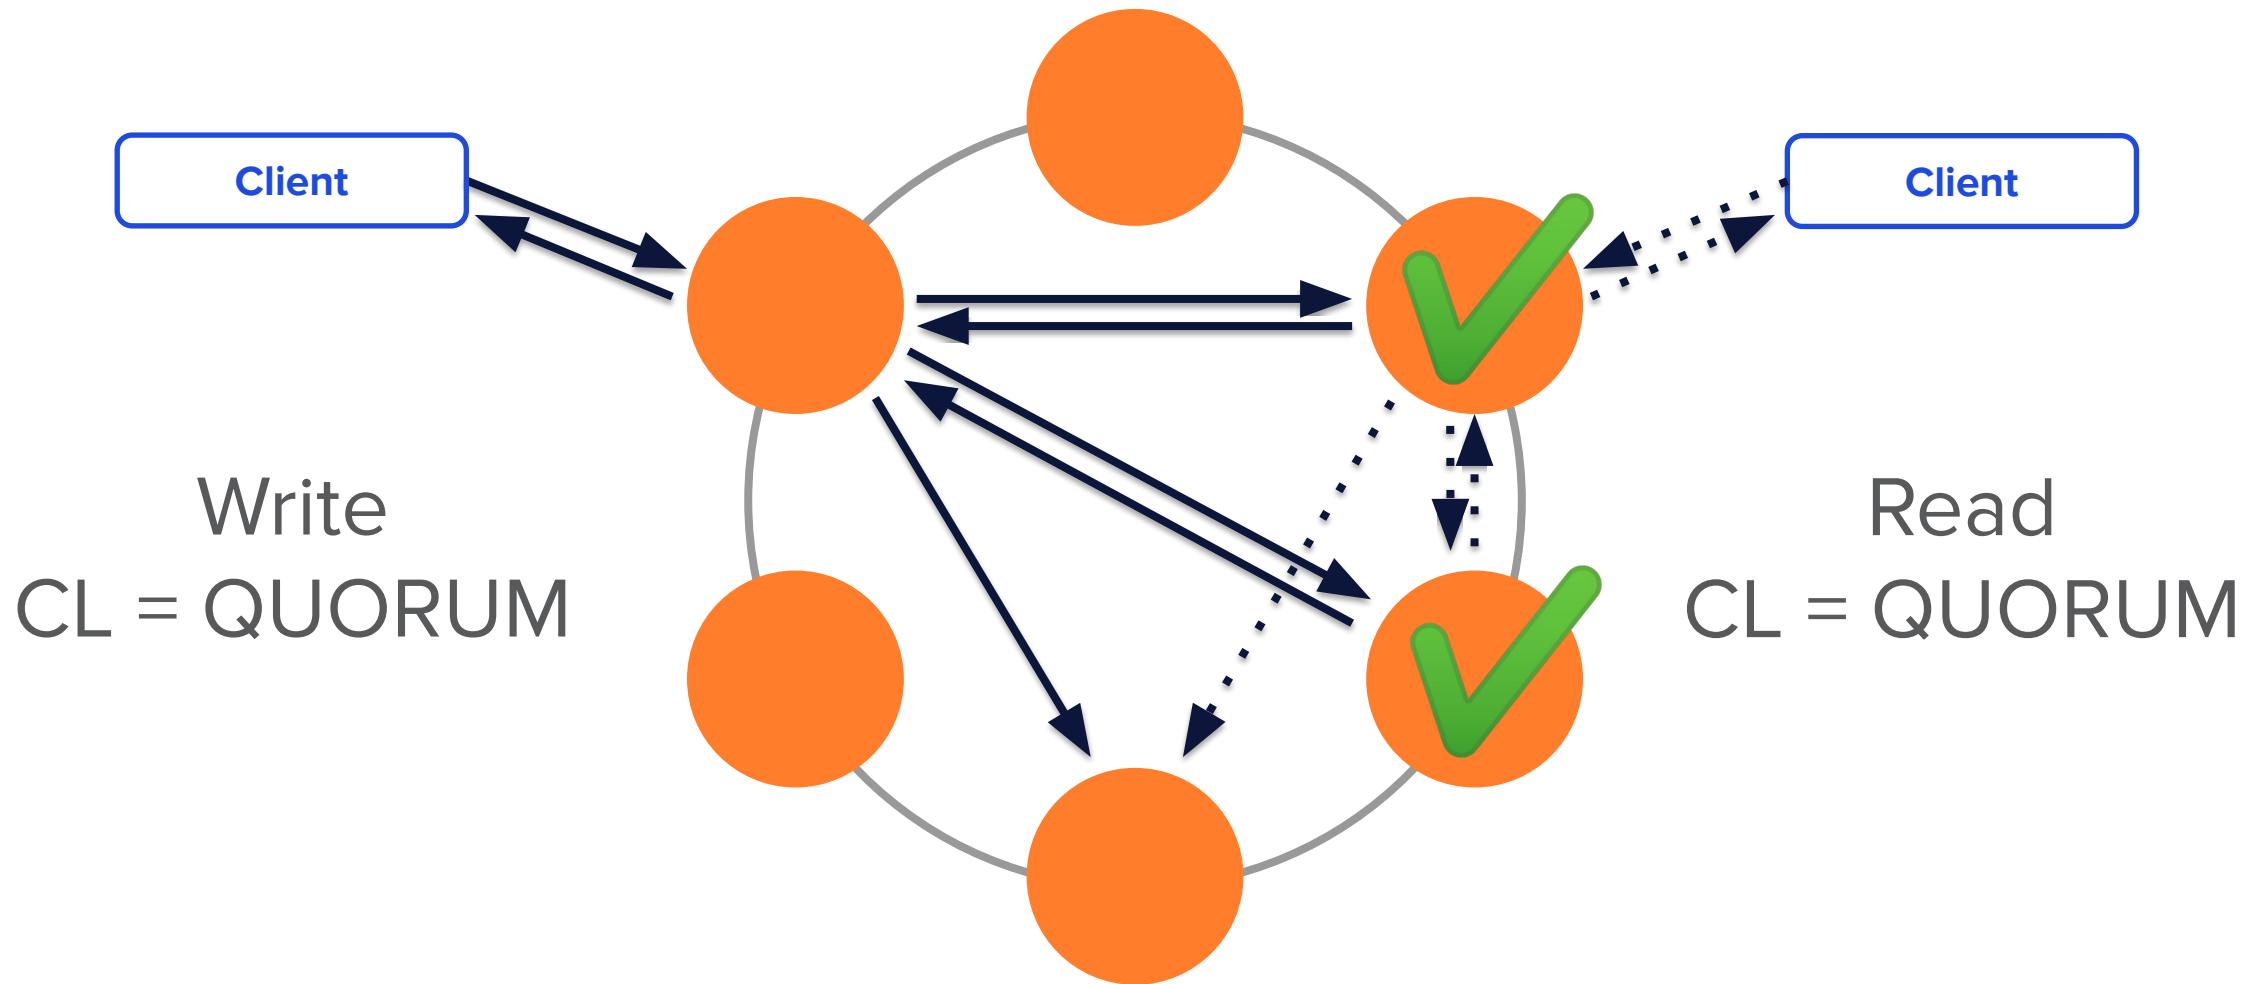
# Weak(er) Consistency



# Immediate Consistency – One Way



# Immediate Consistency – A Better Way



# Consistency Settings

- Weakest to strongest

Setting	Description
ANY	Storing a hint at minimum is satisfactory
ONE, TWO, THREE	Checks closest node(s) to coordinator
QUORUM	Majority vote, $(\text{sum\_of\_replication\_factors} / 2) + 1$
LOCAL_ONE	Closest node to coordinator in same data center
LOCAL_QUORUM	Closest quorum of nodes in same data center
EACH_QUORUM	Quorum of nodes in each data center, applies to writes only
ALL	Every node must participate

# Consistency Tradeoffs

- The higher the consistency, the less chance you may get stale data
  - Pay for this with POSSIBLE latency
  - Depends on your situational needs
- IoT and Sensor systems with write-heavy workloads may benefit from CL=ONE for writes if data isn't required immediately



# Understanding Use Cases

Scalability

High Throughput
High Volume



Heavy Writes
Heavy Reads



Event Streaming	
Log Analytics	
Internet of Things	
Other Time Series	

Availability

Mission-Critical
------------------



No Data Loss
Always-on



Caching	
Pricing	
Market Data	
Inventory	

Distributed

Global Presence
Workload Mobility



Compliance /
GDPR



Banking	
Retail	
Tracking / Logistics	
Customer Experience	

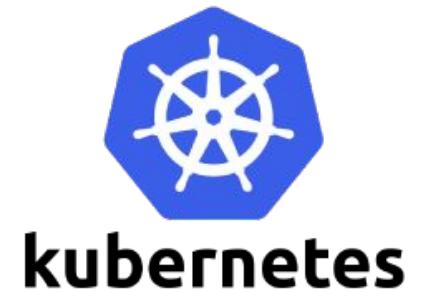
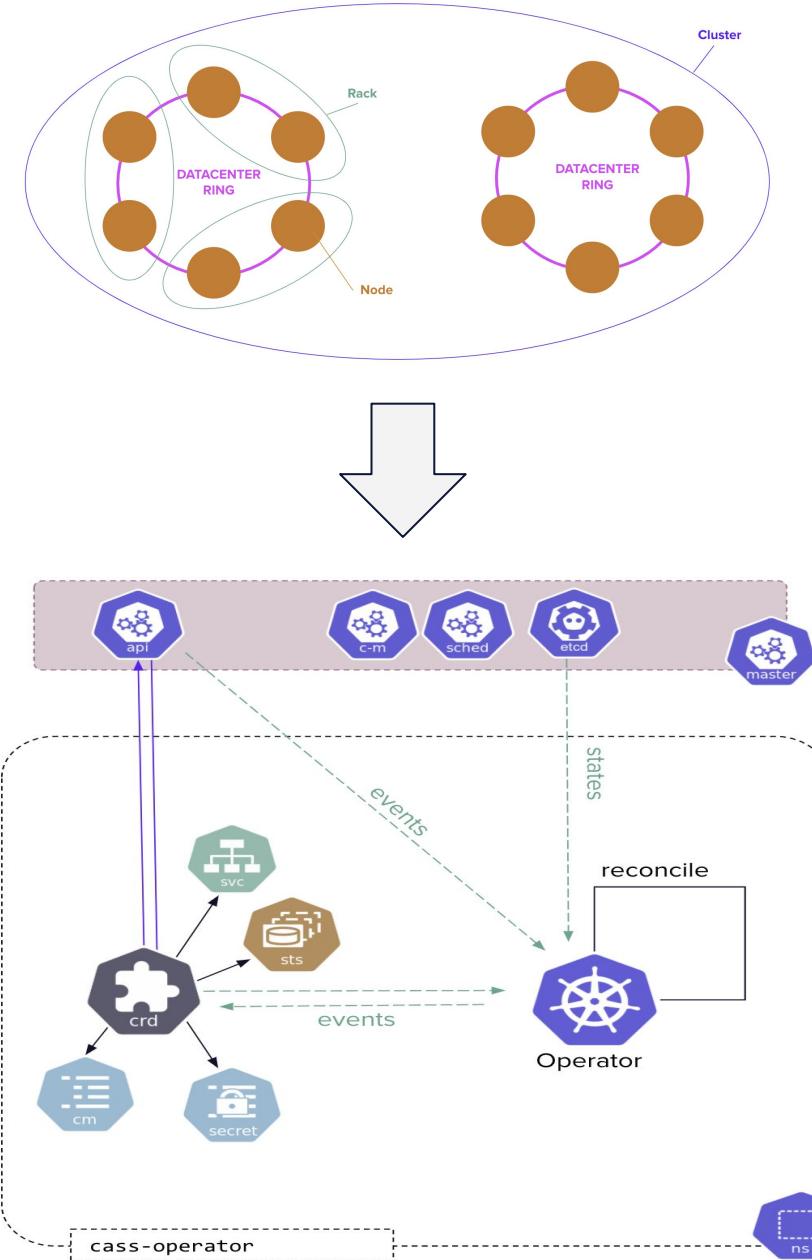
Cloud-native

Modern Cloud Applications
---------------------------



API Layer	
Hybrid-cloud	
Enterprise Data Layer	
Multi-cloud	

# YOUR MISSION SHOULD YOU CHOOSE TO ACCEPT IT



# Online Workshops



## Section #0 : Setup Your Cluster

<https://github.com/DataStax-Academy/kubecon-cassandra-workshop>

- Your Laptop  
**OR**
- Our Cloud Instance

### 5. Install Kind

kind ( kind ) is a tool for running local Kubernetes clusters using Docker container "nodes". kind was primarily designed for testing Kubernetes itself, but may be used for local development or CI. Please refer to [Reference Documentation](#) for more detailed instructions.



: To install on windows please download the [executable](#) and place it on the PATH. You can also use [Chocolatey](#) very clever package manager for windows.

```
choco install kind
```



: To install on MAC OS please use the following [homebrew](#) commands:

```
brew install kubectl
```



: To install on linux (centOS) you can use the following commands

```
curl -Lo ./kind https://github.com/kubernetes-sigs/kind/releases/download/v0.7.0/kind-$(uname)-amd64
chmod +x ./kind
sudo mv ./kind /usr/local/bin/kind
```

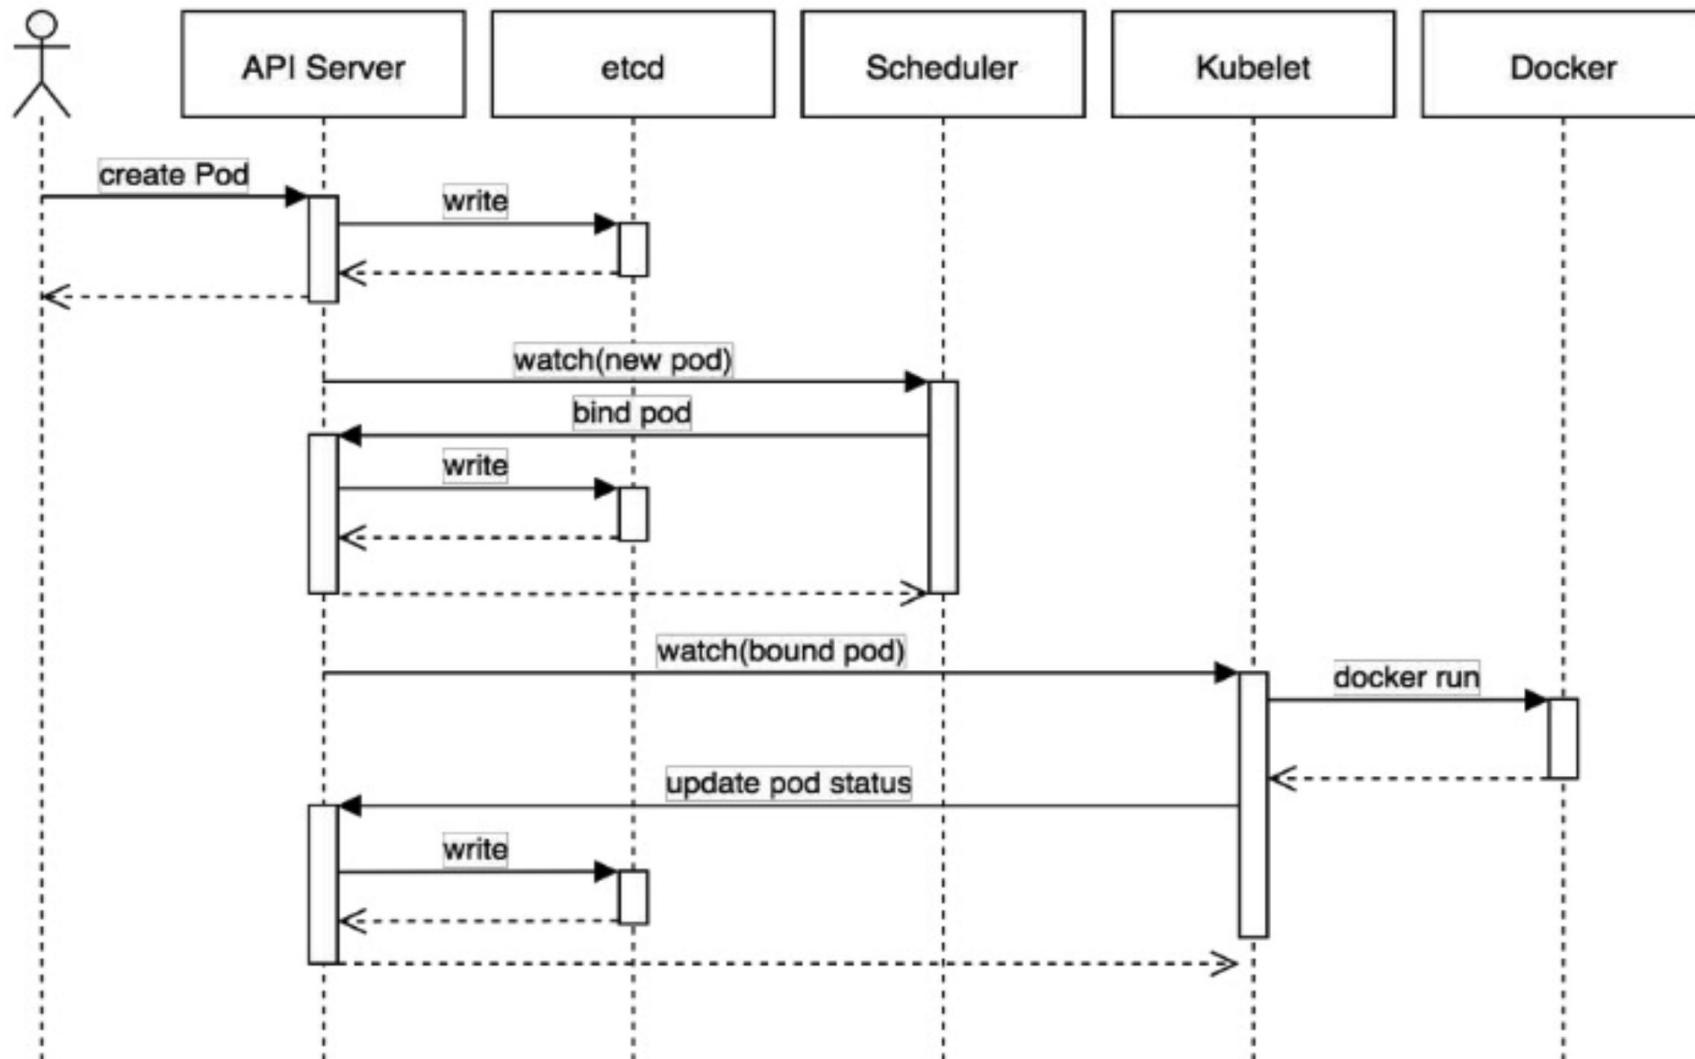
Check that the installation is successful. Starting from now all command will be the same on each platform, as such we will keep providing a single command. We will mark with a blue book the command ( ) and a green book ( ) to show expected result.

# Apache Cassandra™ with Kubernetes

- 1 Housekeeping and Quizz
- 2 Cassandra Basics
- 3 Kubernetes Operators
- 4 Cass Operator in Deep
- 5 Grafana | Prometheus
- 6 Resources



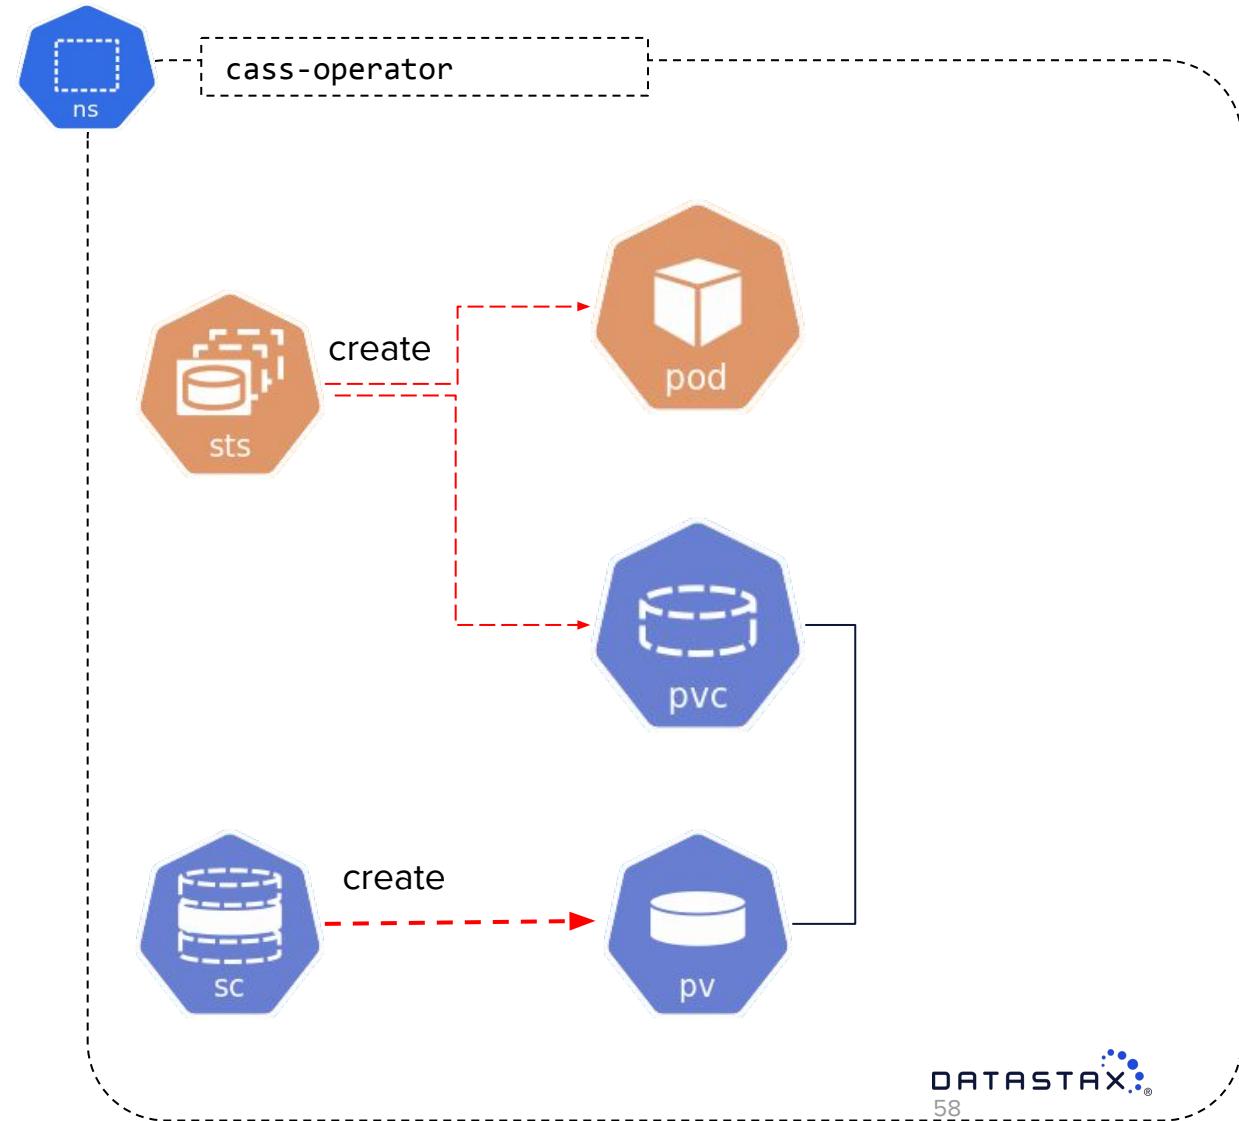
# Pod Lifecycle



# K8s Primitives : StatefulSet



**StatefulSet:** StatefulSet represents a set of pods with consistent identities. Identities are defined as: network, storage.



# K8s Primitives : Service



**Ingress** is a collection of rules that allow inbound connections to reach the endpoints defined by a backend.



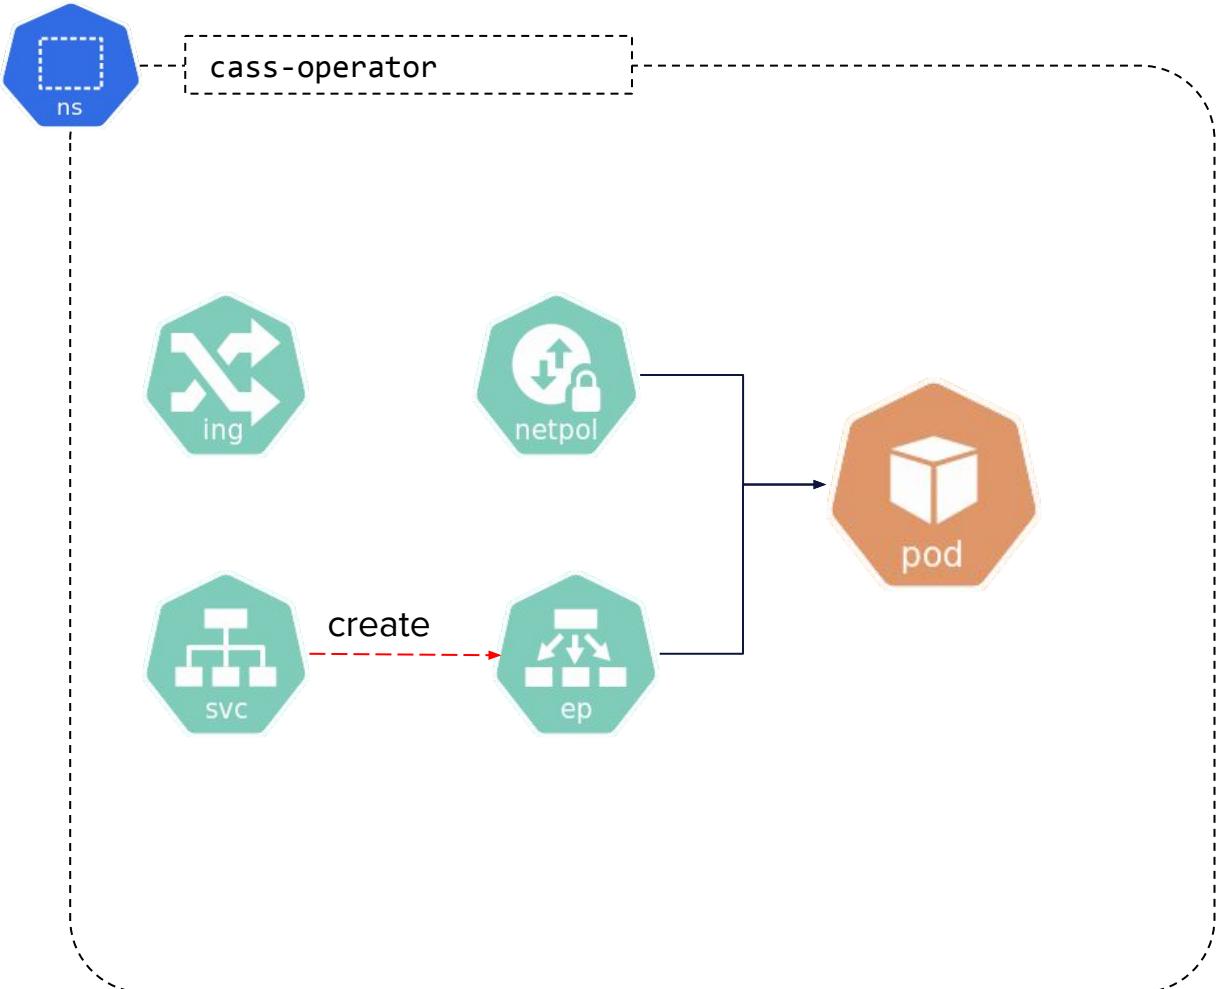
**Service** is a named abstraction of software service with ports to listen on and selector to determine which pods will answer requests.



**EndPoint** is a collection of endpoints that implement the actual service..



**NetworkPolicy**: Describes what network traffic is allowed for a set of Pods.



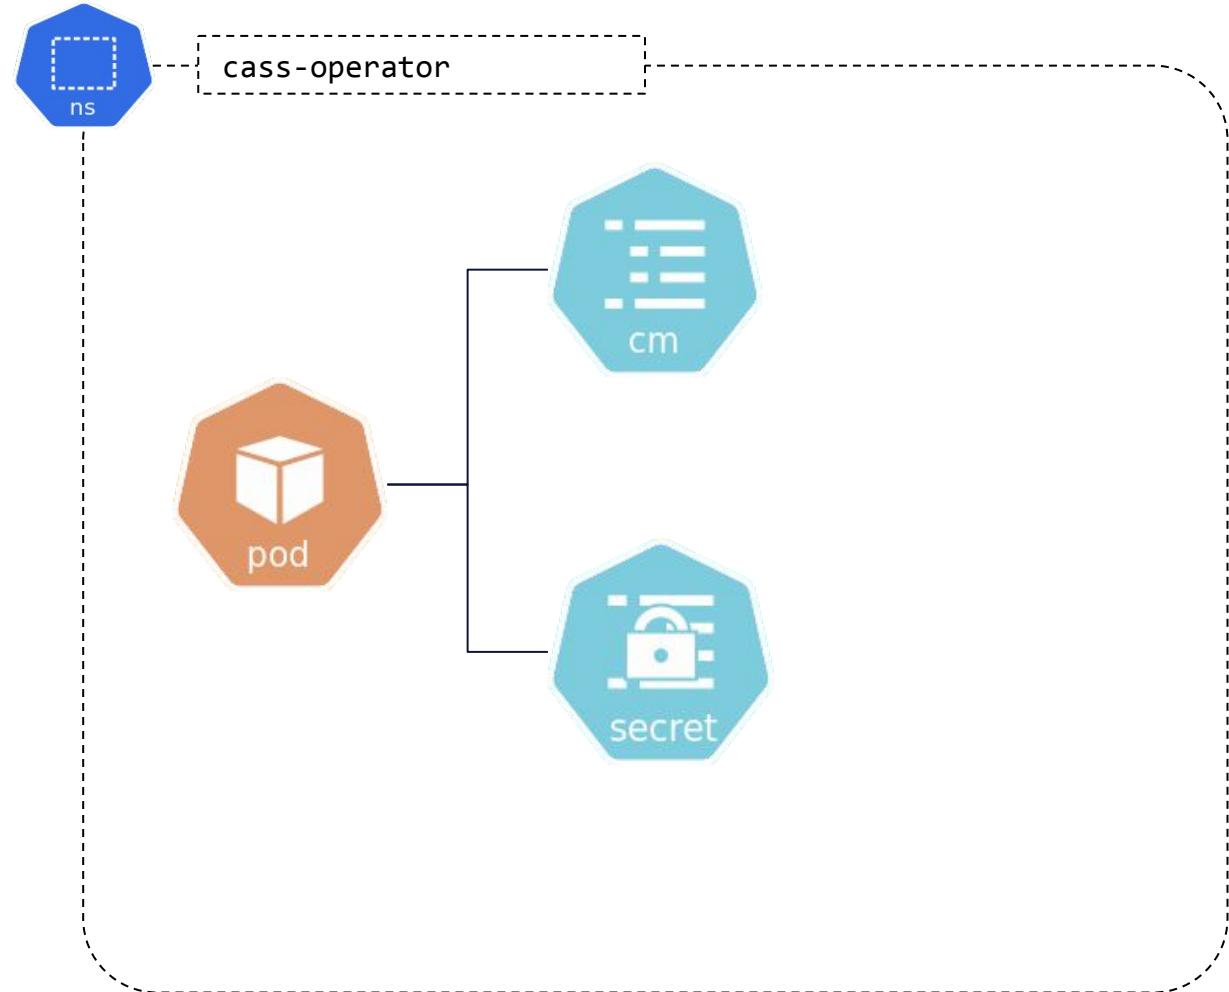
# K8s Primitives : Configuration



**ConfigMap:** ConfigMap holds configuration data for pods to consume..



**Secret:** Secret holds secret data of a certain type..



# K8s Primitives : Custom Resources



**Custom Resource Definition :**  
Extensions of the Kubernetes API.  
Customization making K8s more modular

- **Spec** declares the desired state of a resource
  - **Configuration settings** provided by the user
  - **Default values** expanded by the system
  - **Other properties** initialized by other internal components after resource creation.
- **Status** : describes the object's current, observed state.
  - Kubernetes API server provides a REST API to clients. A Kubernetes object or resource is a REST resource.
  - The status of a Kubernetes resource is typically implemented as a **REST subresource** that can only be modified by internal, system components

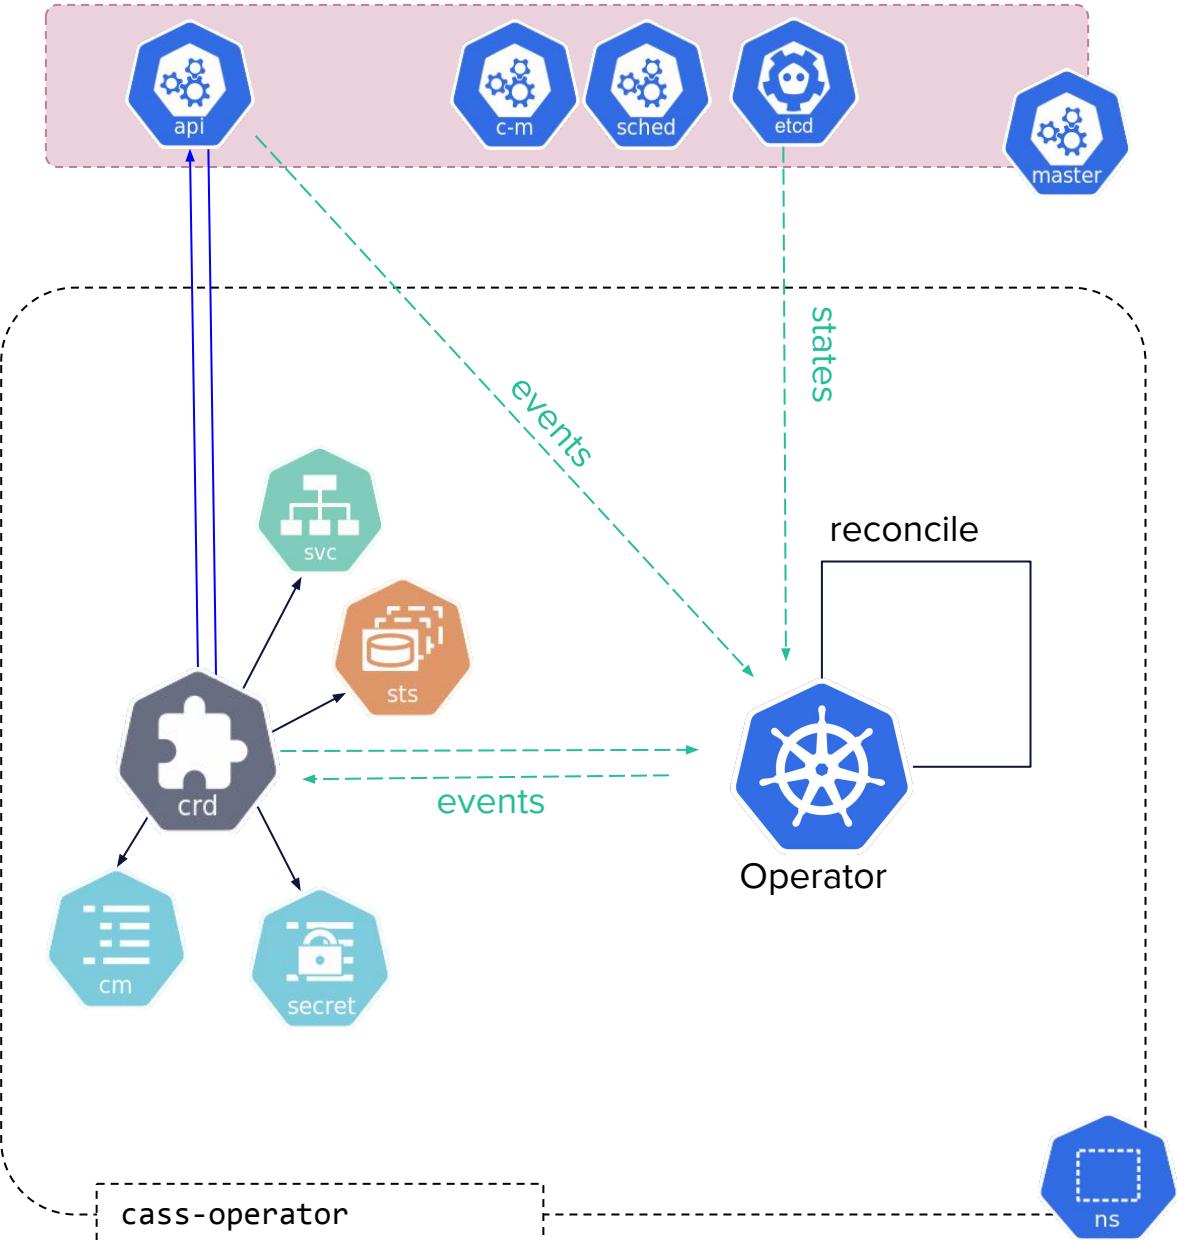
```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name: crontabs.stable.example.com
spec:
  # group name to use for REST API: /apis/<group>/<version>
  group: stable.example.com
  # list of versions supported by this CustomResourceDefinition
  versions:
    - name: v1
      # Each version can be enabled/disabled by Served flag.
      served: true
      # One and only one version must be marked as the storage version.
      storage: true
      schema:
        openAPIV3Schema:
          type: object
          properties:
            spec:
              type: object
              properties:
                cronSpec:
                  type: string
                image:
                  type: string
                replicas:
                  type: integer
                  # either Namespaced or Cluster
                scope: Namespaced
                names:
                [...]
```

# K8s Primitives : Operator

**Building** an application and driving an application on top of Kubernetes, behind Kubernetes APIs

*A Kubernetes Operator helps extend the types of applications that can run on Kubernetes by allowing developers to provide additional knowledge to applications that need to maintain state.” —Jonathan S. Katz*

- **Reconcile** CRD instances which states defined within the “**spec**” attribute.
- **Listen events** and **status evolution** to react accordingly.



# menti.com

# 51 63 695



Available on the iPhone  
App Store

GET IT ON  
Google play

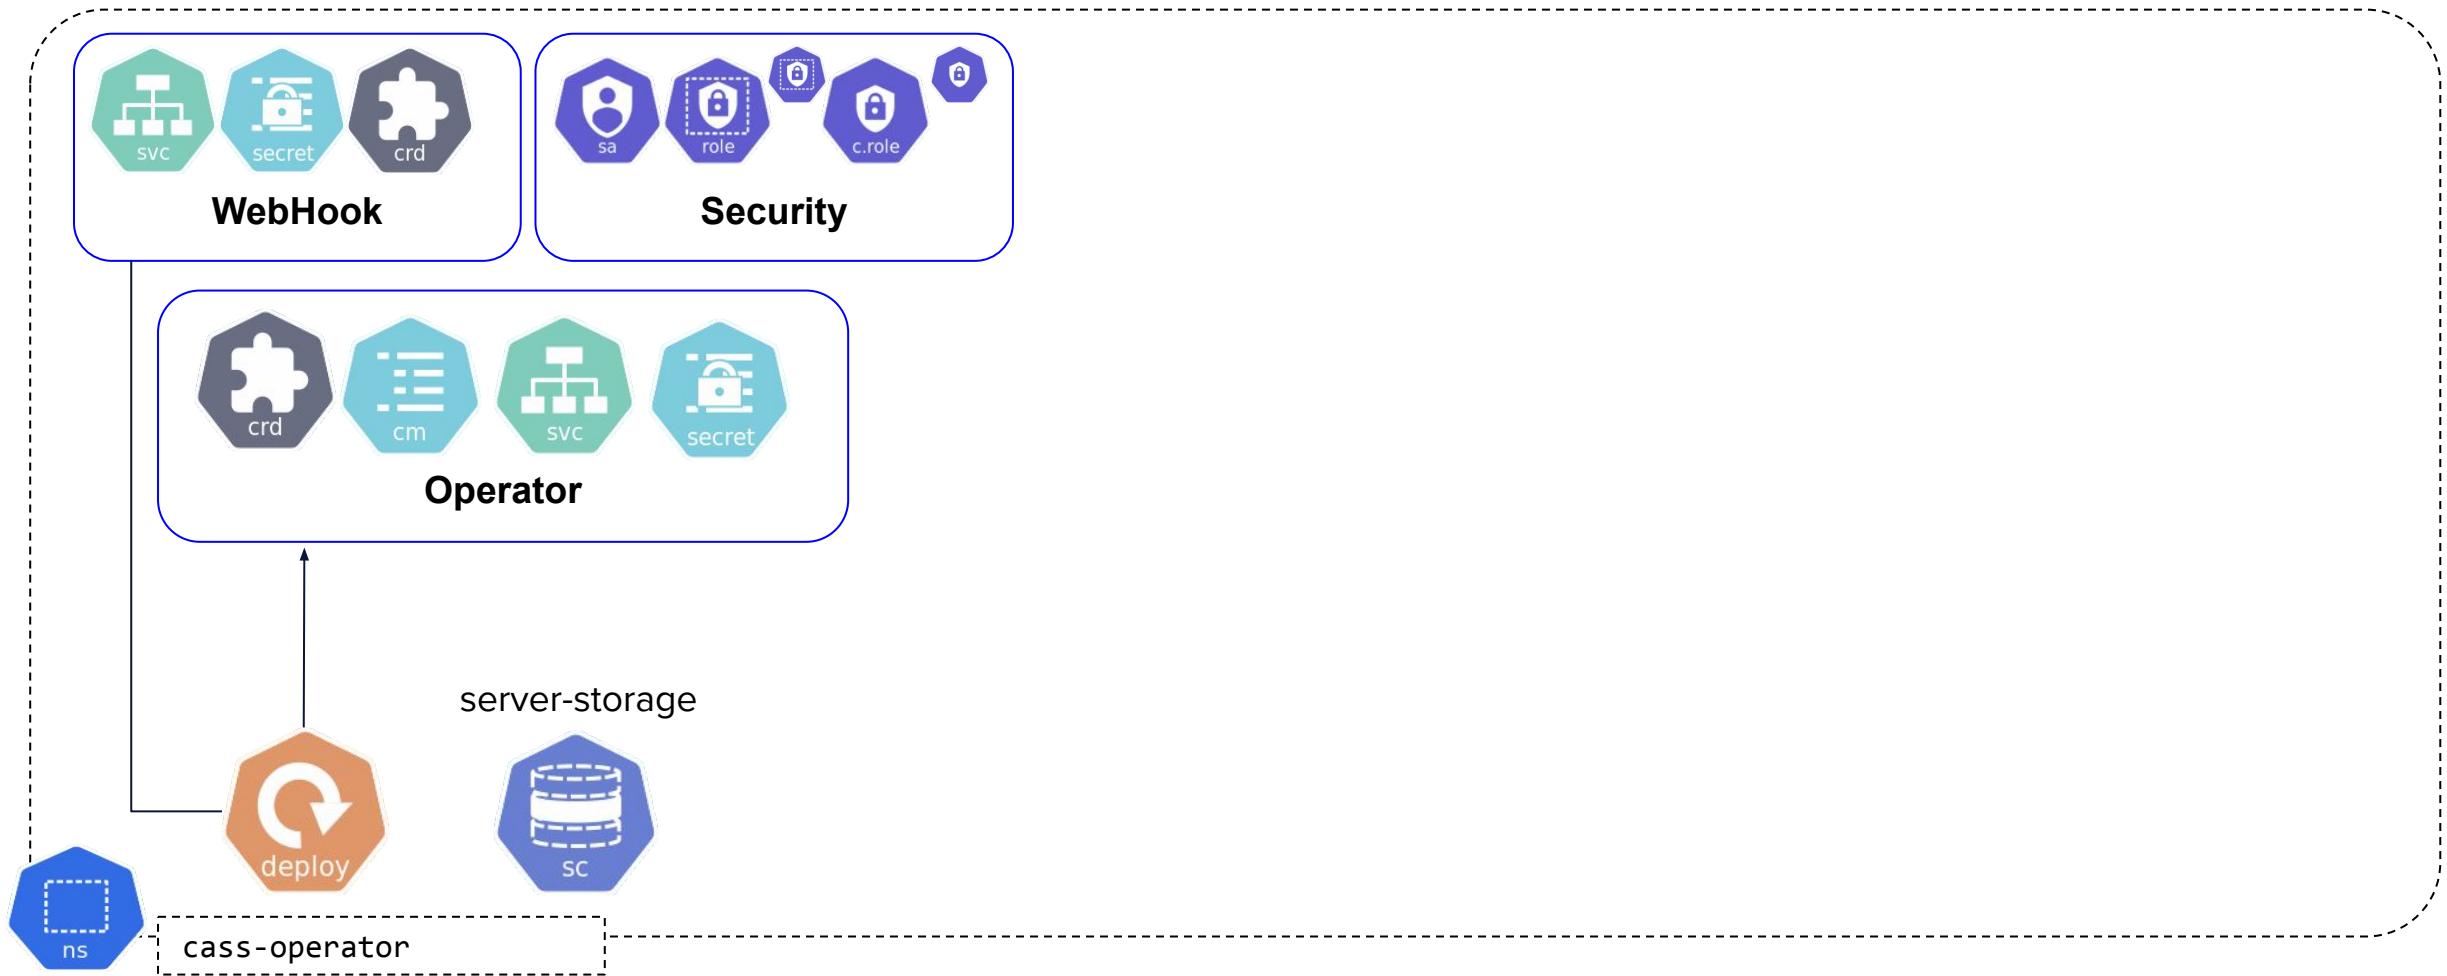
# Apache Cassandra™ with Kubernetes

- 1 Housekeeping and Quizz
- 2 Cassandra Basics
- 3 Kubernetes Operators
- 4 Cass Operator in Deep
- 5 Grafana | Prometheus
- 6 Resources

# Cass Operator : Features

- Proper token **ring initialization**, with only one node bootstrapping at a time
- **Seed node** management -
  - one per rack, or three per datacenter, whichever is more
- Server configuration integrated into the **CassandraDatacenter CRD**
  - Rolling reboot nodes by changing the CRD
  - Store data in a rack-safe way - one replica per cloud AZ
  - Scale up racks evenly with new nodes
  - Replace dead/unrecoverable nodes
- Multi DC clusters (limited to one Kubernetes namespace)

# Installing the Cass Operator Manifest



YAML :

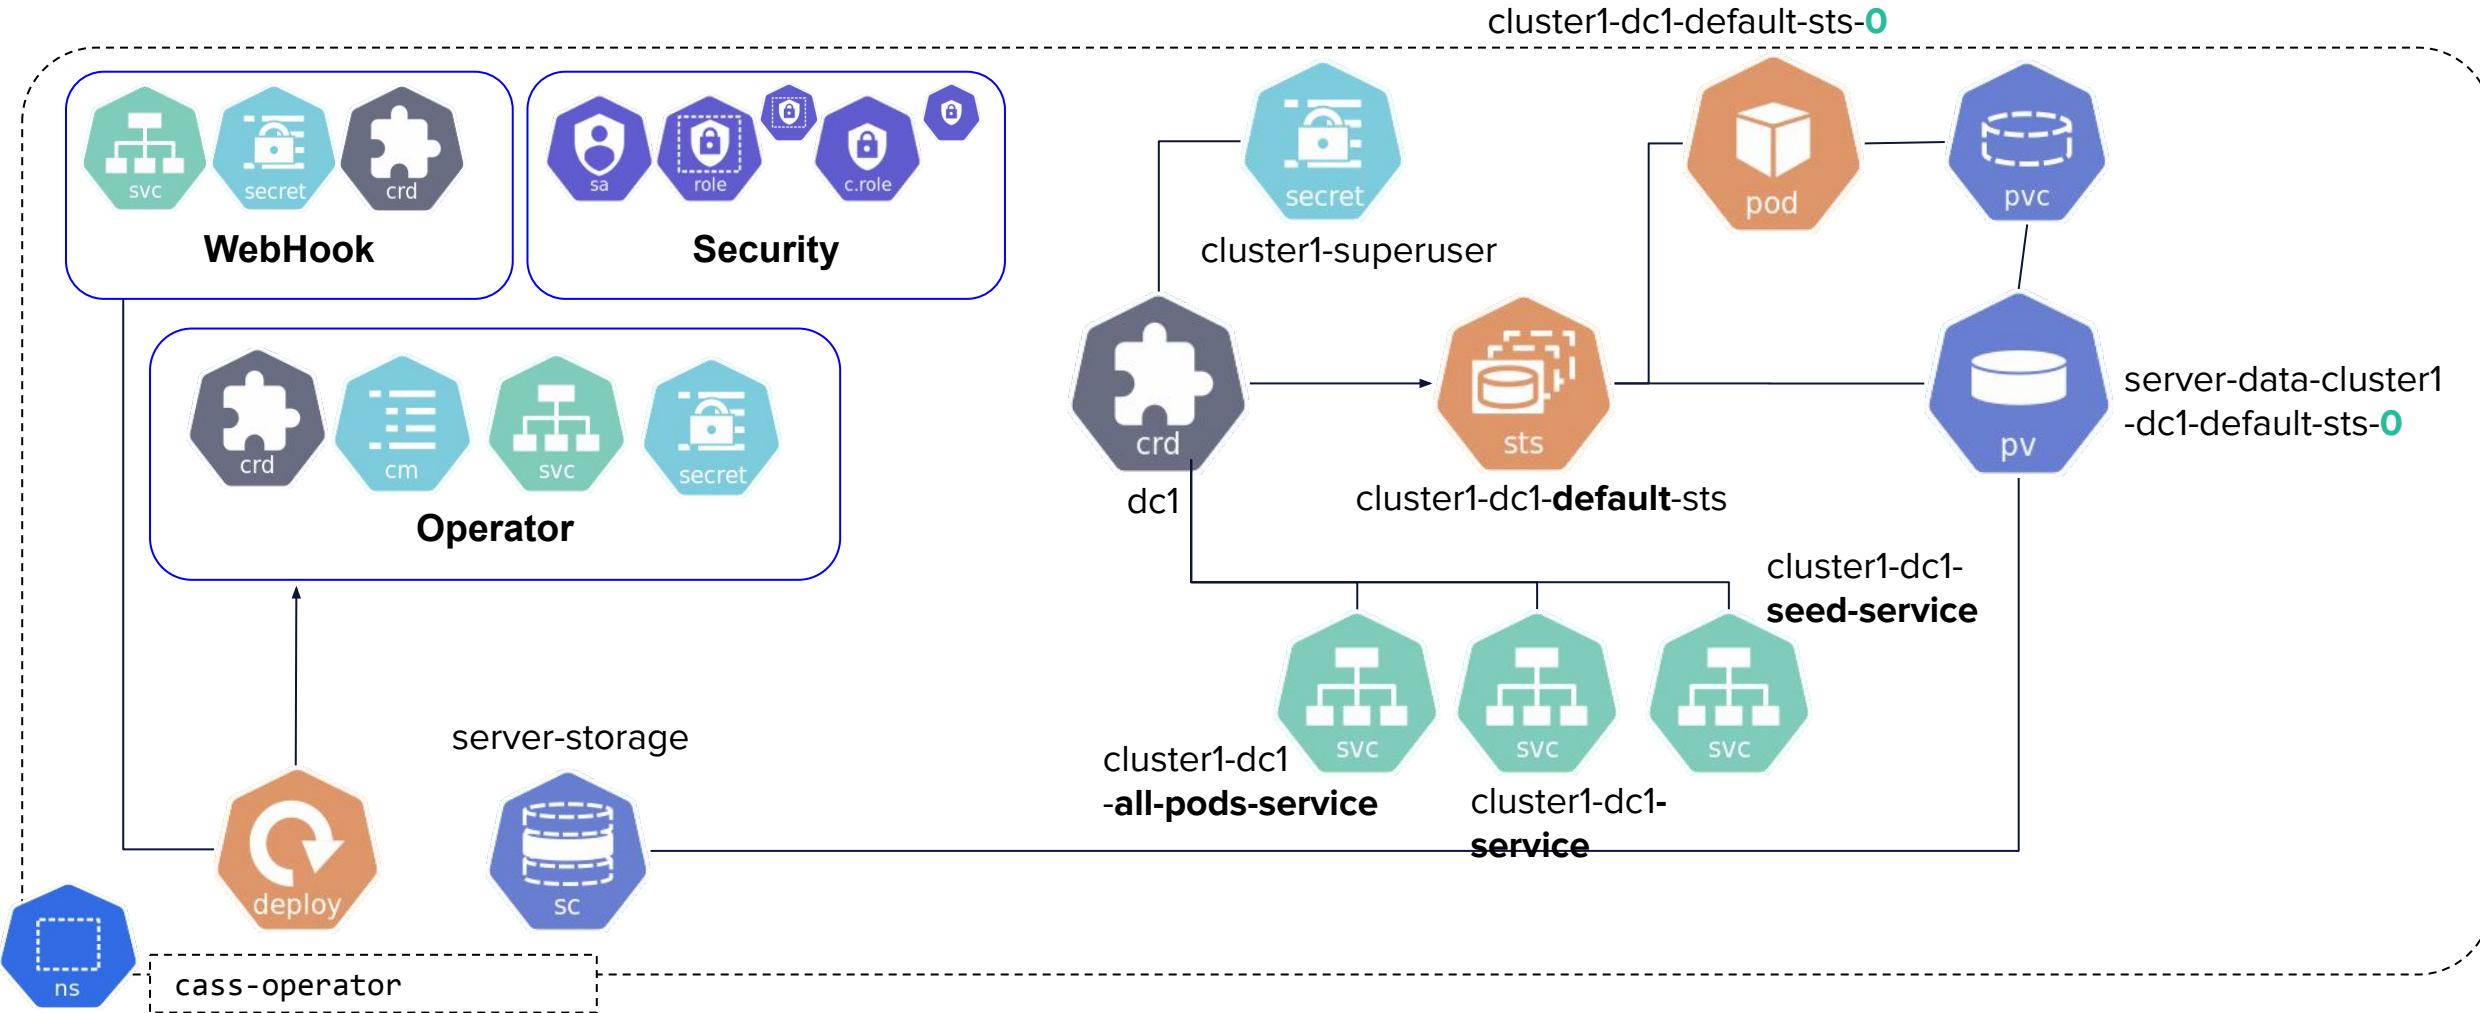
<https://github.com/DataStax-Academy/kubernetes-workshop-online/blob/master/1-cassandra/11-install-cass-operator-v1.1.yaml>

# CRD CassandraDataCenter

```
apiVersion: cassandra.datastax.com/v1beta1
kind: CassandraDatacenter
metadata:
  name: dc1
spec:
  clusterName: cluster1
  serverType: cassandra
  serverVersion: "3.11.6"
  managementApiAuth:
    insecure: {}
  size: 1
  storageConfig:
    cassandraDataVolumeClaimSpec:
      storageClassName: server-storage
      accessModes:
        - ReadWriteOnce
    resources:
      requests:
        storage: 5Gi
config:
  cassandra-yaml:
    authenticator: org.apache.cassandra.auth.PasswordAuthenticator
    authorizer: org.apache.cassandra.auth.CassandraAuthorizer
    role_manager: org.apache.cassandra.auth.CassandraRoleManager
  jvm-options:
    initial_heap_size: "800M"
    max_heap_size: "800M"
```

```
apiVersion: cassandra.datastax.com/v1beta1
kind: CassandraDatacenter
metadata:
  name: multi-rack
spec:
  clusterName: multi-rack
  serverType: cassandra
  serverVersion: 3.11.6
  managementApiAuth:
    insecure: {}
  size: 9
  racks:
    - name: us-east1-b
      zone: us-east1-b
    - name: us-east1-c
      zone: us-east1-c
    - name: us-east1-d
      zone: us-east1-d
  storageConfig:
    cassandraDataVolumeClaimSpec:
      storageClassName: standard
      accessModes:
        - ReadWriteOnce
    resources:
      requests:
        storage: 5Gi
```

# Creating DataCenter dc1



YAML :

<https://github.com/DataStax-Academy/kubernetes-workshop-online/blob/master/1-cassandra/11-install-cass-operator-v1.1.yaml>

# Our Pods



## Cassandra Management API Service

<https://github.com/datastax/management-api-for-apache-cassandra>

### Management API for Apache Cassandra 0.1 OAS3

<https://raw.githubusercontent.com/datastax/management-api-for-apache-cassandra/master/management-api-server/doc/openapi.json>

This is a Restful service for operating Apache Cassandra. You can find out more about the Management API on [Github](#)

Apache 2.0

#### default

`POST /api/v0/ops/auth/role` Creates a new user role

`GET /api/v0/probes/liveness` Indicates whether this service is running

`GET /api/v0/probes/readiness` Indicates whether the Cassandra service is ready to service requests

`GET /api/v0/probes/cluster` Indicated whether the Cassandra cluster is able to achieve the specified consistency

`POST /api/v0/ops/seeds/reload`

`POST /api/v0/ops/keyspace/refresh` Load newly placed SSTables to the system without restart

`POST /api/v0/ops/keyspace/cleanup` Triggers the immediate cleanup of keys no longer belonging to a node. By default, clean all keyspaces

`POST /api/v0/lifecycle/start`

`POST /api/v0/lifecycle/stop`

`POST /api/v0/lifecycle/configure`

`GET /api/v0/lifecycle/pid`

`GET /api/v0/metadata/versions/release` Returns the Cassandra release version

`GET /api/v0/metadata/endpoints` Returns this nodes view of the endpoint states of nodes

`POST /api/v0/ops/node/drain` Drain the node (stop accepting writes and flush all tables)

# Cassandra Management API SideCar

<https://petstore.swagger.io/>

<https://raw.githubusercontent.com/datastax/management-api-for-apache-cassandra/master/management-api-server/doc/openapi.json>

<https://github.com/datastax/management-api-for-apache-cassandra>

The Management API is a sidecar service layer that attempts to build a well supported set of operational actions on Cassandra® nodes that can be administered centrally. It currently works with official Apache Cassandra® 3.11.x and 4.0 via a drop in java agent.

- Lifecycle Management
  - Start Node
  - Stop Node
- Configuration Management (alpha)
  - Change YAML
  - Change jvm-opts
- Health Checks
  - Kubernetes liveness/readiness checks
  - Consistency level checks
- Per node actions
  - All nodetool commands

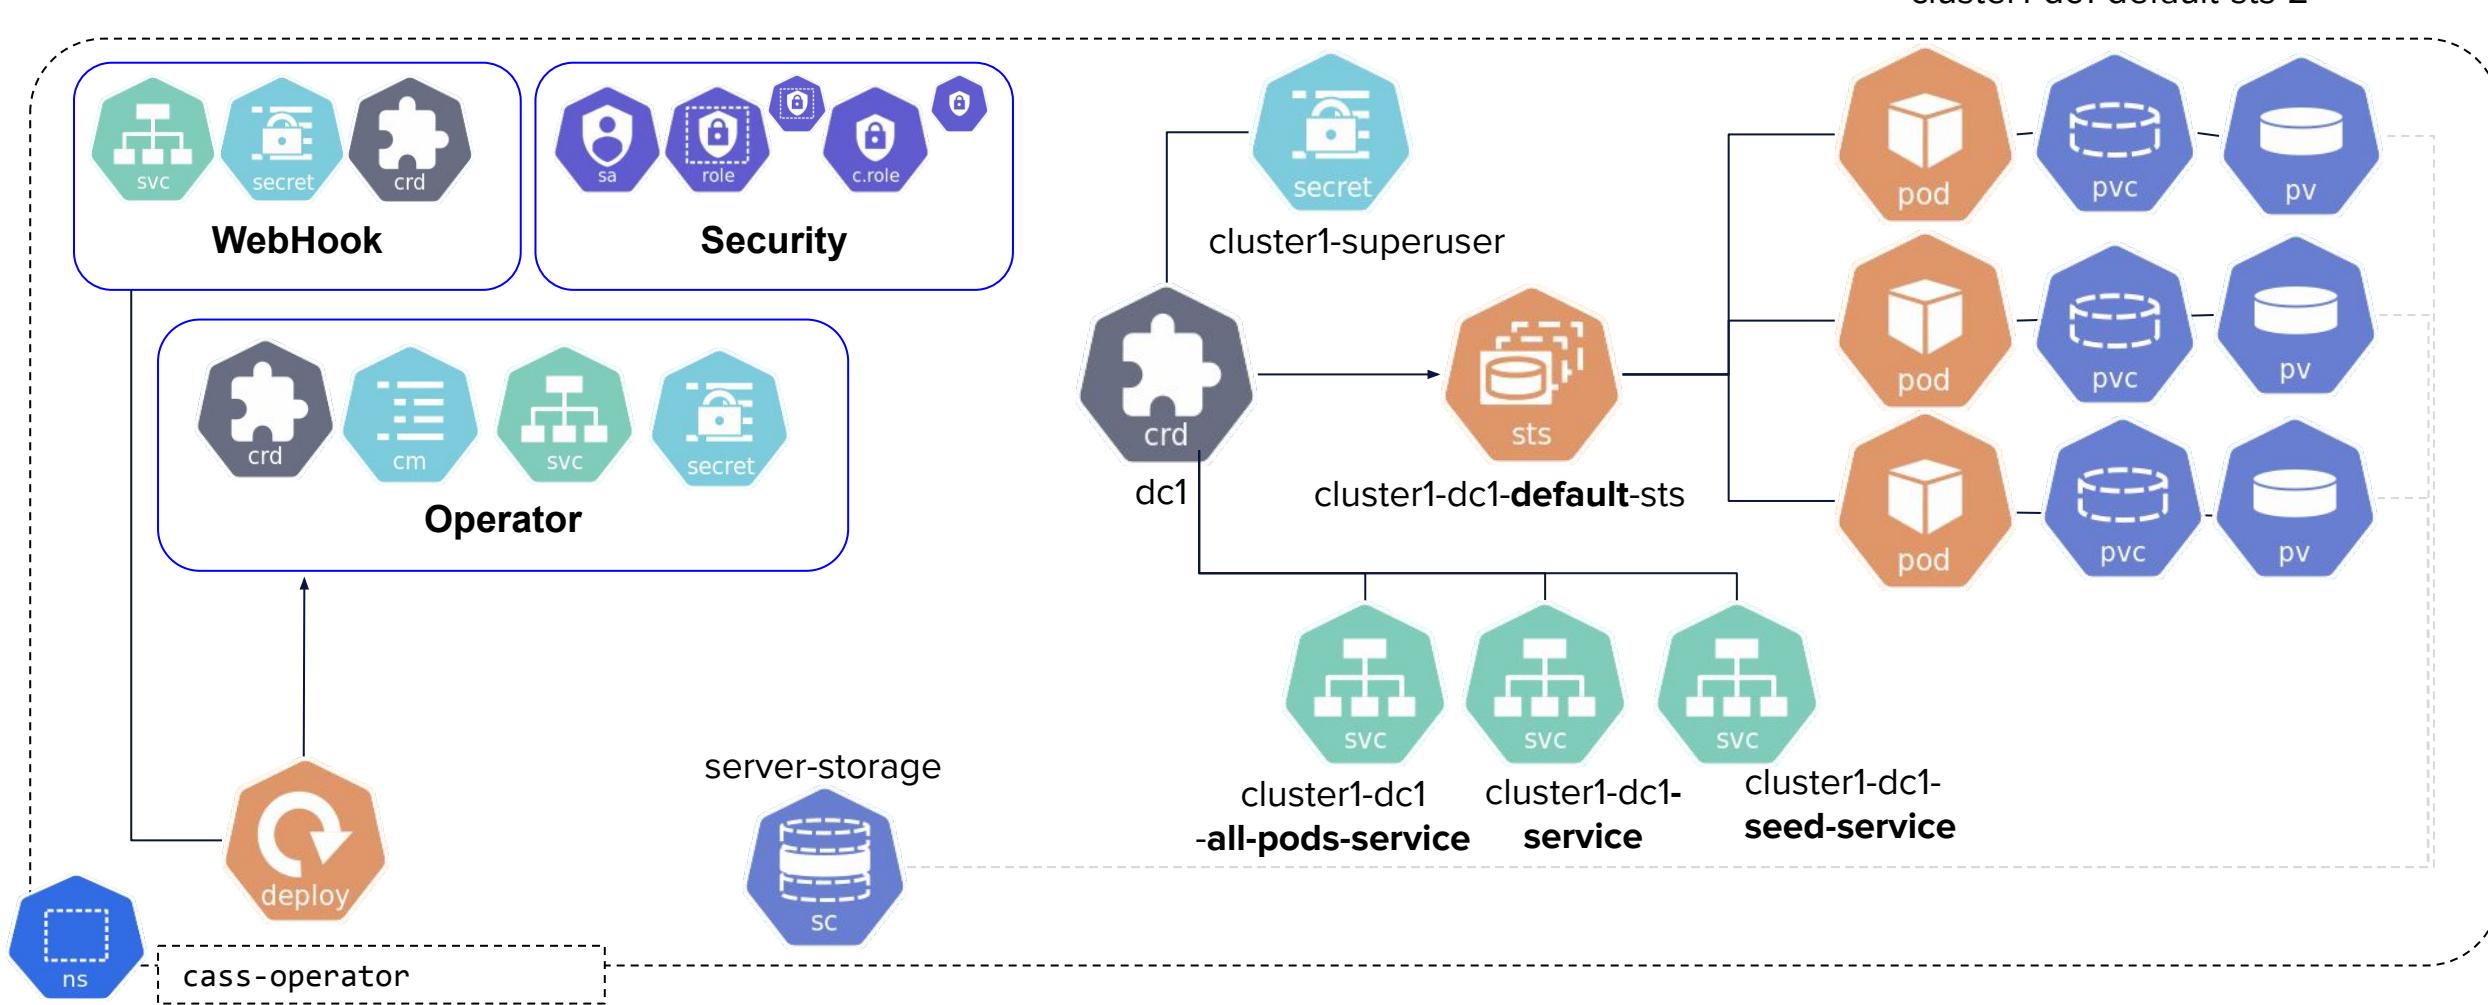
Management API for Apache Cassandra 0.1 OAS3  
<https://raw.githubusercontent.com/datastax/management-api-for-apache-cassandra/master/management-api-server/doc/openapi.json>

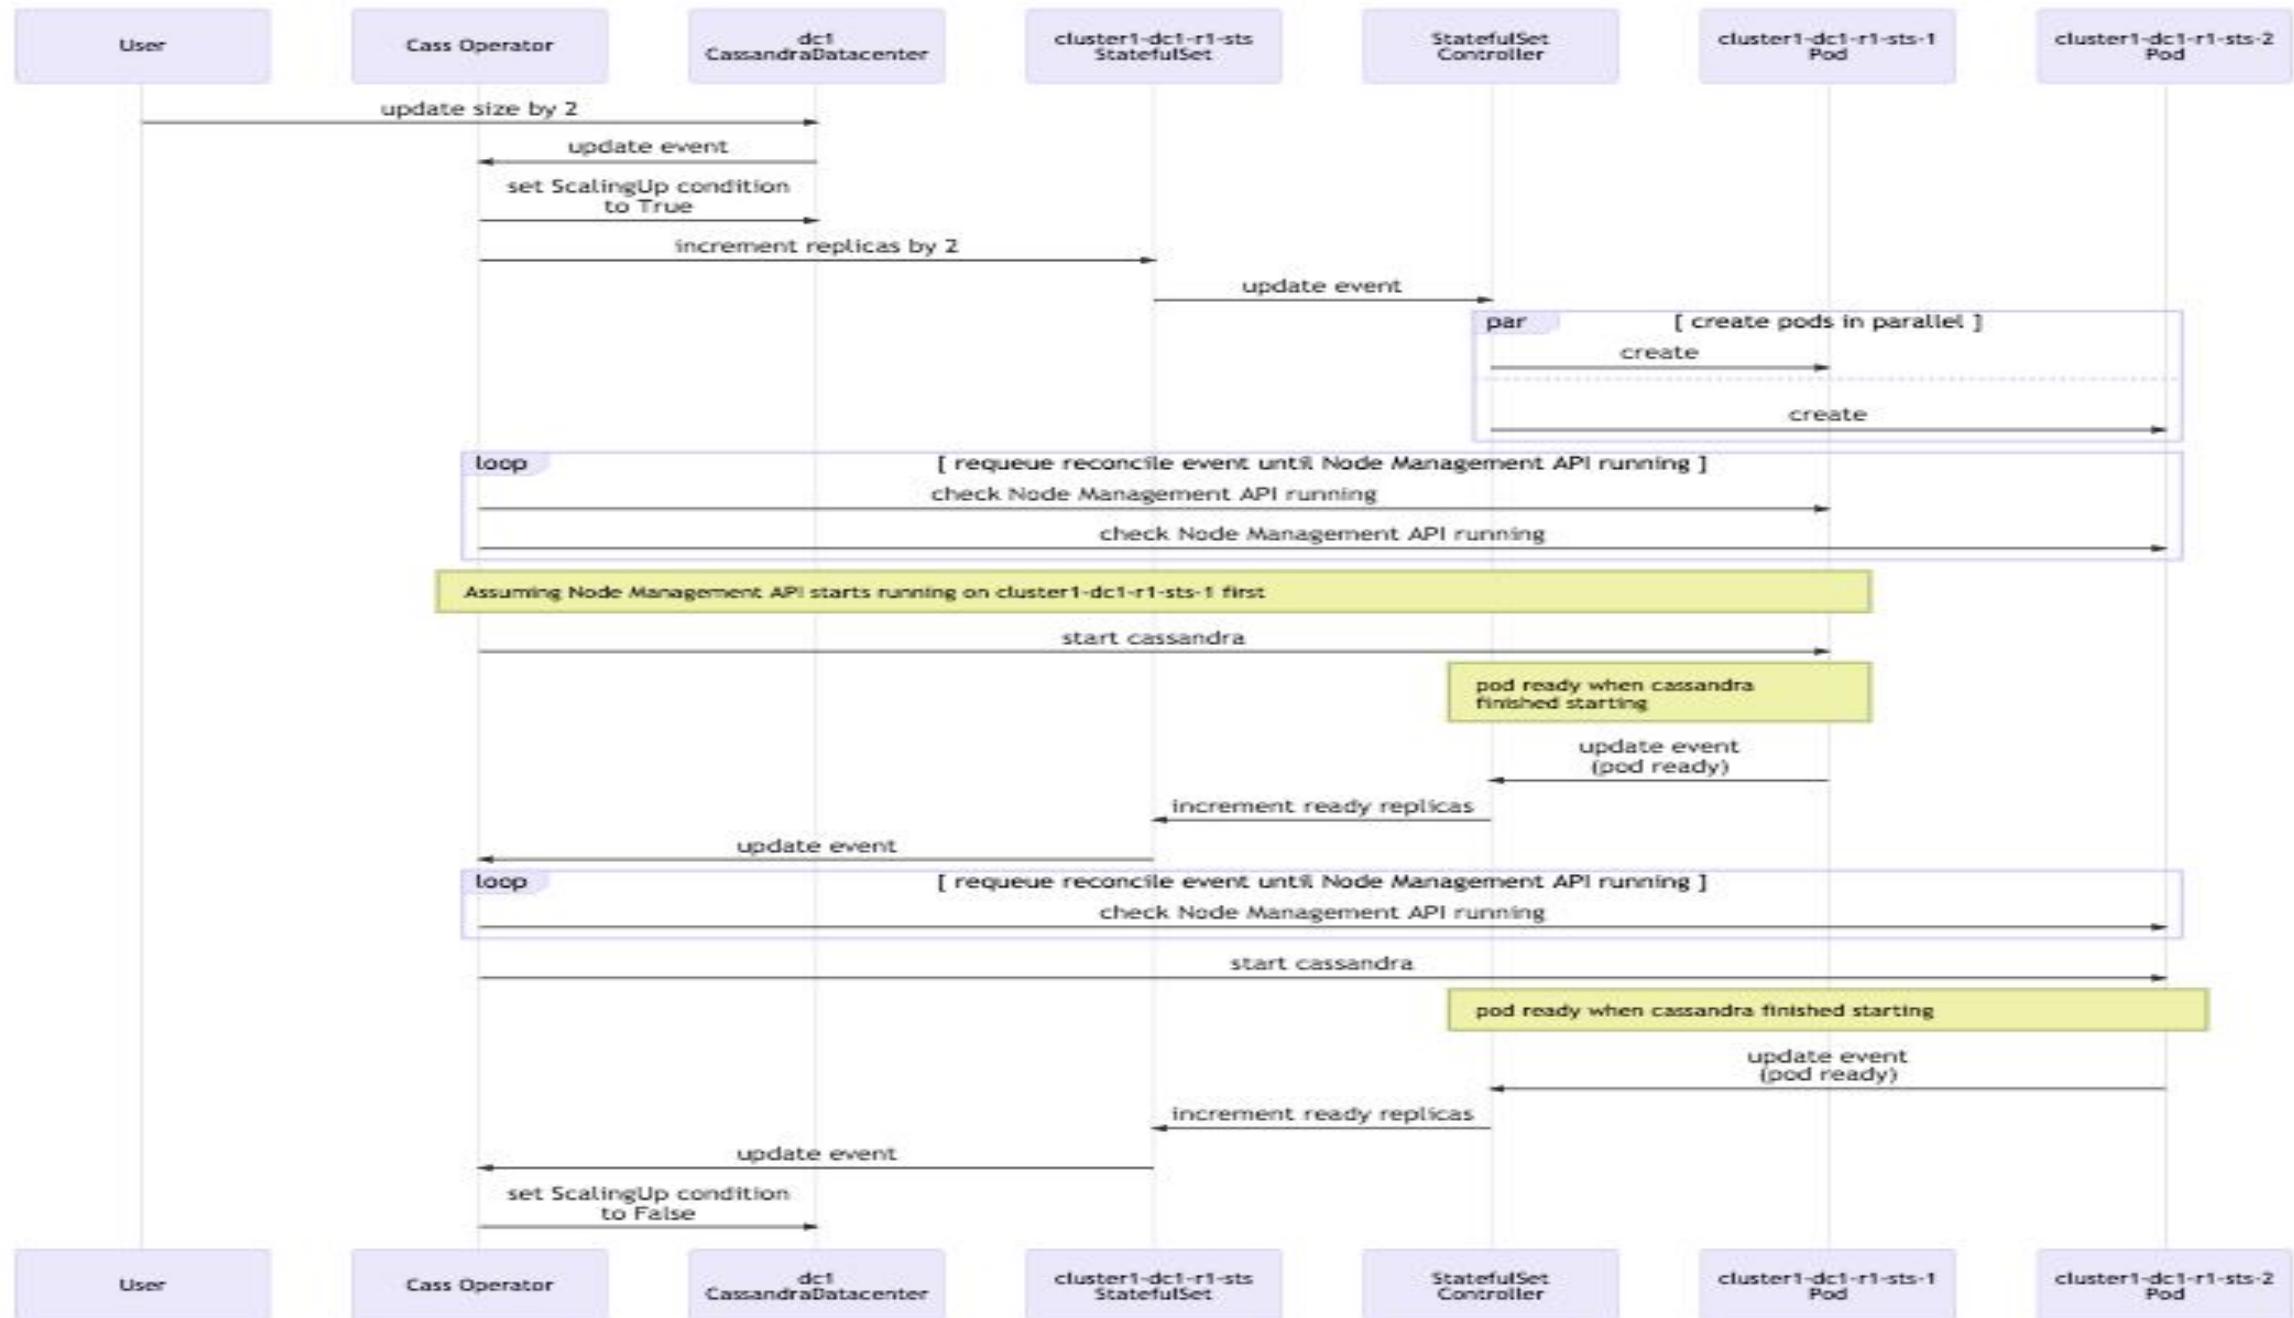
This is a Restful service for operating Apache Cassandra. You can find out more about the Management API on [Github](#).  
Apache 2.0

default

Method	Path	Description
POST	/api/v0/ops/auth/role	Creates a new user role
GET	/api/v0/probes/liveness	Indicates whether this service is running
GET	/api/v0/probes/readiness	Indicates whether the Cassandra service is ready to service requests
GET	/api/v0/probes/cluster	Indicated whether the Cassandra cluster is able to achieve the specified consistency
POST	/api/v0/ops/seeds/reload	
POST	/api/v0/ops/keyspace/refresh	Load newly placed SSTables to the system without restart
POST	/api/v0/ops/keyspace/cleanup	Triggers the immediate cleanup of keys no longer belonging to a node. By default, clean all keyspaces
POST	/api/v0/lifecycle/start	
POST	/api/v0/lifecycle/stop	
POST	/api/v0/lifecycle/configure	
GET	/api/v0/lifecycle/pid	
GET	/api/v0/metadata/versions/release	Returns the Cassandra release version
GET	/api/v0/metadata/endpoints	Returns this nodes view of the endpoint states of nodes
POST	/api/v0/ops/node/drain	Drain the node (stop accepting writes and flush all tables)

# Scale up DataCenter dc1





# Online Workshops



## Section #1 : Cass Operator

[https://github.com/DataStax-Academy/kubecon-cassandra-workshop/  
tree/master/1-cassandra](https://github.com/DataStax-Academy/kubecon-cassandra-workshop/tree/master/1-cassandra)

### 2. Create a single node cluster

Apply this file via `kubectl` and watch the list of pods as the operator deploys them. Completing a deployment may take several minutes per node.

#### 2a. Create the cluster

```
kubectl -n cass-operator apply -f ./1-cassandra/12-cassandra-cluster-1nodes.yaml
```

#### 2b. Watch progression

```
watch kubectl -n cass-operator get pod
```

#### Expected output

NAME	READY	STATUS	RESTARTS	AGE
cass-operator-657cb5c695-q9psl	1/1	Running	0	5m22s
cluster1-dc1-default-sts-0	1/2	Running	0	50s

#### 2c. Execute the command to describe the datacenter

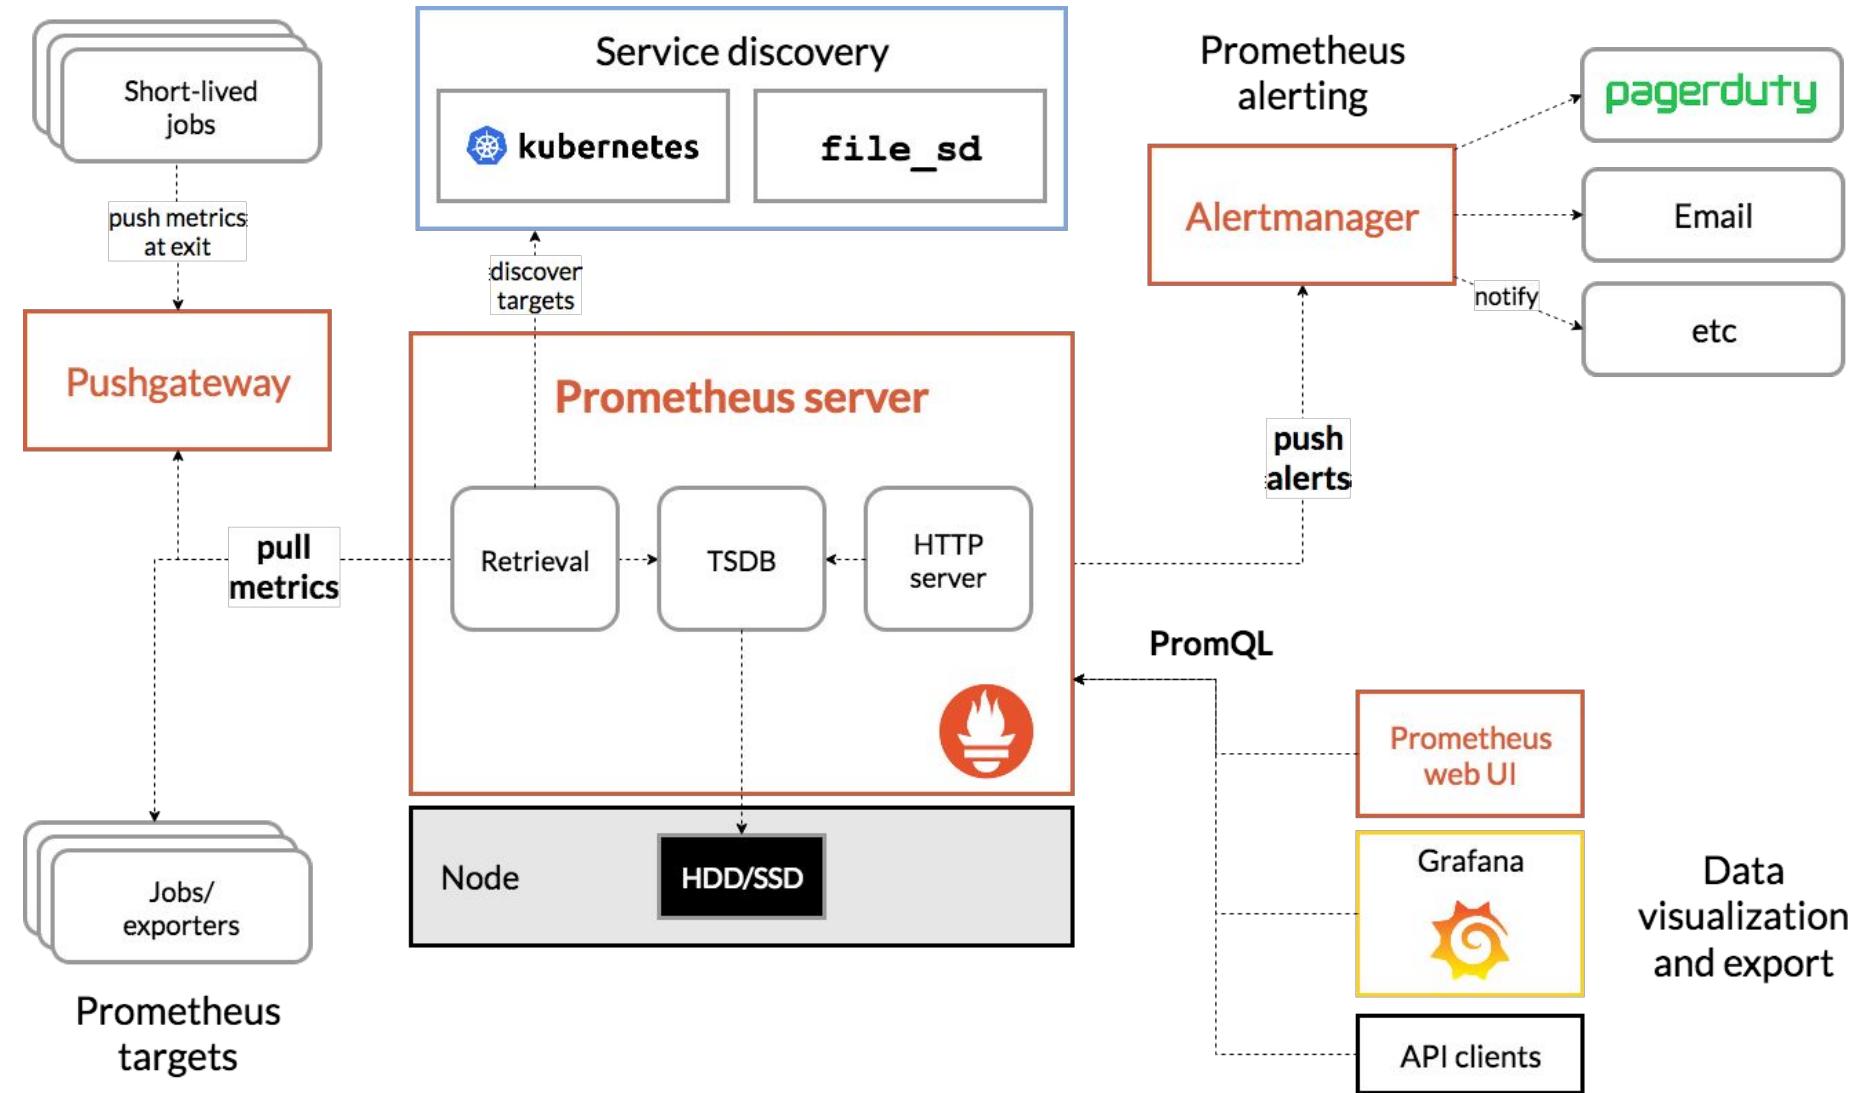
```
kubectl -n cass-operator describe cassdc dc1
```

#### Expected output

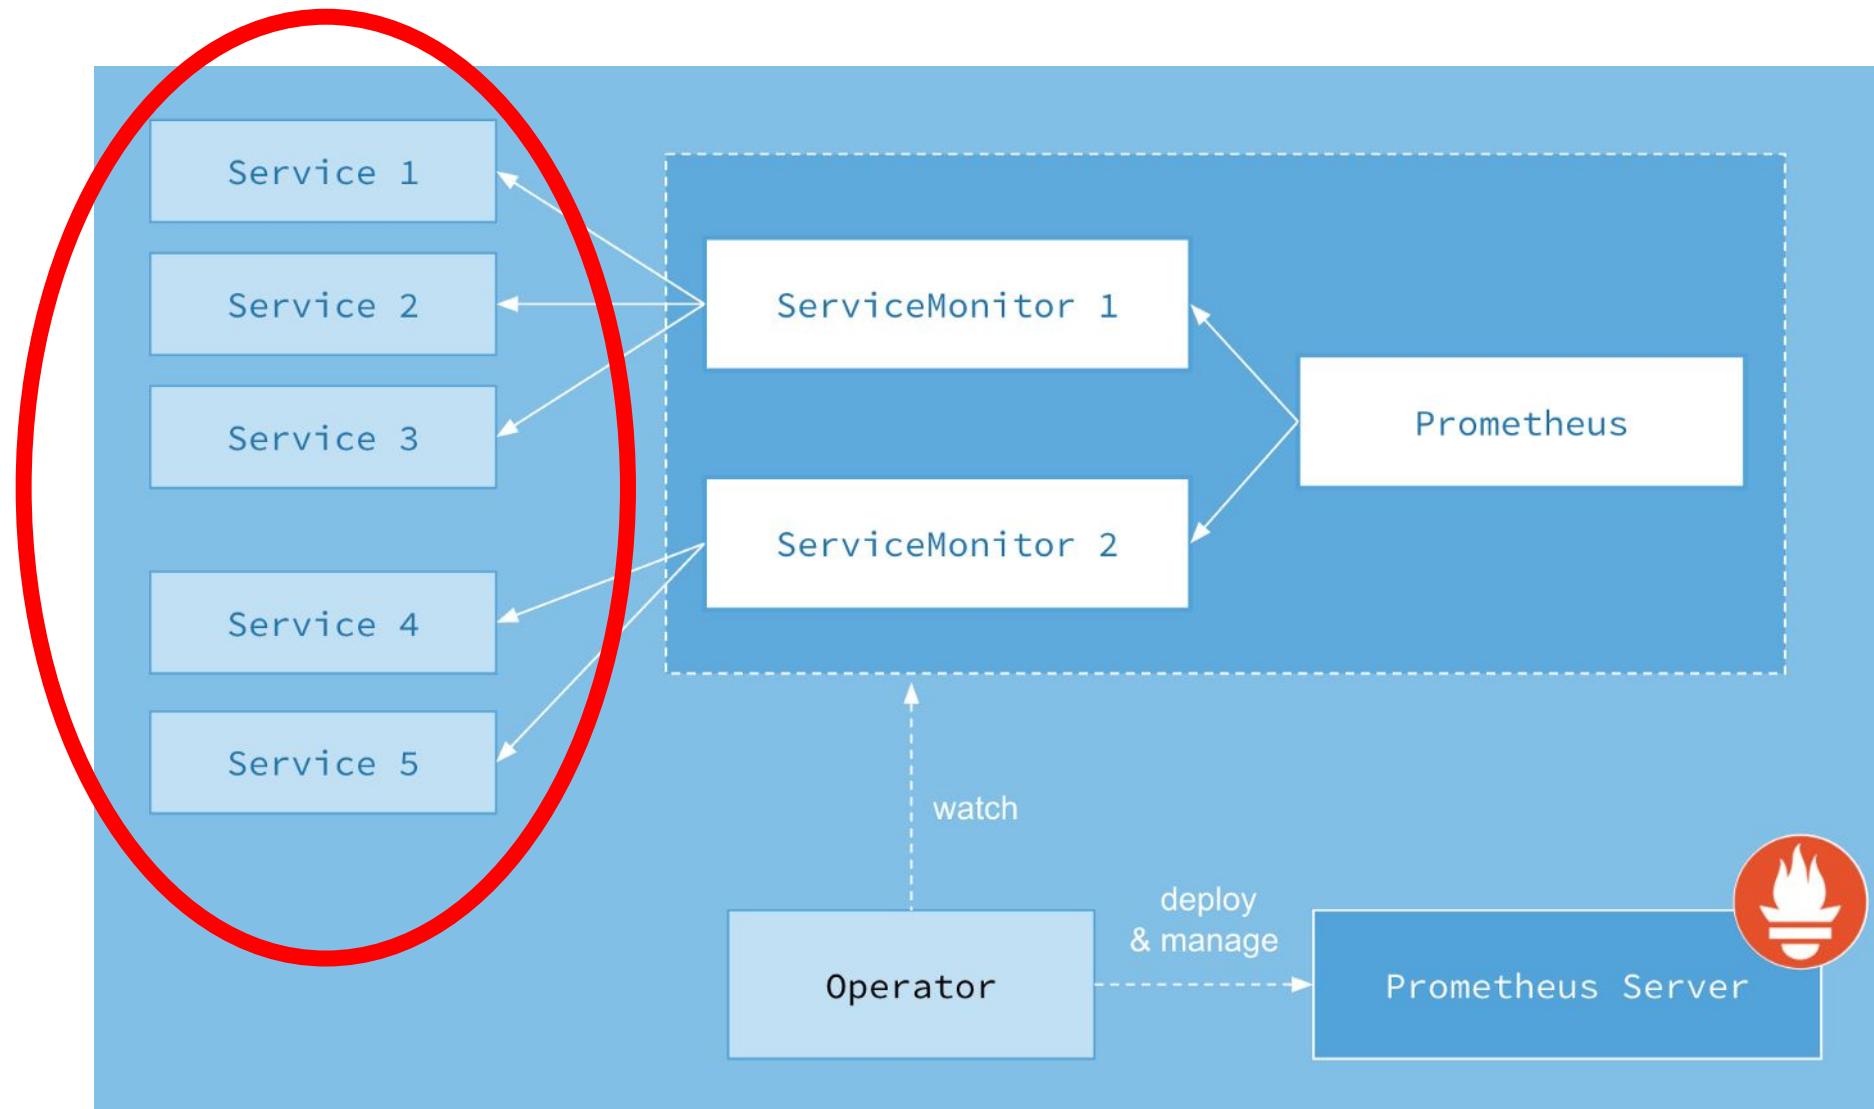
# Apache Cassandra™ with Kubernetes

- 1 Housekeeping and Quizz
- 2 Cassandra Basics
- 3 Kubernetes Operators
- 4 Cass Operator in Deep
- 5 Grafana | Prometheus
- 6 Resources

# Prometheus + Grafana



# Prometheus metrics collection



# DSE Metrics Exporter

## Reference Documentation

[https://docs.datastax.com/en/dse/6.7/dse-dev/datastax\\_enterprise/tools/metricsCollector/mcIntroduction.html](https://docs.datastax.com/en/dse/6.7/dse-dev/datastax_enterprise/tools/metricsCollector/mcIntroduction.html)

- DSE Metrics Collector aggregates DataStax Enterprise (DSE) metrics and integrates with existing monitoring solutions to facilitate problem resolution and remediation.
- **DSE Metrics Collector is built on collectd**, a popular, well-supported, open source metric collection agent. With [over 90 plugins](#), you can tailor the solution to collect metrics most important to your organization.
- When [DSE Metrics Collector](#) is enabled, DSE sends metrics and other structured events to DSE Metrics Collector.

`/etc/dse/collectd.conf.tpl`

```
LoadPlugin load
LoadPlugin memory
LoadPlugin swap
LoadPlugin uptime
LoadPlugin processes
LoadPlugin tcpconns
```

# MCAC Metrics Collector for Apache Cassandra

## Reference Documentation

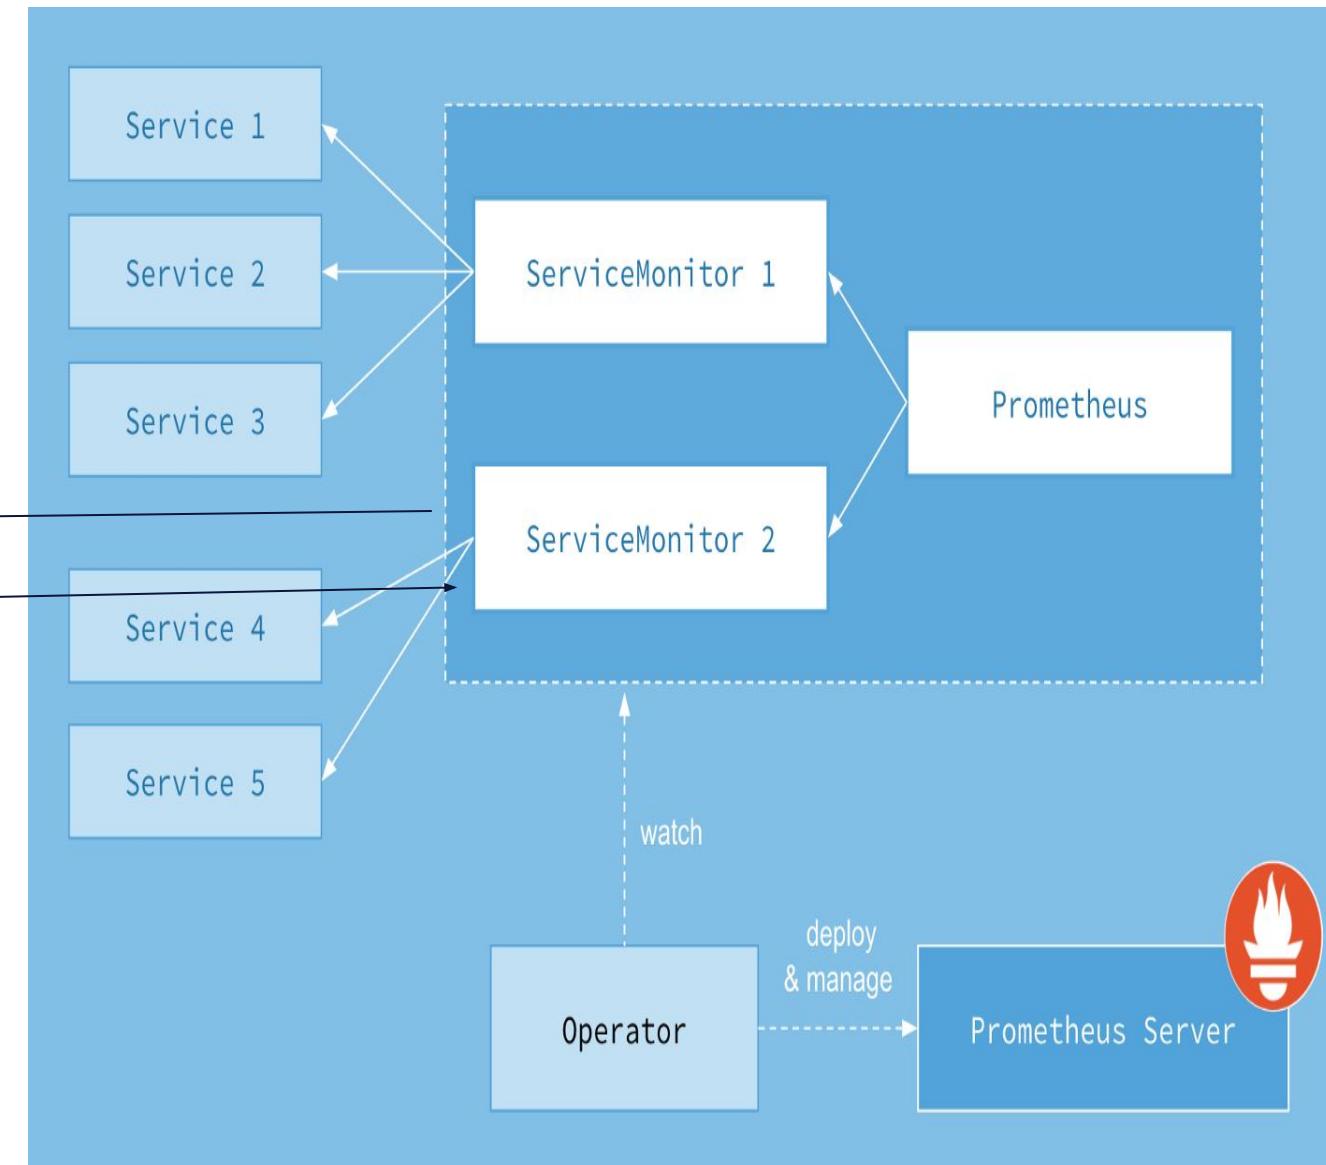
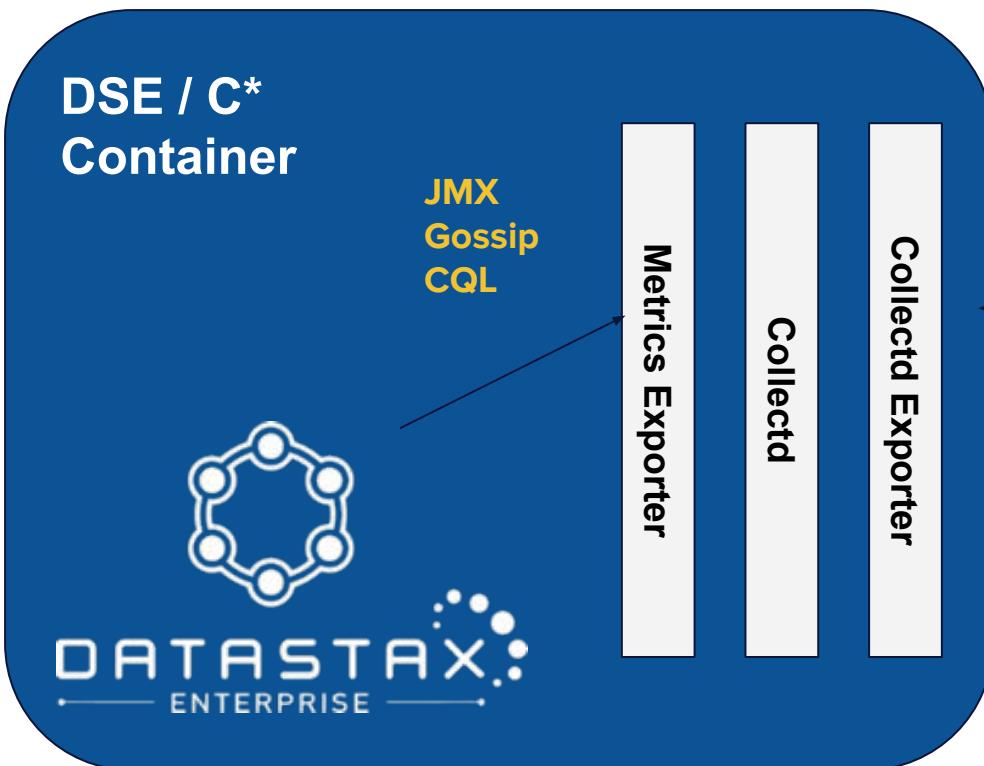
<https://github.com/datastax/metric-collector-for-apache-cassandra>

- MCAC Metrics Collector aggregates Apache Cassandra metrics and integrates with existing monitoring solutions to facilitate problem resolution and remediation.
- MCAC is built on **collectd**, a popular, well-supported, open source metric collection agent. With **over 90 plugins**, you can tailor the solution to collect metrics most important to your organization.
- Easily added to Cassandra nodes as a java agent, Apache Cassandra sends metrics and other structured events to collectd over a local unix socket.

`/etc/dse/collectd.conf tmpl`

```
LoadPlugin load
LoadPlugin memory
LoadPlugin swap
LoadPlugin uptime
LoadPlugin processes
LoadPlugin tcpconns
```

# All Together



# Online Workshops



## Section #2 : Grafana Prometheus

[https://github.com/DataStax-Academy/kubecon-cassandra-workshop/  
tree/master/2-prometheus\\_grafana](https://github.com/DataStax-Academy/kubecon-cassandra-workshop/tree/master/2-prometheus_grafana)



# Apache Cassandra™ with Kubernetes

1

**Housekeeping and Quizz**

2

**Cassandra Basics**

3

**Kubernetes Operators**

4

**Cass Operator in Deep**

5

**Grafana | Prometheus**

6

**Resources**

# Developer Resources

**LEARN**

Join [academy.datastax.com](https://academy.datastax.com)  
Free online courses - Cassandra certifications

**ASK/SHARE**

Join [community.datastax.com](https://community.datastax.com)  
Ask/answer community user questions - share your expertise

**SUBSCRIBE !**

Follow us @DataStaxDevs  
We are on Twitter - Youtube - Twitch!

Slides and code for this course are available at  
<https://github.com/DataStax-Academy/kubecon-cassandra-workshop>

**REVIEW**

Accelerate

<https://www.datastax.com/accelerate/why-tomorrows-cassandra-deployments-will-be-on-kubernetes>



# Training Courses at DataStax Academy

- Free self-paced DSE 6 courses
  - [DS201: Foundations of Apache Cassandra™](#)
  - [DS210: Operations with Apache Cassandra™](#)
  - [DS220: Practical Application Data Modeling with Apache Cassandra™](#)
  - [DS330: DataStax Enterprise 6 Graph](#)
  - [DS332: DataStax Enterprise 6 Graph Analytics](#)
  - **And MUCH MORE!**



# Cassandra The Definitive Guide

<https://www.datastax.com/resources/ebook/oreilly-cassandra-definitive-guide>

O'REILLY®

# Cassandra The Definitive Guide

Distributed Data  
at Web Scale

Third  
Edition



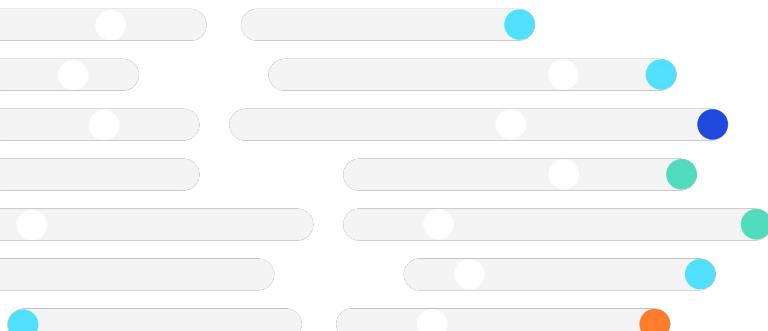
# menti.com

# 51 63 695



Available on the iPhone  
**App Store**

GET IT ON  
**Google play**



**Thank you**

**SURVEY**